

Tackling The Challenges of Big Data

Big Data Storage

Distributed Computing Platforms

Matei Zaharia

Assistant Professor
Massachusetts Institute of Technology



PROFESSIONAL
EDUCATION



Tackling The Challenges of Big Data

Big Data Storage

Distributed Computing Platforms

Introduction

Matei Zaharia

Assistant Professor
Massachusetts Institute of Technology



PROFESSIONAL
EDUCATION



Motivation


- **Large datasets are inexpensive to collect, but require high parallelism to process**
 - 1 TB of disk space = \$50
 - Reading 1 TB from disk = 6 hours
 - **Reading 1 TB from 1000 disks = 20 seconds**
- **Not only queries, but *all* computations will need to scale (data loading, complex analytics, etc)**
- **How do we program large clusters?**



Tackling the Challenges of Big Data

© 2014 Massachusetts Institute of Technology





Traditional Network Programming

- **Message-passing between nodes**
- **Very difficult to use at scale:**
 - How to split problem across nodes?
***Must consider network, data locality**
 - How to deal with failures?
***1 server fails every 3 years => 10K nodes see 10 faults per day**
 - Even worse: stragglers (node is not failed, but slow)

Almost nobody directly uses this model!

Data-Parallel Models

- **Restrict the programming interface so that the system can do more automatically**
- **"Here's an operation, run it on all of the data"**
 - I don't care *where* it runs (you schedule that)
 - In fact, feel free to run it *twice* on different nodes
- **Biggest example: MapReduce**

Software in this Space

Hadoop Dryad Spark
Hive Pig DryadLINQ Shark
Pregel Dremel
Giraph Storm
Impala
GraphLab Tez Samza



Tackling the Challenges of Big Data

© 2014 Massachusetts Institute of Technology



Applications

- **Extract, transform and load ("ETL")**
- **Web indexing (Google)**
- **Spam filtering (Yahoo!)**
- **Product recommendation (Netflix)**
- **Ad-hoc queries (Facebook)**
- **Fraud detection**



Tackling the Challenges of Big Data

© 2014 Massachusetts Institute of Technology



This Lecture

- **Challenges in large-scale environments**
- **MapReduce model**
- **Limitations and extensions of MapReduce**
- **Other types of platforms**



Tackling the Challenges of Big Data

© 2014 Massachusetts Institute of Technology



Tackling The Challenges of Big Data

Big Data Storage

Distributed Computing Platforms

Introduction

THANK YOU



PROFESSIONAL
EDUCATION



Tackling The Challenges of Big Data

Big Data Storage

Matei Zaharia

Assistant Professor

Massachusetts Institute of Technology



PROFESSIONAL
EDUCATION



Tackling The Challenges of Big Data

Big Data Storage

Distributed Computing Platforms

Large-Scale Computing Environments

Matei Zaharia

Assistant Professor

Massachusetts Institute of Technology



PROFESSIONAL
EDUCATION



Goals of Large-Scale Data Platforms

- **Scalability in number of nodes:**
 - Parallelize I/O to quickly scan large datasets
- **Cost-efficiency:**
 - Commodity nodes (cheap, but unreliable)
 - Commodity network (low bandwidth)
 - Automatic fault-tolerance (fewer admins)
 - Easy to use (fewer programmers)




Tackling the Challenges of Big Data

© 2014 Massachusetts Institute of Technology



Typical Hadoop Cluster



- **40 nodes/rack, 1000-4000 nodes in cluster**
- **10 Gbps bandwidth in rack, 40 Gbps out of rack**
- **Node specs (Cloudera):**
8-16 cores, 64-512 GB RAM, 12×2 TB disks



Tackling the Challenges of Big Data

© 2014 Massachusetts Institute of Technology



Typical Hadoop Cluster



Source: Yahoo!



Tackling the Challenges of Big Data

© 2014 Massachusetts Institute of Technology



Challenges of Cluster Environment

- **Cheap nodes fail, especially when you have many**
 - Mean time between failures for 1 node = 3 years
 - MTBF for 1000 nodes = 1 day
 - **Solution:** Build fault tolerance into system
- **Commodity network = low bandwidth**
 - **Solution:** Push computation to the data
- **Programming distributed systems is hard**
 - **Solution:** Data-parallel model: users write map/reduce functions, system handles work distribution and faults



Tackling the Challenges of Big Data

© 2014 Massachusetts Institute of Technology



Typical Software Components

- **Distributed file system (e.g. Hadoop's HDFS)**
 - Single namespace for entire cluster
 - Replicates data 3x for fault-tolerance
- **MapReduce system (e.g. Hadoop MapReduce)**
 - Runs jobs submitted by users
 - Manages work distribution & fault-tolerance
 - Colocated with file system




Tackling the Challenges of Big Data

© 2014 Massachusetts Institute of Technology



Hadoop Distributed File System

- Files split into blocks
- Blocks replicated across several nodes (often 3)
- Master stores metadata (file names, locations, ...)
- Optimized for large files, sequential reads



Tackling the Challenges of Big Data

© 2014 Massachusetts Institute of Technology



Tackling The Challenges of Big Data

Big Data Storage
Distributed Computing Platforms
Large-Scale Computing Environments

THANK YOU



PROFESSIONAL
EDUCATION



Tackling The Challenges of Big Data

Big Data Storage

Matei Zaharia

Assistant Professor
Massachusetts Institute of Technology



PROFESSIONAL
EDUCATION



Tackling The Challenges of Big Data

Big Data Storage
Distributed Computing Platforms
Introduction to MapReduce

Matei Zaharia

Assistant Professor
Massachusetts Institute of Technology



PROFESSIONAL
EDUCATION



MapReduce Programming Model

- A MapReduce job consists of two user functions
- Both operate on key-value records

- **Map function:**

$(K_{in}, V_{in}) \rightarrow \text{list}(K_{inter}, V_{inter})$

- **Reduce function:**

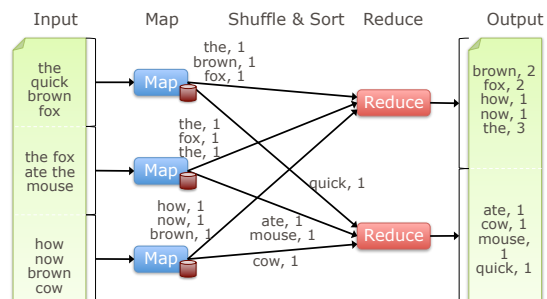
$(K_{inter}, \text{list}(V_{inter})) \rightarrow \text{list}(K_{out}, V_{out})$

Example: Word Count

```
def map(line):  
    for word in line.split():  
        output(word, 1)
```

```
def reduce(key, values):  
    output(key, sum(values))
```

Word Count Execution



MapReduce Execution

- **Automatically split work into many small tasks**
- **Send map tasks to nodes based on data locality**
 - Typically, files are replicated so that there are 3 copies of each block
- **Load-balance dynamically as tasks finish**



Tackling the Challenges of Big Data

© 2014 Massachusetts Institute of Technology



Fault Recovery

1. **If a task crashes:**
 - Retry on another node
 - ***OK for a map because it had no dependencies**
 - ***OK for reduce because map outputs are on disk**
 - If the same task repeatedly fails, end the job

Requires user code to be
deterministic and idempotent



Tackling the Challenges of Big Data

© 2014 Massachusetts Institute of Technology



Fault Recovery

2. **If a node crashes:**
 - Relaunch its current tasks on other nodes
 - Relaunch any maps the node previously ran
 - ***Necessary because their output files were lost along with the crashed node**



Tackling the Challenges of Big Data

© 2014 Massachusetts Institute of Technology



Fault Recovery

3. If a task is going slowly (straggler):

- Launch second copy of task on another node
- Take the output of whichever copy finishes first, and cancel the other one

By offering a data-parallel model, MapReduce can handle many distribution issues automatically



Tackling the Challenges of Big Data

© 2014 Massachusetts Institute of Technology



Tackling The Challenges of Big Data

Big Data Storage

Distributed Computing Platforms

Introduction to MapReduce

THANK YOU



Tackling The Challenges of Big Data

Big Data Storage

Matei Zaharia

Assistant Professor

Massachusetts Institute of Technology



Tackling The Challenges of Big Data

Big Data Storage Distributed Computing Platforms MapReduce Examples

Matei Zaharia

Assistant Professor
Massachusetts Institute of Technology




1. Search

- **Input:** (lineNumber, line) records
- **Output:** lines matching a given pattern
- **Map:**

```
if (line matches pattern):  
    output(line, lineNumber)
```
- **Reduce:** identity function
 - Alternative: no reducer (map-only job)

2. Sort

- **Input:** (key, value) records
- **Output:** same records, sorted by key
- **Map:** identity function
- **Reduce:** identity function
- **Pick partitioning function p so that $k_1 < k_2 \Rightarrow p(k_1) < p(k_2)$**




3. Inverted Index

- **Input:** (filename, text) records
- **Output:** list of files containing each word
- **Map:**

```
for word in text.split():  
    output(word, filename)
```
- **Reduce:**

```
def reduce(word, filenames):  
    output(word, unique(filenames))
```

Inverted Index Example



4. Most Popular Words

- **Input:** (filename, text) records
- **Output:** the 100 words occurring in most files
- **Two-stage solution:**
 - Job 1:
 - Create inverted index, giving (word, list(file)) records
 - Job 2:
 - Map each (word, list(file)) to (count, word)
 - Sort these records by count as in sort job

Tackling The Challenges of Big Data

Big Data Storage

Distributed Computing Platforms

MapReduce Examples

THANK YOU



PROFESSIONAL
EDUCATION



Tackling The Challenges of Big Data

Big Data Storage

Matei Zaharia

Assistant Professor

Massachusetts Institute of Technology



PROFESSIONAL
EDUCATION



Tackling The Challenges of Big Data

Big Data Storage

Distributed Computing Platforms

Models Built on MapReduce

Matei Zaharia

Assistant Professor

Massachusetts Institute of Technology



PROFESSIONAL
EDUCATION



Limitations of MapReduce

- **MapReduce is great at single-pass analysis, but most applications require multiple MR steps**
- **Examples:**
 - Google indexing pipeline: 21 steps
 - Analytics queries (e.g. sessions, top K): 2-5 steps
 - Iterative algorithms (e.g. PageRank): 10's of steps
- **Two problems: programmability & performance**



Tackling the Challenges of Big Data

© 2014 Massachusetts Institute of Technology



Programmability

- **Multi-step jobs lead to convoluted code**
 - 21 MR steps -> 21 mapper and reducer functions
- **Repeated code for common patterns**



Tackling the Challenges of Big Data

© 2014 Massachusetts Institute of Technology



Performance

- **MR only provides one pass of computation**
 - Must write out data to file system in-between jobs
- **Expensive for apps that need to *reuse* data**
 - Multi-pass algorithms (e.g. PageRank)
 - Interactive data mining (many queries on same data)



Tackling the Challenges of Big Data

© 2014 Massachusetts Institute of Technology



Reactions

- **Higher-level interfaces over MapReduce**
 - Translate a higher-level language into MR steps
 - May merge steps to optimize performance
 - Examples: Hive, Pig
- **Generalizations of the model**
 - Examples: Dryad, Spark

Hive


- **Relational data warehouse over Hadoop**
 - Maintains a catalog of tables with schemas
 - Runs queries in a subset of SQL
 - Supports traditional query optimization, as well as complex data types and MapReduce scripts
- **Developed at Facebook**
- **Used for most of Facebook's MR jobs**



Hive Example

- **Given a tables of page views and user info, find top 5 pages visited by users under 20**

```
SELECT url, COUNT(*) AS views
FROM users u, pageviews v
WHERE u.name = v.username
AND u.age < 20
ORDER BY views DESC
LIMIT 5
```



Pig Example

```
Users = load 'users.txt' as (name, age);
Filtered = filter Users by age < 20;
Pages = load 'pageviews.txt' as (user, url);
Joined = join Filtered by name, Pages by user;
Grouped = group Joined by url;
Summed = foreach Grouped generate group,
count(Joined) as clicks;
Sorted = order Summed by clicks desc;
Top5 = limit Sorted 5;

store Top5 into 'top5sites';
```

- Scripting-like language offering SQL operators
- Developed and widely used at Yahoo!

Example from <https://cwiki.apache.org/confluence/download/attachments/26805800/HadoopSummit2009.ppt>
Tackling the Challenges of Big Data © 2014 Massachusetts Institute of Technology

Tackling The Challenges of Big Data

Big Data Storage

Distributed Computing Platforms

Models Built on MapReduce

THANK YOU



Tackling The Challenges of Big Data

Big Data Storage

Matei Zaharia

Assistant Professor

Massachusetts Institute of Technology



Tackling The Challenges of Big Data

Big Data Storage Distributed Computing Platforms Generalizations of MapReduce


Matei Zaharia

Assistant Professor
Massachusetts Institute of Technology



Motivation

- **Compiling high-level operators to MapReduce helps, but multi-step apps still have limitations**
 - Sharing data *between* MapReduce jobs is slow (requires writing to replicated file system)
 - Some apps need to do multiple passes over *same* data
- **E.g., word count + sort:**




Tackling the Challenges of Big Data

© 2014 Massachusetts Institute of Technology



Dryad Model

- **General graphs of tasks instead of two-level MapReduce graph**



- **Similar recovery mechanisms (replay parent tasks in graph to recover, replicate slow tasks)**

[Isard et al, EuroSys 2007]




Tackling the Challenges of Big Data

© 2014 Massachusetts Institute of Technology



Spark Model

- Dryad supports richer operator graphs, but data flow is still acyclic
- Many applications need to efficiently *reuse* data
 - Interactive data mining: run multiple user queries (not known in advance) on same subset of data
 - Iterative algorithms: make multiple passes over data



Tackling the Challenges of Big Data © 2014 Massachusetts Institute of Technology

Spark Model

- Let users explicitly build and persist distributed datasets
- Key idea: Resilient Distributed Datasets (RDDs)
 - Collections of objects partitioned across cluster that can be stored on disk or in memory
 - Built through graphs of parallel transformations (e.g. *map*, *reduce*, *group-by*)
 - Automatically rebuilt on failure
- High-level APIs in Java, Scala, Python

Tackling the Challenges of Big Data © 2014 Massachusetts Institute of Technology


Example: Log Mining

Load error messages from a log file in memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")
errors = lines.filter(lambda x: x.startswith("ERROR"))
messages = errors.map(lambda x: x.split('\t')[2])
messages.persist()
```

```
messages.filter(lambda x: "foo" in x).count()
messages.filter(lambda x: "bar" in x).count()
...
```

Result: full-text search of Wikipedia in 1 sec (vs 30 sec for on-disk data)




Tackling the Challenges of Big Data © 2014 Massachusetts Institute of Technology

Fault Recovery

RDDs track their *lineage graph* to rebuild lost data


```
file.map(lambda rec: (rec.type, 1))  
    .reduceByKey(lambda x, y: x + y)  
    .filter(lambda (type, count): count > 10)
```



Fault Recovery

RDDs track their *lineage graph* to rebuild lost data

```
file.map(lambda rec: (rec.type, 1))  
    .reduceByKey(lambda x, y: x + y)  
    .filter(lambda (type, count): count > 10)
```



Benefits of RDDs

- **Can express general computation graphs, similar to Dryad**
- **No need to replicate data for fault tolerance**
 - Write at memory speed instead of network speed
 - Use less memory than replicated systems
- **Persisting is only a hint**
 - If there's too little memory, can spill to disk or just drop old data

Iterative Algorithms



Tackling The Challenges of Big Data
Big Data Storage
Distributed Computing Platforms
Generalizations of MapReduce

THANK YOU

Tackling The Challenges of Big Data
Big Data Storage

Matei Zaharia

Assistant Professor
Massachusetts Institute of Technology

Tackling The Challenges of Big Data

Big Data Storage
Distributed Computing Platforms
Spark Demo

Matei Zaharia

Assistant Professor
Massachusetts Institute of Technology



PROFESSIONAL
EDUCATION



Tackling The Challenges of Big Data

Big Data Storage
Distributed Computing Platforms
Spark Demo

THANK YOU



PROFESSIONAL
EDUCATION



Tackling The Challenges of Big Data

Big Data Storage

Matei Zaharia

Assistant Professor
Massachusetts Institute of Technology



PROFESSIONAL
EDUCATION



Tackling The Challenges of Big Data

Big Data Storage
Distributed Computing Platforms
Other Platforms

Matei Zaharia

Assistant Professor
Massachusetts Institute of Technology



Introduction

- **We've covered MapReduce and its extensions, which are the most popular models today**
- **This area is still seeing rapid change & innovation**
- **We'll briefly sketch some related questions:**
 - How do these efforts relate to databases and SQL?
 - What are some very different models from MapReduce?
*Asynchronous computation, streaming



Tackling the Challenges of Big Data

© 2014 Massachusetts Institute of Technology



1. Convergence with SQL

- **One of the first things users wanted to run in MapReduce clusters was SQL!**
- **Some systems build SQL engines over MapReduce cluster architectures**
 - Google Dremel, Cloudera Impala, Apache Drill & Tez
 - Usually lack fault tolerance but target short queries
- **Others implement many of the optimizations in database engines on MapReduce-like runtimes**
 - Google Tenzing, Shark (Hive on Spark)



Tackling the Challenges of Big Data

© 2014 Massachusetts Institute of Technology



How This Works

- **Several things make analytical databases fast**
 - Efficient storage format (e.g. column-oriented)
 - Precomputation (e.g. indices, partitioning, statistics)
 - Query optimization
- **Shark (Hive on Spark)**
 - Implements column-oriented storage and processing *within* Spark records
 - Supports data partitioning and statistics on load
- **Result: can also combine SQL with Spark code**
 - From Spark: `rdd = shark.sqlRdd("select * from users")`
 - From SQL: `generate KMeans(select lat, long from tweets)`



Tackling the Challenges of Big Data

© 2014 Massachusetts Institute of Technology



In the Other Direction

- **Some databases adding support for MapReduce**
- **Greenplum, Aster Data: MapReduce executes within database engine, on relational data**
- **Hadapt: hybrid database / MapReduce system**
 - Queries sent to Hadoop cluster with DB on each node
 - Supports “schema on read” functionality (e.g. run SQL over JSON records)



Tackling the Challenges of Big Data

© 2014 Massachusetts Institute of Technology



2. Asynchronous Computing

- **MapReduce & related models use deterministic, synchronized computation for fault recovery**
- **Some applications (e.g. numerical optimization, machine learning) can converge without this**
 - Each step makes “progress” towards a goal
 - Losing state just moves you slightly away from goal
- **Asynchronous systems use this to improve speed**
 - Don’t wait for all nodes to advance / communicate
 - Examples: GraphLab, Hogwild, Google DistBelief



Tackling the Challenges of Big Data


© 2014 Massachusetts Institute of Technology



Asynchronous System Example


3. Stream Processing

- Processing data in real-time can require very different designs
- Example: Storm (continuous, stateful operators)
 - Build graph of operators
 - System plays each record through at least once in case of failures
- Several options for determinism on failures
 - Transactional storage (Trident)
 - Run as sequence of batch jobs (Spark Streaming)



Resource Sharing

- Given all these programming models, how do we get them to coexist
- Cross-framework *resource managers* allow dynamically sharing a cluster between them
 - Let applications launch work through common API
 - Hadoop YARN, Apache Mesos



Source: Hortonworks (<http://hortonworks.com/hadoop/yarn/>)

Conclusion

- Large-scale cluster environments are difficult to program with traditional methods
- Result is a profusion of new programming models that aim to simplify this
- Main idea: capture computation in a declarative manner to let system handle distribution



Tackling the Challenges of Big Data

© 2014 Massachusetts Institute of Technology



Tackling The Challenges of Big Data
Big Data Storage
Distributed Computing Platforms
Other Platforms

THANK YOU



PROFESSIONAL
EDUCATION



Tackling The Challenges of Big Data
Big Data Storage
Distributed Computing Platforms

Matei Zaharia

Assistant Professor

Massachusetts Institute of Technology



PROFESSIONAL
EDUCATION

