Tackling The Challenges of Big Data
Big Data Systems

Nickolai Zeldovich

Associate Professor

Massachusetts Institute of Technology

Tackling The Challenges of Big Data
Big Data Systems
Security
Introduction

Nickolai Zeldovich

Associate Professor

Massachusetts Institute of Technology

Security is a Negative Goal

- **No way for adversary to violate security policy**
- **Difficult to achieve: many avenues of attack**

Example: Confidential Database

Approach: Encryption

- Server has no access to decryption key
  - Cannot decrypt data, even if server is compromised
  - Broad threat model
- **Challenge:** how to process queries without decrypting?

Outline: Encrypted Query Processing

- **Theoretical results: fully homomorphic encryption**
- **Practical approach: CryptDB**
  - Specialized encryption schemes
  - Order-preserving encryption
  - Onions of encryption
- **Results from a prototype of CryptDB**

Tackling The Challenges of Big Data
Big Data Systems
Security
Introduction

THANK YOU

Tackling The Challenges of Big Data
Big Data Systems

Nickolai Zeldovich

Associate Professor

Massachusetts Institute of Technology

Tackling The Challenges of Big Data
Big Data Systems
Security
Fully homomorphic encryption

Nickolai Zeldovich

Associate Professor

Massachusetts Institute of Technology

Fully Homomorphic Encryption (FHE)

- **Recent breakthrough result: Craig Gentry, 2009**

FHE in a Database Context

name	salary
Alice	60
Bob	100
Carl	800
Doug	100

Tackling The Challenges of Big Data

Big Data Systems

Security

Fully homomorphic encryption

THANK YOU

PROFESSIONAL EDUCATION

CSAIL

Tackling The Challenges of Big Data
Big Data Systems

Nickolai Zeldovich

Associate Professor

Massachusetts Institute of Technology

Tackling The Challenges of Big Data
Big Data Systems
Security
CryptDB approach

Nickolai Zeldovich

Associate Professor

Massachusetts Institute of Technology

Goal: Comparable Performance to Plaintext Database

- **Server must use efficient index data structures**
- **Server must determine if row matches query**

CryptDB's Approach

- **Trade off some generality, security for performance**
  - Expose database structure to server: rows, columns
  - Encrypt individual cell values
  - Use encryption schemes that enable specific functions
- **Reveals some information to the server**
- **Necessary for performance**
  - Server must decide if a row matches a query

Example

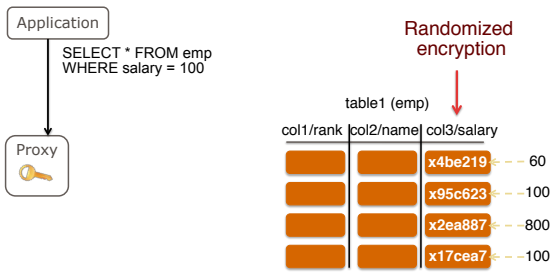

Example

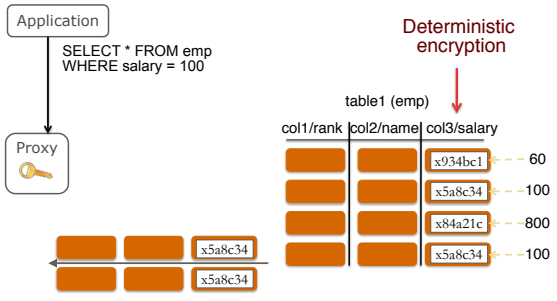

Example

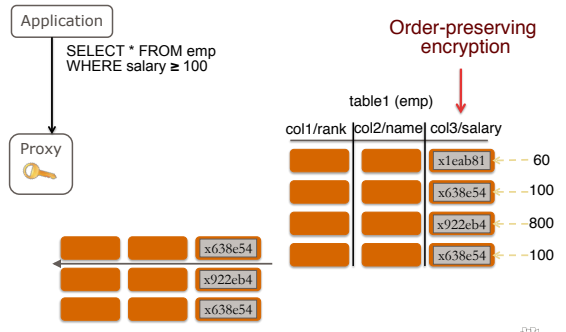

Tackling The Challenges of Big Data

Big Data Systems

Security

CryptDB approach

THANK YOU

PROFESSIONAL EDUCATION

CSAIL

Tackling The Challenges of Big Data

Big Data Systems

Nickolai Zeldovich

Associate Professor

Massachusetts Institute of Technology

PROFESSIONAL EDUCATION

CSAIL

Tackling The Challenges of Big Data

Big Data Systems

Security

Order-preserving encryption

Nickolai Zeldovich

Associate Professor

Massachusetts Institute of Technology

Order-Preserving Encryption Goal

- **Functionality: order-preserving**
- **Security: indistinguishability**

Sketch of OPE Construction (mOPE)

E(K, 20)

E(K, 18)

E(K, 100)

E(K, 3)

K=

Client | **Server**

Tackling The Challenges of Big Data
Big Data Systems
Security
Order-preserving encryption

THANK YOU

Tackling The Challenges of Big Data
Big Data Systems

Nickolai Zeldovich

Associate Professor

Massachusetts Institute of Technology

Tackling The Challenges of Big Data
Big Data Systems
Security
Multiple encryption schemes

Nickolai Zeldovich

Associate Professor

Massachusetts Institute of Technology

Multiple Encryption Schemes

- **Semantic security (AES-CBC)**
- **Homomorphic (Paillier, ElGamal)**
- **Searchable encryption**
- **Deterministic (AES-CMC)**
- **JOIN**
- **Order-preserving encryption**

How to Encrypt Data?

- **Encryption schemes depend on queries**
- **May not know queries ahead of time**

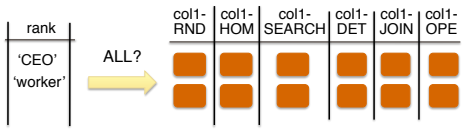

Onions of Encryption

Confidentiality Guarantees

- Never reveal plaintext data to the server
- Queries → Encryption schemes → Leakage
- Reveal most secure scheme that supports query
- Use thresholds to limit leakage (e.g., no OPE)

Tackling The Challenges of Big Data

Big Data Systems

Security

Multiple encryption schemes

THANK YOU

PROFESSIONAL EDUCATION

CSAIL

Tackling The Challenges of Big Data

Big Data Systems

Nickolai Zeldovich

Associate Professor

Massachusetts Institute of Technology

PROFESSIONAL EDUCATION

CSAIL

Tackling The Challenges of Big Data

Big Data Systems

Security

CryptDB results

Nickolai Zeldovich

Associate Professor

Massachusetts Institute of Technology

Implementation

- **Proxy rewrites query to run over encrypted data**
  - User-defined functions for crypto (onion adjustment, …)
- **Prototype: 26,000 lines of C++ code**

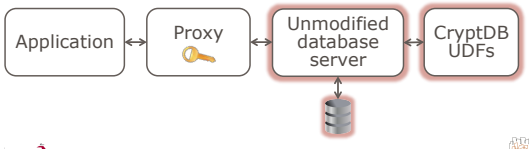

Evaluation Questions

- **Can CryptDB support real queries and applications?**
- **What is the resulting level of confidentiality?**
- **What is the performance overhead of CryptDB?**

CryptDB supports real applications

Application	Total columns	Encrypted columns	# cols not supported
phpBB	563	23	0
HotCRP	204	22	0
grad-apply	706	103	0
TPC-C	92	92	0
sql.mit.edu	128,840	128,840	1,094

SELECT 1/log(series_no+1.2) …
… WHERE sin(latitude + PI()) …

Onions provide high confidentiality

Application	Total columns	Encrypted columns	Min level is RND	Min level is DET	Min level is OPE
phpBB	563	23	21	1	1
HotCRP	204	22	18	1	2
grad-apply	706	103	95	6	2
TPC-C	92	92	65	19	8
sql.mit.edu	128,840	128,840	80,053	34,212	13,131

Performance overheads are modest

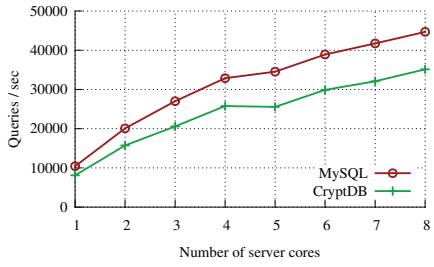

- **TPC-C database benchmark: small transactions**

Summary of CryptDB Evaluation

- **CryptDB supports most SQL queries in practice**
- **CryptDB provides high confidentiality for most data**
- **CryptDB's performance overheads are modest**

Tackling The Challenges of Big Data

Big Data Systems

Security

CryptDB results

THANK YOU

PROFESSIONAL EDUCATION

CSAIL

Tackling The Challenges of Big Data

Big Data Systems

Nickolai Zeldovich

Associate Professor

Massachusetts Institute of Technology

PROFESSIONAL EDUCATION

CSAIL

Tackling The Challenges of Big Data

Big Data Systems

Security

Conclusion

Nickolai Zeldovich

Associate Professor

Massachusetts Institute of Technology

Conclusion

- **Security is hard to achieve: negative goal**
  - Encryption can address a broad threat model
  - Challenge: computing on encrypted data
- **CryptDB: computing on encrypted data is practical**
  - Specialized encryption schemes
  - Onions of encryption
- **Beyond CryptDB**
  - Partitioning to handle complex computations
  - Push all encryption/decryption into the web browser

Tackling The Challenges of Big Data

Big Data Systems

Security

Conclusion

THANK YOU

Tackling The Challenges of Big Data

Module: Big Data Systems

Topic: Security

THANK YOU

Nickolai Zeldovich

Associate Professor

Massachusetts Institute of Technology