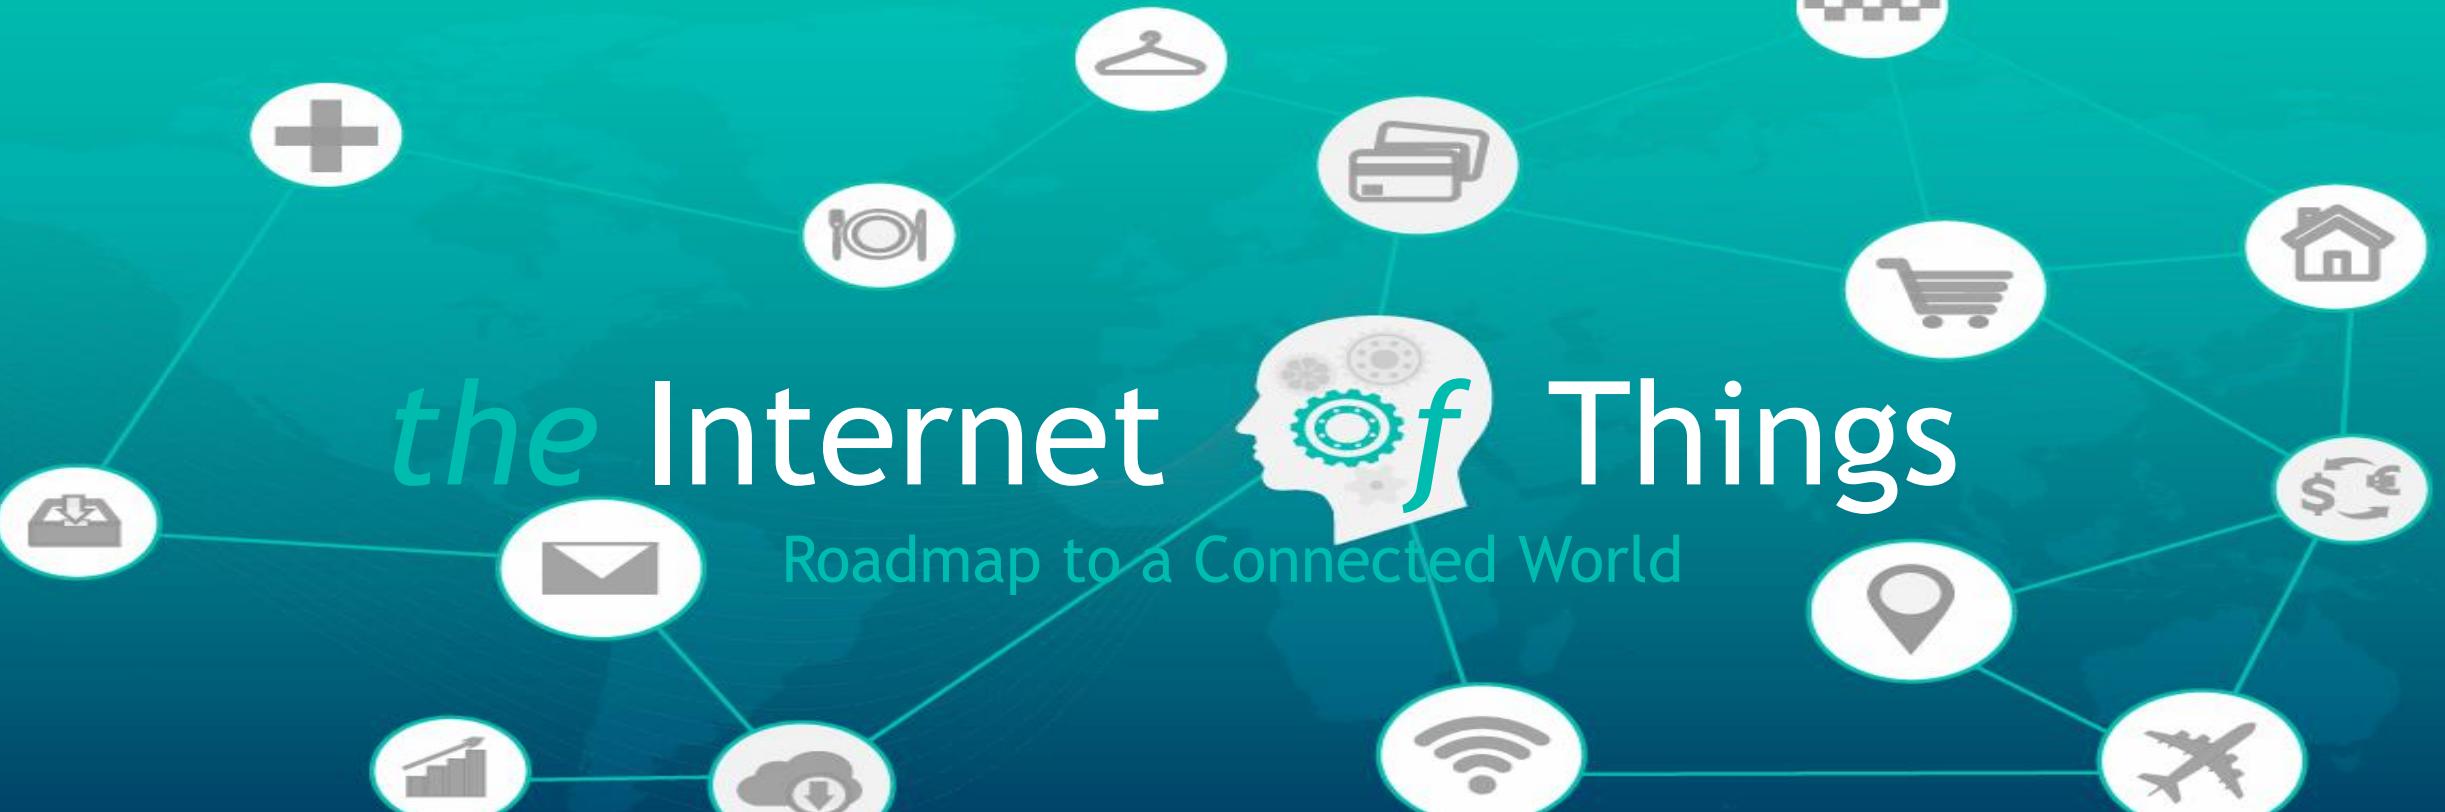




the Internet *of* Things

Roadmap to a Connected World



MODULE

Technologies

The Internet of Things: Roadmap to a Connected World

NETWORK CONNECTIVITY FOR IOT

HARI BALAKRISHNAN

Fujitsu Professor, EECS Department

Computer Science and Artificial Intelligence Laboratory (CSAIL)
Massachusetts Institute of Technology



The Internet of Things: Roadmap to a Connected World

© 2016 Massachusetts Institute of Technology



NETWORKING: “GLUE” FOR THE IOT

IoT’s “technology push” from the convergence of

- Embedded computing
- Sensing & actuation
- Wireless networks



The Internet of Things: Roadmap to a Connected World

© 2016 Massachusetts Institute of Technology



THE IOT CONNECTIVITY SOUP



Brand names, logos and trademarks used herein remain the property of their respective owners.

The Internet of Things: Roadmap to a Connected World

© 2016 Massachusetts Institute of Technology



NETWORKING: “GLUE” FOR THE IOT

(Too) many different approaches, (too) many different standards,
(too) much confusion

One size does not fit all: best network depends on application

What are the key organizing principles and ideas?

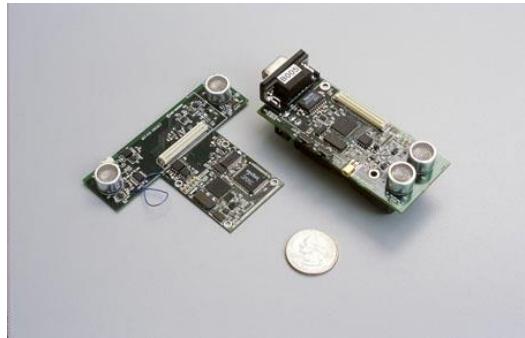


The Internet of Things: Roadmap to a Connected World

© 2016 Massachusetts Institute of Technology



MY IOT EXPERIENCE: CRICKET

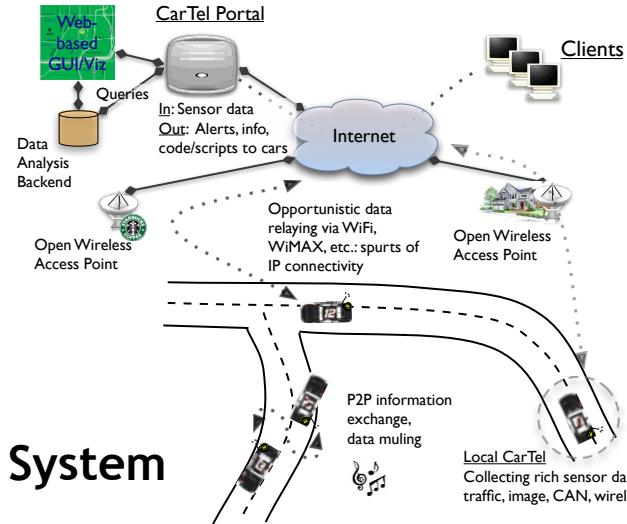


Cricket Indoor Location System

Started 1999-00, ran until 2005
Key component of MIT's Project Oxygen
3-5 cm accuracy for indoor location
Beacons in building transmit RF and ultrasound
Listeners on mobile devices and sensors estimate distance and trilateral position
Beacons self-configure using multihop protocol



MY IOT EXPERIENCE: CARTEL



CarTel Mobile Sensing System

Started 2005, ran until 2011

One of the first mobile sensing systems

Transportation focus: precursor to today's

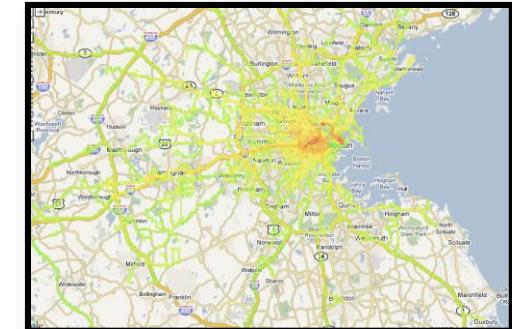
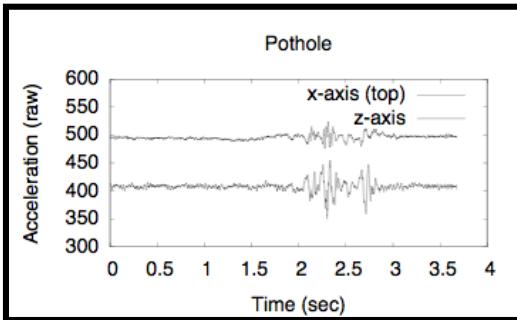
"connected car" and smartphone sensing efforts

Cabernet network: QuickWiFi for vehicular Wi-Fi

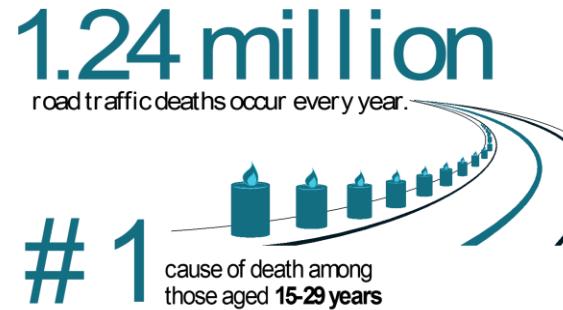
New map-matching algorithms

Traffic-aware routing

Pothole patrol



MY IOT EXPERIENCE: DRIVEWELL



DriveWell Safe Driving System

Started 2010-11

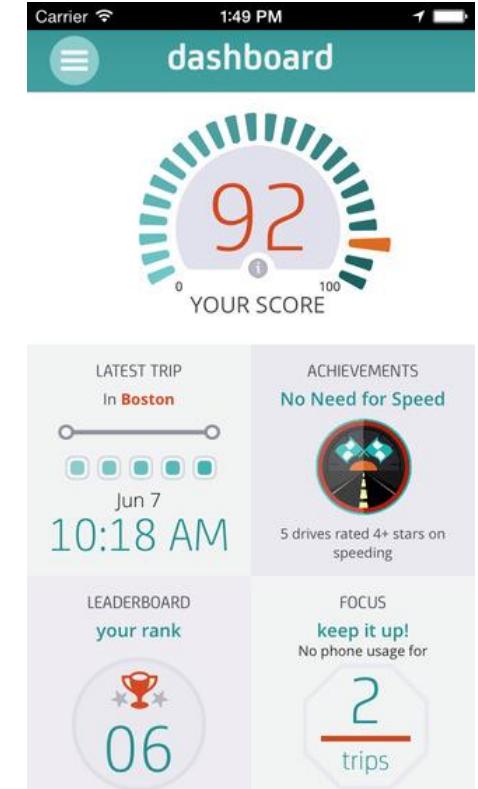
Commercially available (Cambridge Mobile Telematics)

Goal: Make roads safer by making drivers better

Smartphone sensing to assess driving safety

Feedback and incentives (behavioral economics)

App augmented with BLE DriveWell Tag
for real-time collision detection & roadside assistance



HOW TO MAKE SENSE OF THIS SOUP?



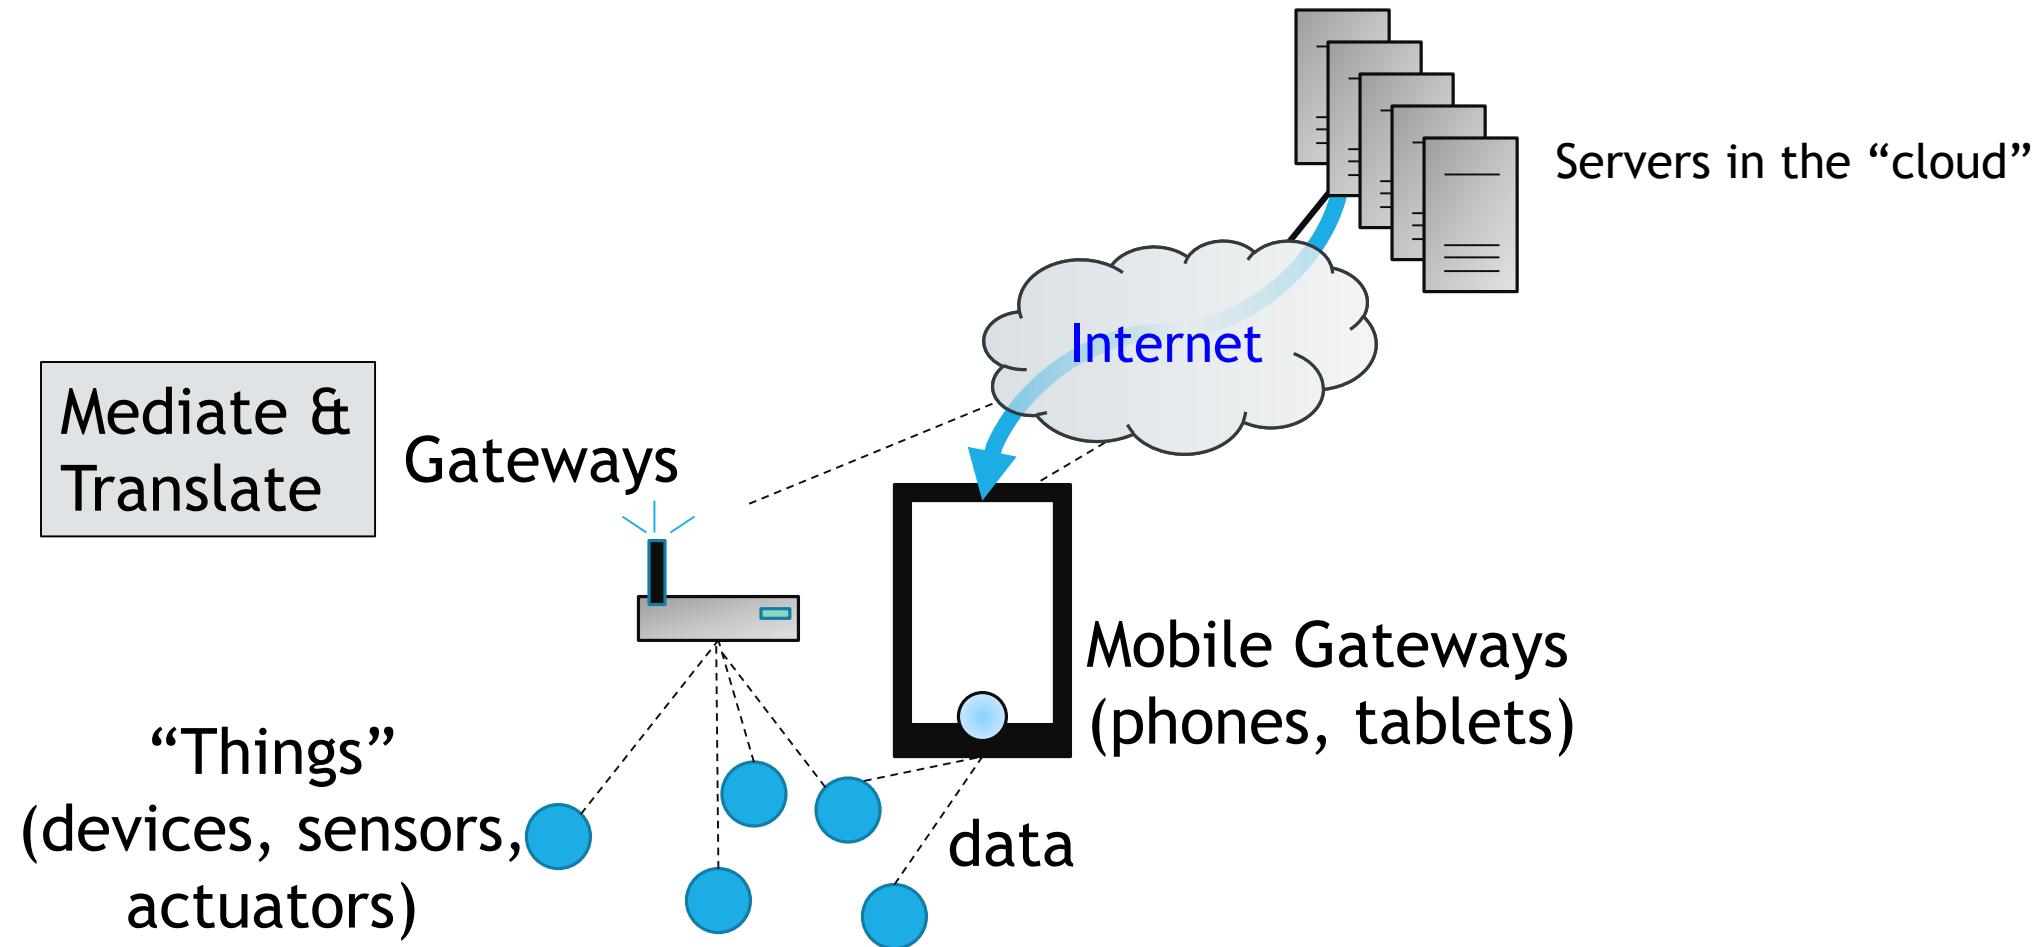
Brand names, logos and trademarks used herein remain the property of their respective owners.

The Internet of Things: Roadmap to a Connected World

© 2016 Massachusetts Institute of Technology



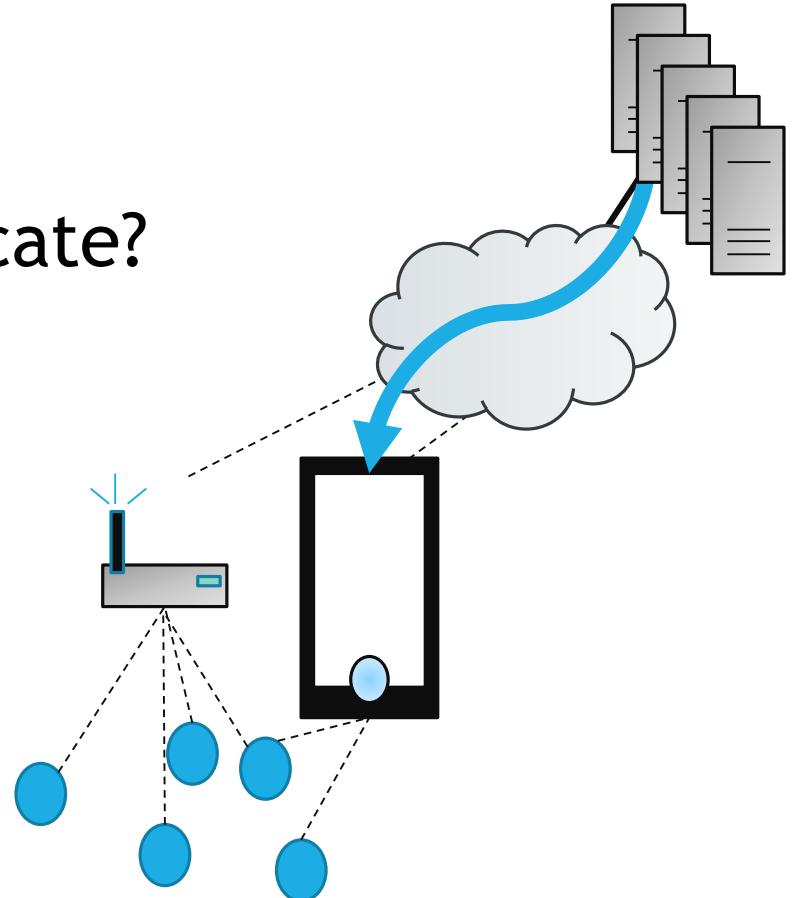
ARCHITECTURE, SIMPLIFIED



BUT, IN FACT, A RICH DESIGN SPACE

How should gateways and things communicate?

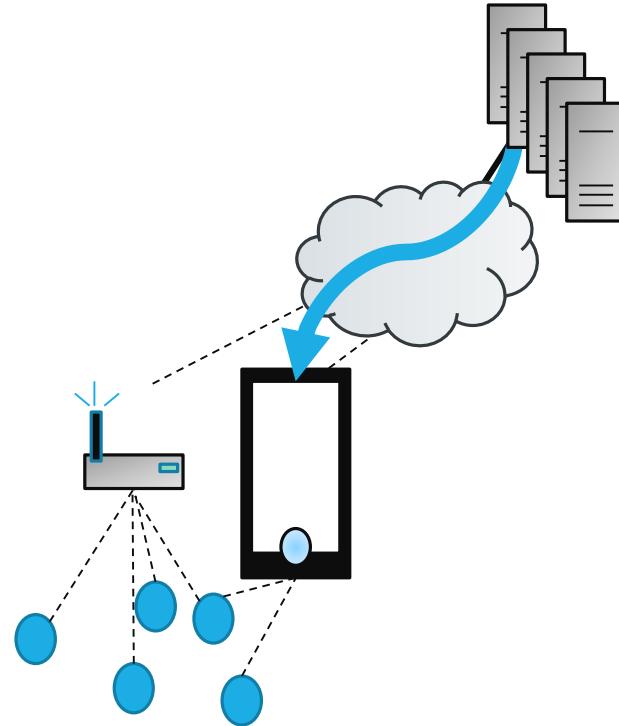
Many answers, many approaches



CAN'T WE JUST USE THE WIRELESS INTERNET?

Cellular and Wi-Fi

Yes, we can...
except when we can't!



WIRELESS INTERNET FOR IOT?

Cellular (LTE/4G, 3G, 2G) and Wi-Fi are

- + Widely available
- + High bandwidth (for most purposes), so can support high-rate apps

But, each has two big drawbacks

- **High power:** not suitable for battery-operated scenarios
- Cellular: often high cost (esp. per byte if usage-per-thing is low)
- Wi-Fi: OK in most buildings, but not for longer range

Wi-Fi: In-building powered things (speakers, washers, refrigerators, ...)

Cellular: High-valued powered things (e.g., “connected car”)

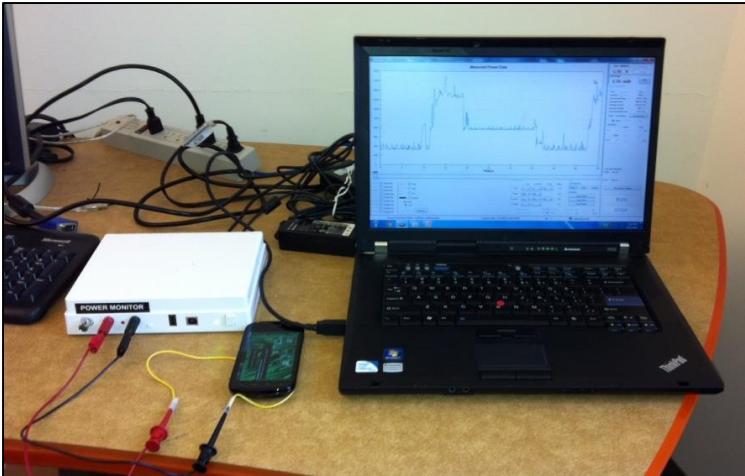


The Internet of Things: Roadmap to a Connected World

© 2016 Massachusetts Institute of Technology

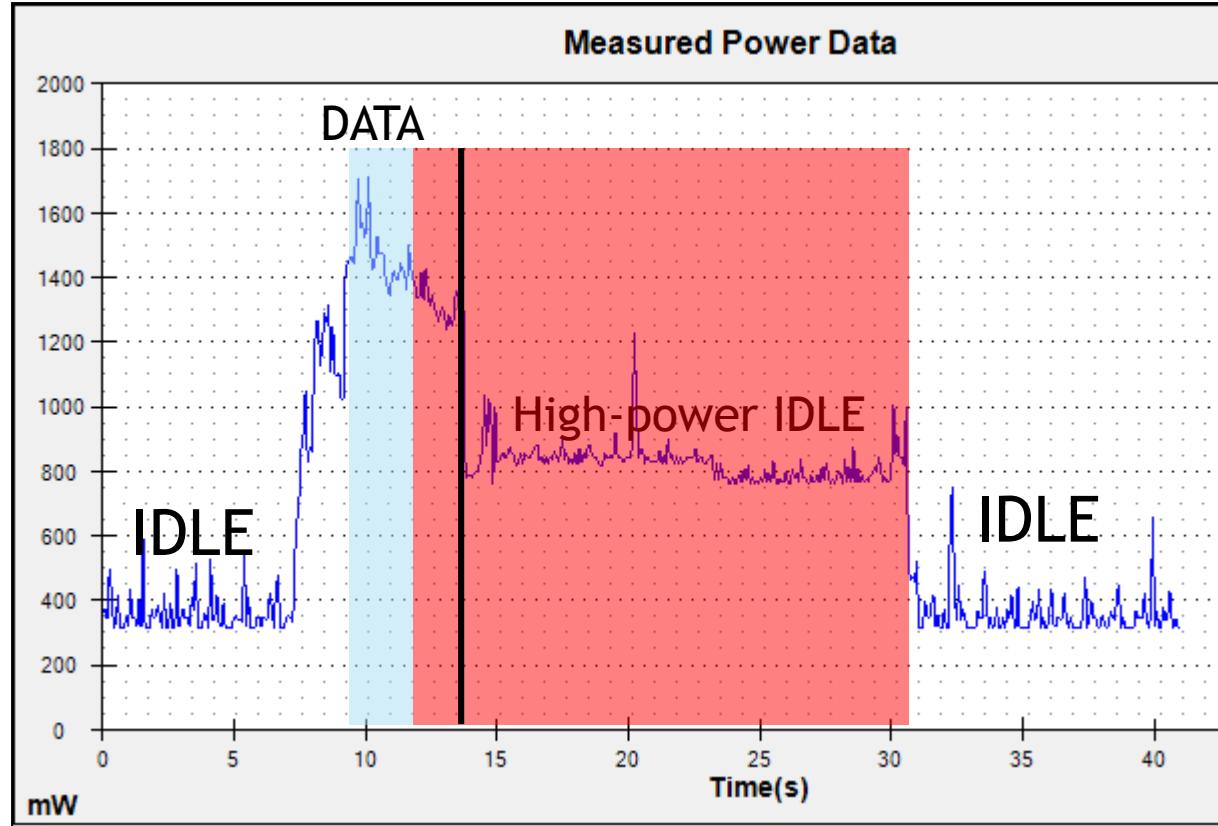


CELLULAR POWER CONSUMPTION

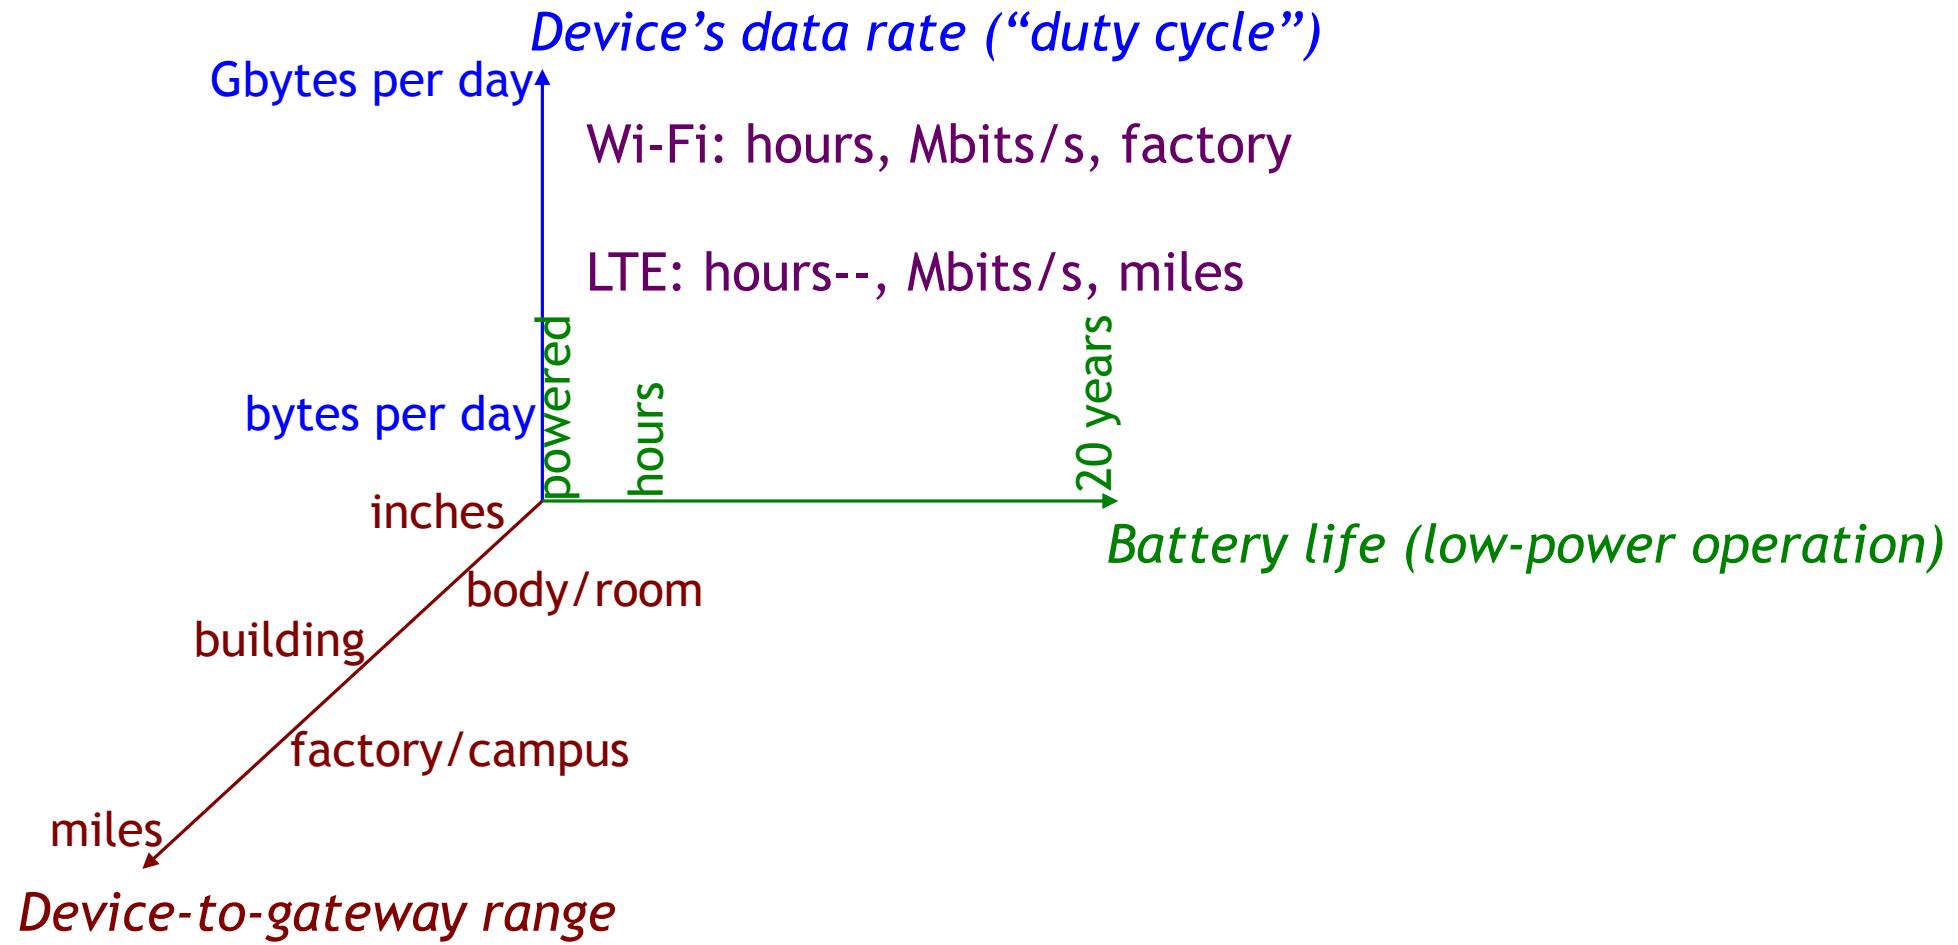


Power monitor apparatus

Deng & Balakrishnan, “Traffic-Aware Techniques to Reduce 3G/LTE Energy Consumption,” CoNext 2012.



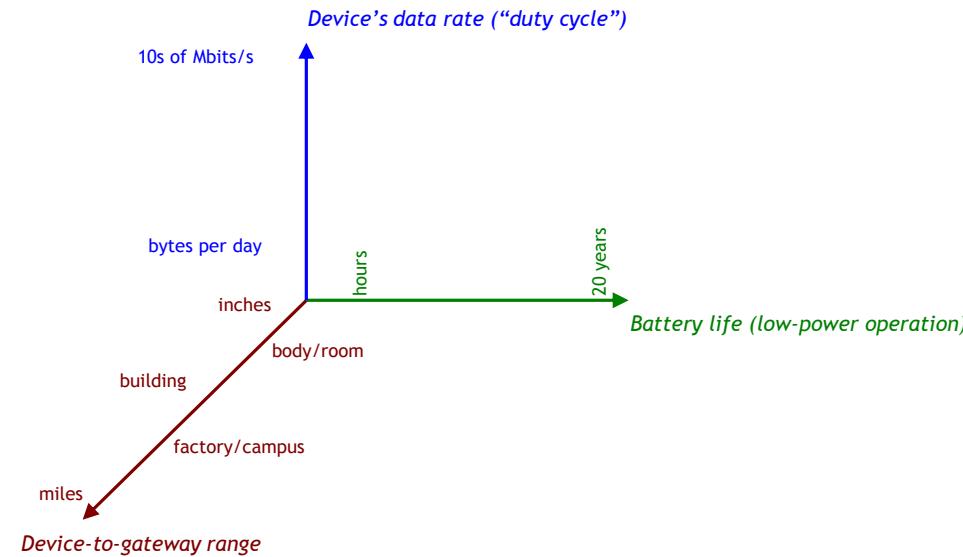
IOT NETWORK DESIGN SPACE



WHY SO MANY IOT NETWORKS?

Because engineers love inventing technologies!

But really because you can pick many interesting regions from this design space

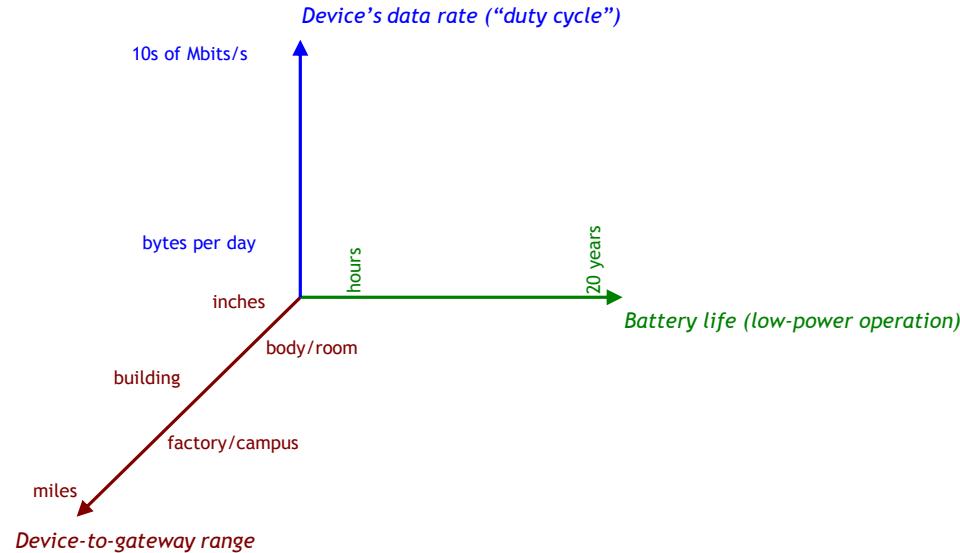


WHY SO MANY IOT NETWORKS?

Because engineers love inventing technologies!

But really because you can pick many interesting regions from this design space

- Note, axes aren't independent
- And technology evolves fast

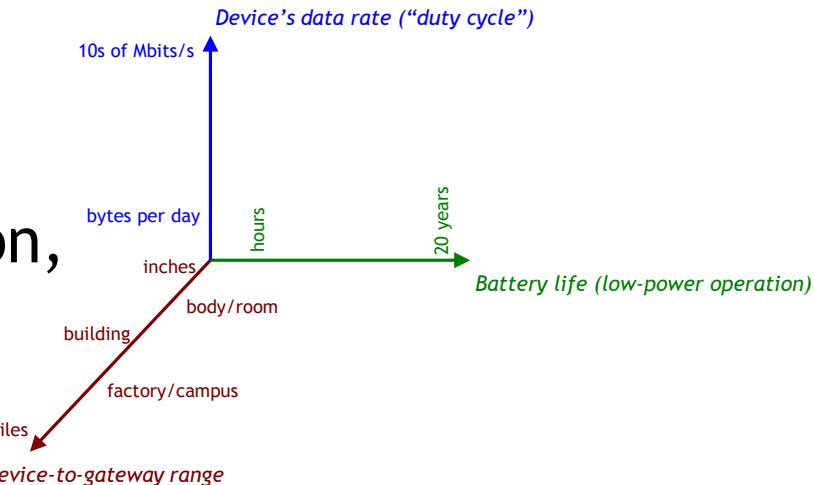


WHY SO MANY IOT NETWORKS?

Because engineers love inventing technologies!

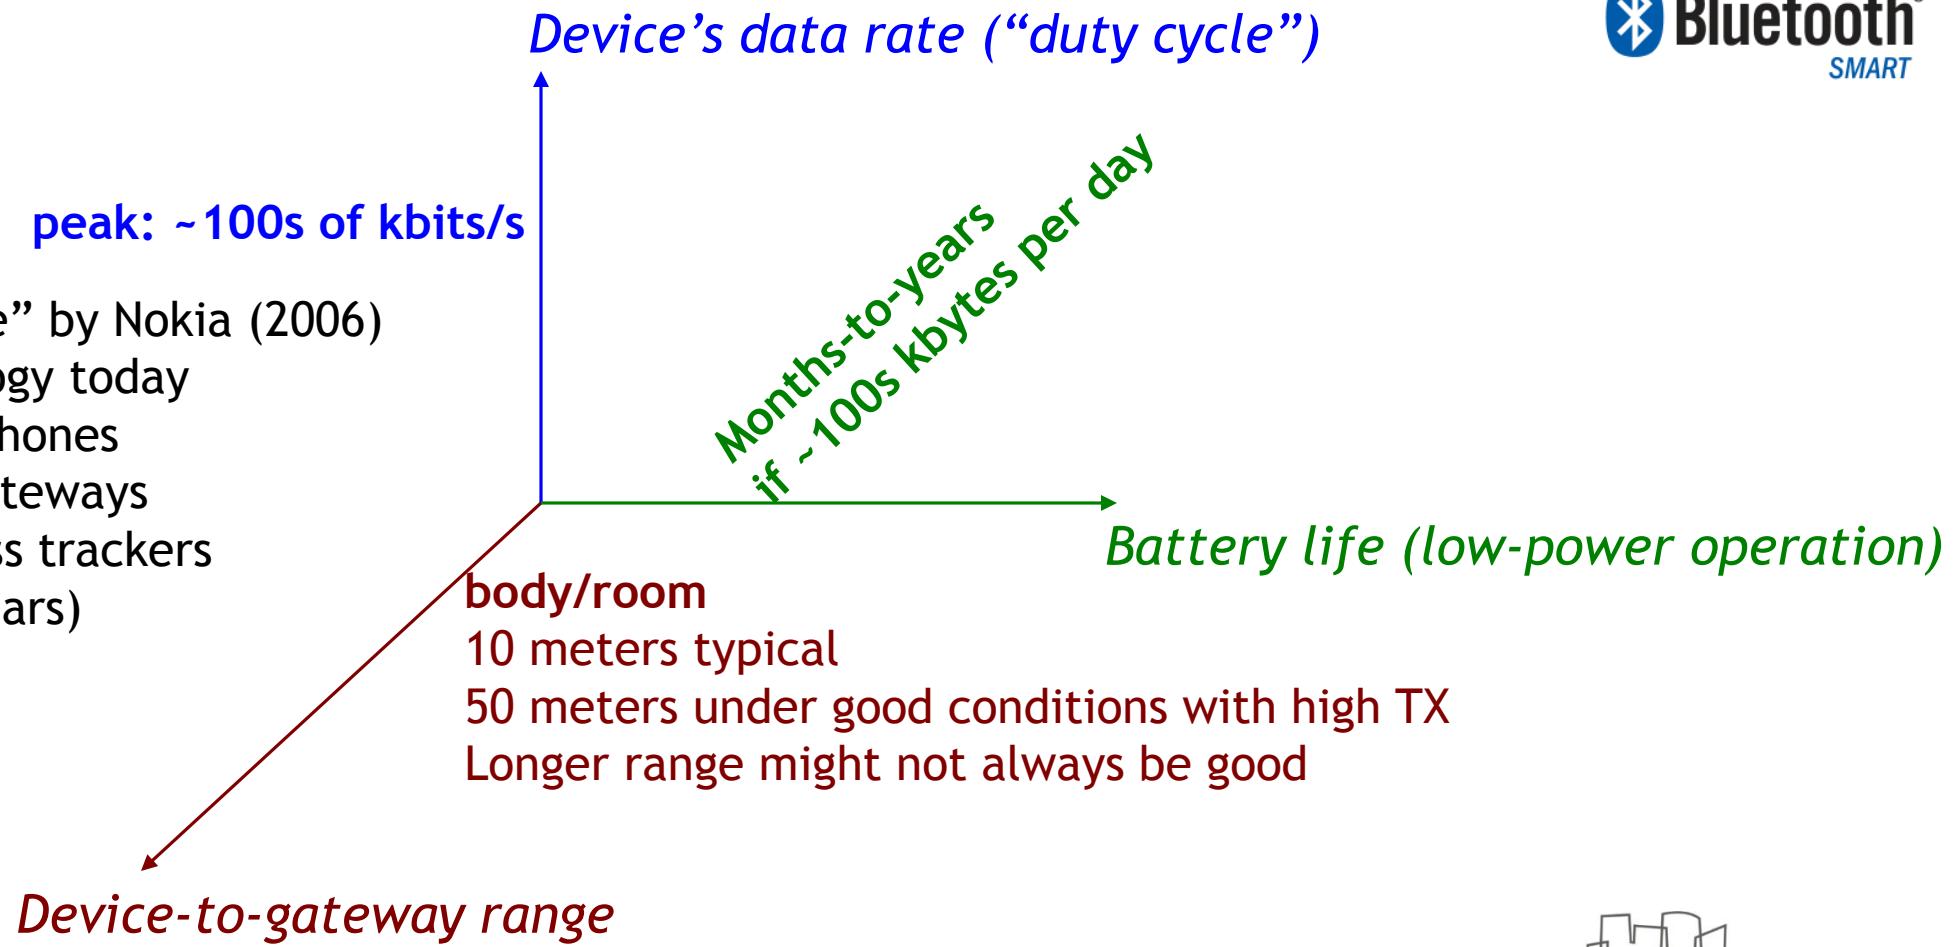
But really because you can pick many interesting regions from this design space

- Note, axes aren't independent
- And technology evolves fast
- And bundling into popular devices speeds-up adoption, changing the economics
 - Cf. Wi-Fi → laptops (without external cards)
 - Bluetooth classic → cell phones → wireless headsets
 - Bluetooth Low Energy (BLE) → iPhone then Android smartphones → “body/room” with months-to-years at low duty cycles



BODY/ROOM-AREA EXAMPLE: BLE

Started as “Wibree” by Nokia (2006)
Dominant technology today
Because of smartphones
Smartphones as gateways
Wearables, fitness trackers
Vehicles (bikes, cars)



HOW DOES BLE WORK?

Two parts:

1. Advertisements (aka “beaconing”) for device discovery
2. Connection phase to exchange data

Peripheral: device with data
Central: gateway



BLE ADVERTISEMENTS ARE PERIODIC



Typical period: 100 ms (“iBeacon”)
Less frequent is fine
Triggered advertisements are often a good idea

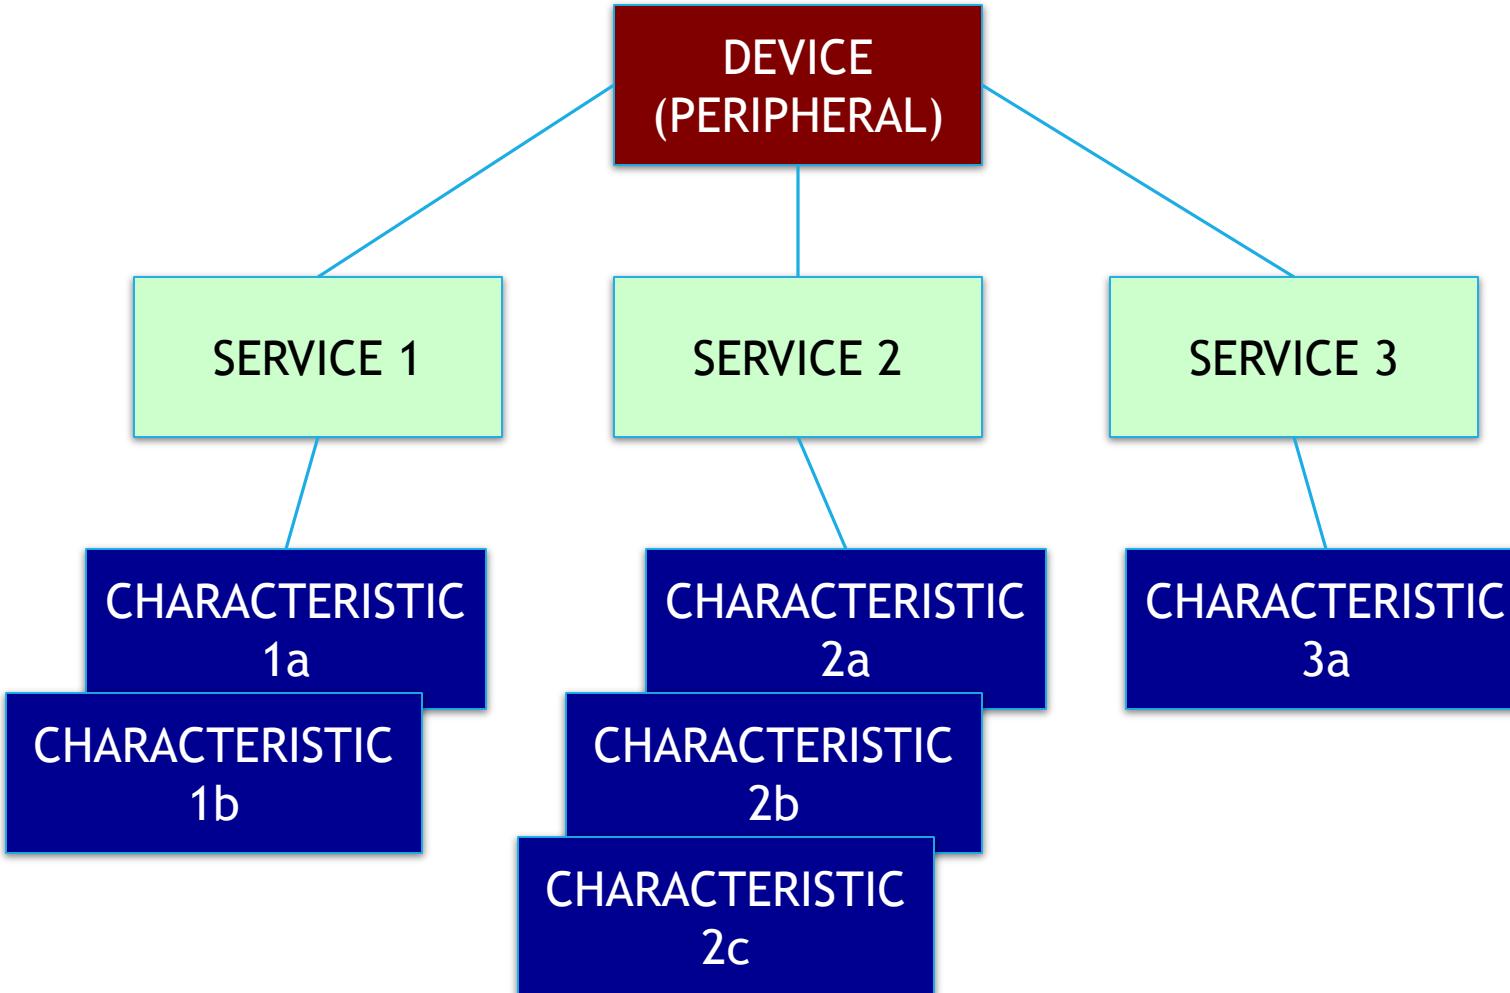


Trade-off between energy consumed
and discovery latency

ON CONNECTION

READABLE
READ/WRITE
NOTIFICATIONS

Usually support
OTA (over-the-air
upgrades)



The Internet of Things: Roadmap to a Connected World

© 2016 Massachusetts Institute of Technology



ON CONNECTION: MAC PROTOCOL

Central orchestrates data communication

Key idea: time-schedule to reduce energy consumption

On connect: exchange parameters

- Frequency hopping sequence
- Connection interval, i.e., periodicity of data exchange (T milliseconds)

Every T milliseconds, Central and Peripheral exchange up to 4 packets,
alternating turns

Then Peripheral can go back to sleep until next interval



The Internet of Things: Roadmap to a Connected World

© 2016 Massachusetts Institute of Technology



BATTERY LIFETIME CALCULATION

Consider an IoT system with coin-cell battery-powered nodes

Battery: 250 mAh (milliamp-hours) capacity; 3 Volts

Recall that power = voltage * current and energy = power * time

So this battery has $0.75 \text{ amp-hour-volts} = 0.75 * 3600 \text{ Joules} = 2.7 \text{ kJ}$ of energy

Example of BLE current draw:

Standby: 1 microAmp (typically in the 1-10 microAmp range)

Receive (RX): 3.3 mA

Transmit (TX): 4 mA

Suppose device transmits every second: how long does the battery last?



BATTERY CALCULATION (CONT.)

Consider an IoT system with coin-cell battery-powered nodes

Battery: 250 mAh (milliamp-hours) capacity; 3 Volts

Recall that power = voltage * current and energy = power * time

So this battery has 0.75 amp-hour-volts = 0.75×3600

Joules = 2.7 kJ of energy

Example of BLE current draw:

Standby: 1 microAmp (typically in the 1-10 microAmp range)

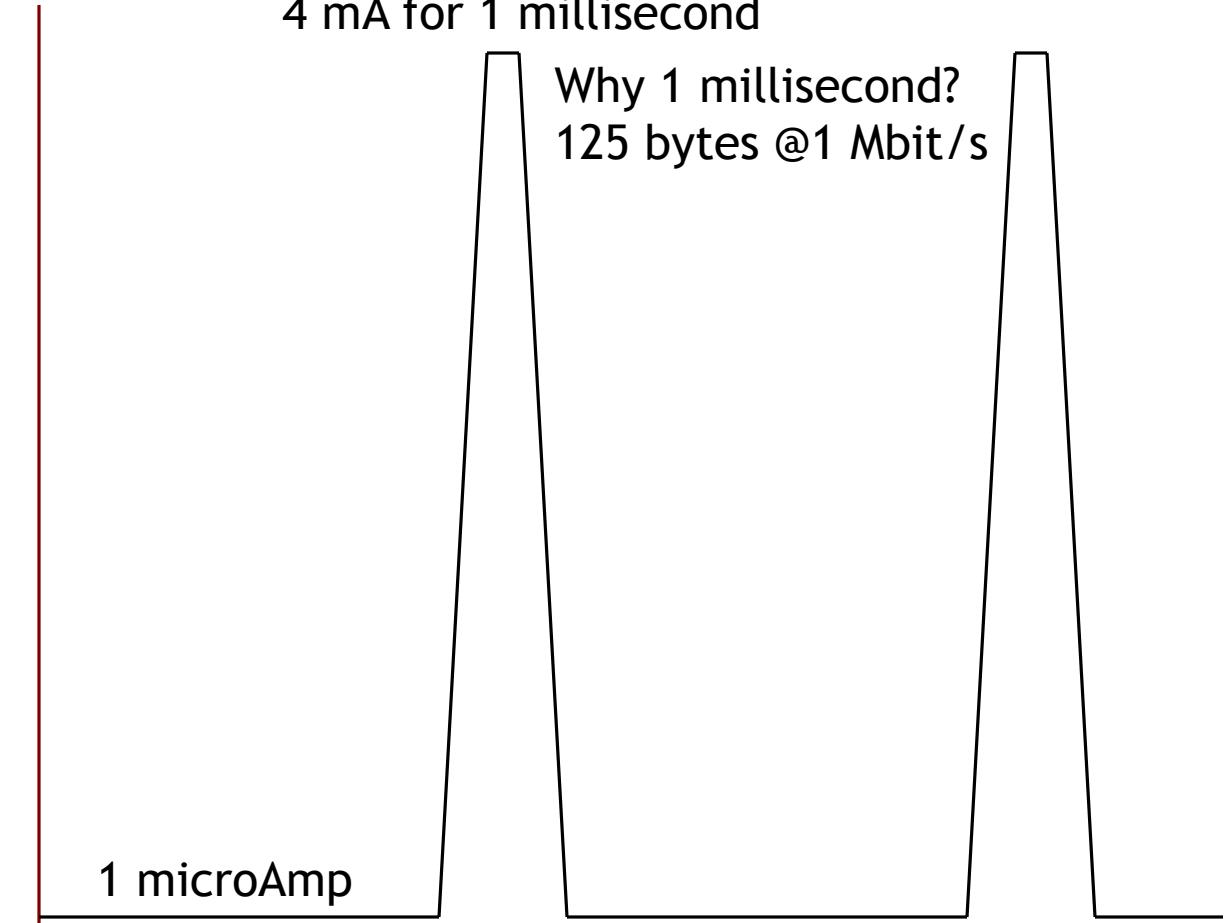
Receive (RX): 3.3 mA

Transmit (TX): 4 mA

Suppose device transmits every second: how long does the battery last?

4 mA for 1 millisecond

Why 1 millisecond?
125 bytes @1 Mbit/s



BATTERY CALCULATION (CONT.)

4 milliAmp for 1 millisecond

Why 1 millisecond?

125 bytes @1 Mbit/s

1 microAmp

Ramp-up and down: 1 milliAmp for 5 milliseconds

Over a 1 second interval, average current is:

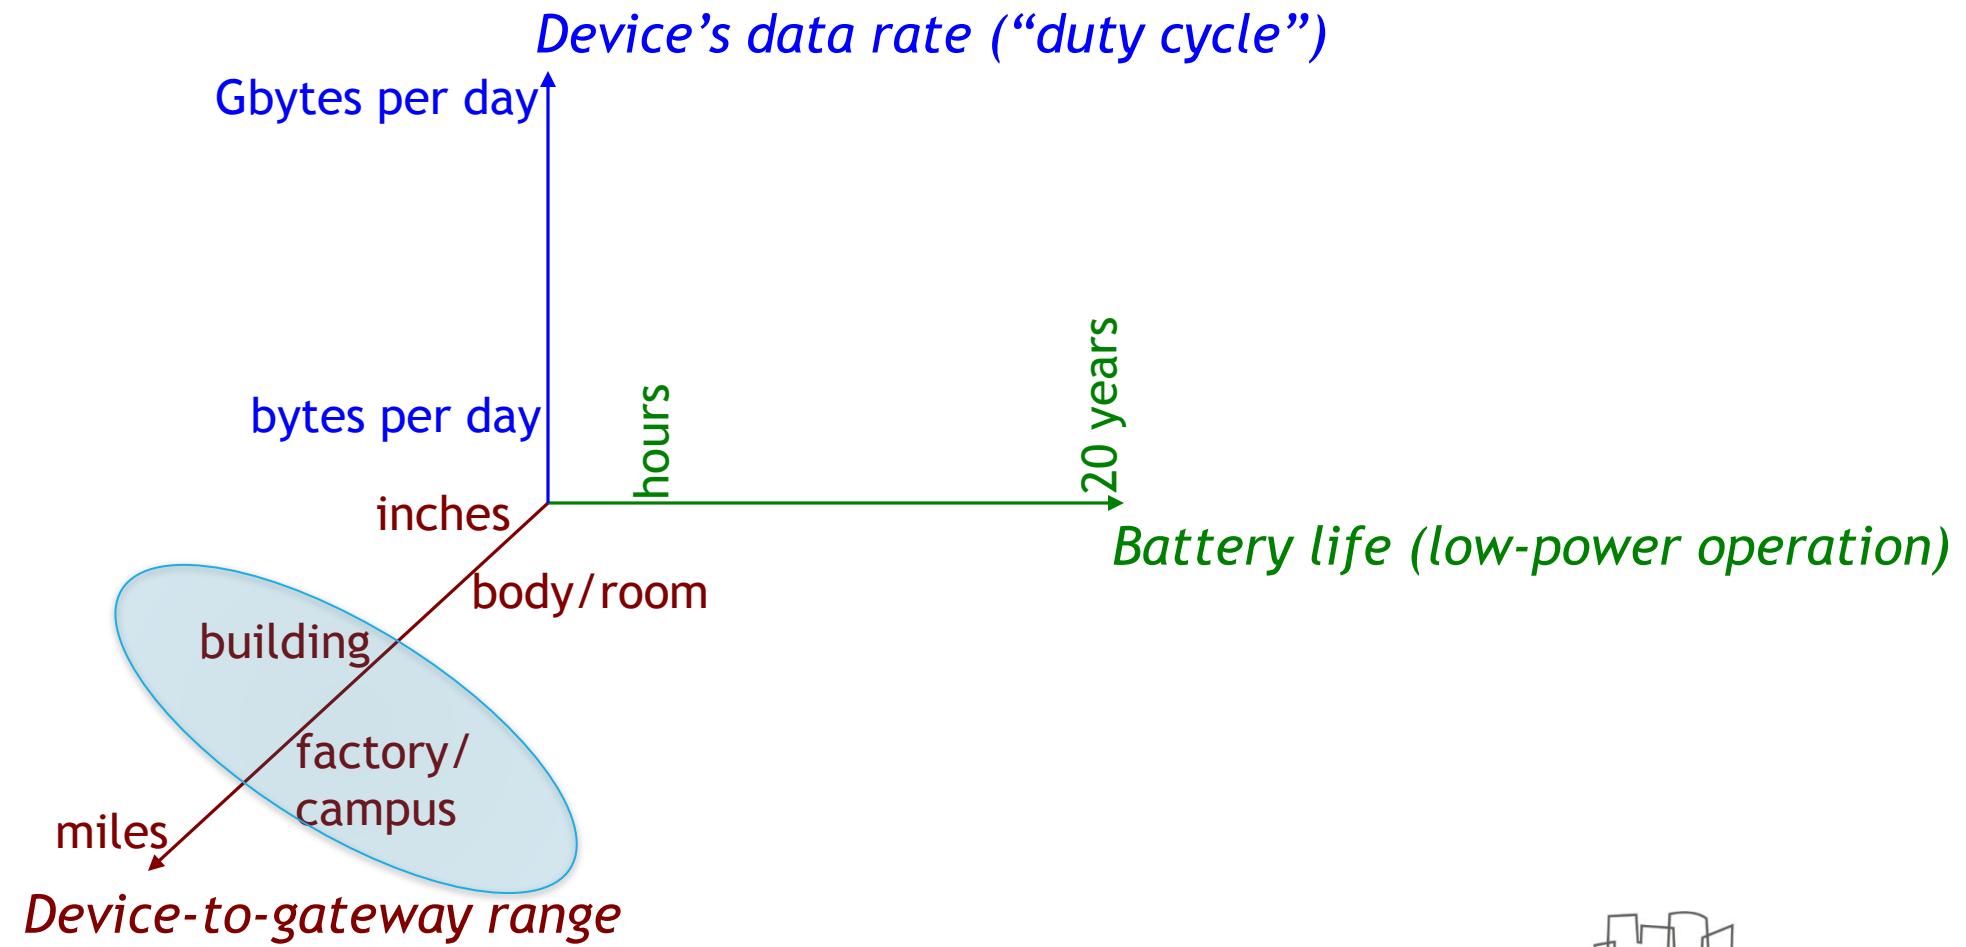
$$\begin{aligned} & 4 \text{ microAmps (xmit)} + \\ & 5 \text{ microAmps (ramping)} + \\ & 1 \text{ microAmp (standby)} \\ & = 10 \text{ microAmps} \end{aligned}$$

Therefore, battery lifetime

$$\begin{aligned} & = 250 \text{ mAh} / 10 \text{ microAmps} \\ & = 250 \text{ mAh} / 0.01 \text{ mA} \\ & = 25,000 \text{ hours} \\ & = 2 \text{ years and 10 months} \end{aligned}$$

This works because it's sleeping most of the time!

EXTENDING COMMUNICATION RANGE



EXTENDING RANGE: MESH NETWORKS

1980s: DARPA packet radio networks

The DARPA Packet Radio Network Protocols

JOHN JUBIN AND JANET D. TORNOW, ASSOCIATE, IEEE

Invited Paper

In this paper we describe the current state of the DARPA packet radio network. Fully automated algorithms and protocols to organize, control, maintain, and move traffic through the packet radio network have been designed, implemented, and tested. By means of protocols, networks of up to 50 packet radio nodes can be organized and managed under a fully distributed mode of control. We have described the algorithms and illustrated how the PRNET provides highly reliable network transport and datagram service, by dynamically determining optimal routes, intelligently controlling congestion, and fairly allocating the channel in the face of changing link conditions, mobility, and varying traffic loads.

I. INTRODUCTION

In 1973, the Defense Advanced Research Projects Agency (DARPA) initiated research on the feasibility of using packet-switched, store-and-forward radio communications to provide reliable computer communications [1]. This development was motivated by the need to provide computer network access to mobile hosts and terminals, and to provide computer communications in a mobile environment. Packet radio networking offers a highly efficient way of using a multiple-access channel, particularly with bursty traffic [2]. The DARPA Packet Radio Network (PRNET) has evolved through the years to be a robust, reliable, operational experimental network [3]. The development process has been of an incremental, evolutionary nature [4]; as algorithms were designed and implemented, new versions of the PRNET with increased capabilities were demonstrated. The PRNET has been in daily operation for experimental purposes for nearly ten years. In this paper we describe the current state of the DARPA PRNET.

We begin by providing a synopsis of the PRNET system concepts, attributes, and physical components in Section II. In Section III, we illustrate the mechanisms by which a packet radio automatically keeps track of a potentially continuously changing network topology. In Section IV, we de-

Manuscript received February 1, 1986; revised July 30, 1986. This work of J. Jubin was supported by the Defense Advanced Research Projects Agency of the Department of Defense under contract MDA903-85-C-0205. The work of J. D. Tornow was supported by the Defense Advanced Research Projects Agency of the Department of Defense under Contract MDA903-85-C-0254.

J. Jubin is with Collins Defense Communications, Rockwell International, Richardson, TX 75081, USA.

J. D. Tornow is with SRI International, Menlo Park, CA 94025, USA.

scribe the algorithms used to route a packet through the packet radio communications subnet. In Section V, we examine the protocols for transmitting packets. In Section VI, we describe some of the hardware capabilities of the packet radio that strongly influence the design and characteristics of the PRNET protocols. We conclude by looking briefly at some applications of packet radio networks and by summarizing the state of the current technology.

II. DESCRIPTION OF THE PACKET RADIO SYSTEM

A. Broadcast Radio

The PRNET provides, via a common radio channel, the exchange of data between computers that are geographically separated. As a communications medium, broadcast radio (as opposed to wires and antenna-directed radio) provides important advantages to the user of the network. One of the benefits is mobility; a packet radio (PR) can operate while in motion. Second, the network can be installed or deployed quickly; there are no wires to set up. A third advantage is the ease of reconfiguration and redeployment. The PRNET protocols take advantage of broadcasting and common-channel properties to allow the PRNET to be expanded or contracted automatically and dynamically. A group of packet radios leaving the original area simply departs. Having done so, it can function as an autonomous group and may later rejoin the original network or join another group.

The broadcasting and common channel properties of radio do have disadvantages too. These properties, for all practical purposes, prohibit the building of a radio that is able to transmit and receive at the same time. Therefore, the PRNET protocols must attempt to schedule each transmission when the intended PR is not itself transmitting. Also, transmissions often reach unintended PRs and interfere with intended receptions. Therefore, the protocols must attempt to schedule each transmission when the intended PR is not receiving another PR's transmission.

B. Automated Network Management

The PRNET features fully automated network management. It is self-configuring upon network initialization, reconfigures upon gain or loss of packet radios, and has dy-

1990s: mobile ad hoc networks (MANET)

A Performance Comparison of Multi-Hop Wireless Ad Hoc Network Routing Protocols

Josh Broch David A. Maltz David B. Johnson Yih-Chun Hu Jorjeta Jetcheva

Computer Science Department
Carnegie Mellon University
Pittsburgh, PA 15213

<http://www.monarch.cs.cmu.edu/>

Abstract

An ad hoc network is a collection of wireless mobile nodes dynamically forming a temporary network without the use of any existing network infrastructure. In such networks, each node is responsible for the routing of wireless network interfaces, multiple network "hops" may be needed for one node to exchange data with another across the network. In recent years, a variety of new routing protocols targeted specifically at this environment have been developed, but little quantitative comparison on their protocol and performance has been made. This paper presents the results of a detailed packet-level simulation comprising four multi-hop wireless ad hoc network routing protocols that cover a range of design choices: DSDV, TORA, DSR, and AODV. We have extended the 2-hop model introduced in [1] to include the MAC and physical-layer behavior of the IEEE 802.11 wireless LAN standard, including a realistic wireless transmission channel model, and present the results of simulations of networks of 50 mobile nodes.

1 Introduction

In areas in which there is little or no communication infrastructure or the existing infrastructure is expensive or inconvenient to use, wireless mobile nodes will still be able to form a temporary network for the formation of an *ad hoc network*. Such a network, each mobile node operates not only as a host but also as a router, forwarding packets for other mobile nodes in the network that may not be within direct wireless transmission range of each other. Each node participates in an ad hoc routing protocol that allows it to discover "multi-hop" paths to other nodes in the network. The routing protocol is responsible for finding the shortest path to the destination. The routing protocol is sometimes also called a *routing algorithm*. The nodes in the network dynamically establish routing among themselves to form their own network "on the fly". Some examples of the possible uses of ad hoc networks include students using laptop computers to pass information to each other in a classroom, business people sharing information during a meeting, soldiers relaying information for situational awareness on the battlefield [12, 21], and emergency disaster relief personnel coordinating efforts after a hurricane or earthquake.

This work was supported in part by the National Science Foundation (NSF) under CAREER Award NCR-9204723, by the Air Force Material Command (AFMC) under DARPA contract number F105-93-8-0061, and by the AT&T Foundation under a Special Purpose Grant in Science and Engineering. David Maltz was also supported under an IBM Graduate Fellowship. This work was partially funded by an NSF Graduate Fellowship. The views and conclusions contained here are those of the authors and should not be interpreted as necessarily representing the official policies or endorsements, either expressed or implied, of NSR, AFMC, DARPA, the AT&T Foundation, IBM, Carnegie Mellon University, or the U.S. Government.

2.1 Physical and Data Link Layer Model
To accurately model the attenuations of radio waves between antennas close to the ground, radio engineers typically use a model that attenuates the power of a signal as $1/r^2$ at short distances (r is the distance between the antennas), and as $1/r^4$ at longer distances. The crossover point is called the *reference distance*, and is typically around 100 meters for outdoor low-gain antennas 1.5m above the ground plane operating in the 1–2GHz band [20]. Following this practice, our signal propagation model combines both a free-space propagation model and a two-ray ground reflection model. When a transmitter is within the reference distance of the receiver, we use

ns to allow accurate simulation of mobile wireless networks.

2.1 Physical and Data Link Layer Model

To accurately model the attenuations of radio waves between anten-

nas close to the ground, radio engineers typically use a model that attenuates the power of a signal as $1/r^2$ at short distances (r is the distance between the antennas), and as $1/r^4$ at longer distances. The crossover point is called the *reference distance*, and is typically around 100 meters for outdoor low-gain antennas 1.5m above the ground plane operating in the 1–2GHz band [20]. Following this practice, our signal propagation model combines both a free-space propagation model and a two-ray ground reflection model. When a trans-

mitter is within the reference distance of the receiver, we use



The Internet of Things: Roadmap to a Connected World

© 2016 Massachusetts Institute of Technology



0018-9219/87/0100-0021\$01.00 © 1987 IEEE

EXTENDING RANGE: MESH NETWORKS

Late 90s, 2000s: Sensor networks

Next Century Challenges: Scalable Coordination in Sensor Networks

660

IEEE TRANSACTIONS ON WIRELESS COMMUNICATIONS, VOL. 1, NO. 4, OCTOBER 2002

Deborah Estrin
Ramesh Govindan
John Heidemann
Satish Kumar

USC/Information Sciences Institute
4676 Admiralty Way
Marina del Rey, CA 90292, USA
estrin,govindan,johnh,kumar@isi.edu

Abstract

Networked sensors—those that coordinate amongst themselves to achieve a larger sensing task—will revolutionize information gathering and processing both in urban environments and in inhospitable terrain. The sheer numbers of these sensors and their distributed nature present unique challenges in the design of unattended autonomous sensor networks. These challenges lead us to hypothesize that sensor network coordination applications may need to be structured differently from traditional network applications. In particular, we believe that *localized algorithms* (in which simple local node behavior achieves a desired global objective) may be necessary for sensor network coordination. In this paper, we describe localized algorithms, and then discuss *distributed diffusion*, a simple communication model for describing localized algorithms.

1 Introduction

Integrated low-power sensing devices will permit remote object monitoring and tracking in many different contexts: in the field (vehicles, equipment, personnel), the office building (projectors, furniture, books, people), the hospital ward (patients, medical equipment, floors), and finally, in small robotic devices. Networking these sensors—empowering them with the ability to *coordinate* amongst themselves on a larger sensing task—will revolutionize information gathering and processing in many situations. Large scale, dynamically changing, and robust sensor colonies can be deployed in inhospitable physical environments such as remote geographic regions or toxic urban locations. They will also enable low maintenance sensing in most benign, but less accessible environments such as deep space, underground, etc.

To motivate the challenges in designing these sensor networks, consider the following scenario. Several thousand sensors are rapidly deployed (*e.g.*, thrown from an aircraft) in remote terrain. The sensors coordinate to establish a communication network, divide the task of mapping and monitoring the terrain amongst themselves in an energy-

efficient manner, adapt their overall sensing accuracy to the remaining total resources, and re-organize upon sensor failure. When additional sensors are added or old sensors fail, the sensors re-organize themselves to take advantage of the added system resources.

Several aspects of this scenario present systems design challenges that are unique to sensor network coordination networks (Section 2). The sheer numbers of these devices, and their unattended deployment, will preclude reliance on broadcast communication or the configuration currently needed to deploy and operate networked devices. Devices may be battery constrained or subject to hostile environments, so individual device failure will be a regular or common event. In addition, the configuration devices will frequently change in terms of position, reachability, power availability, and even task details. Finally, because these devices interact with the physical environment, they, and the network as a whole, will experience a significant range of task dynamics.

The WINES project [1] has considered device-level communication primitives needed to satisfy these requirements. However, these requirements potentially affect many other aspects of network design: routing and addressing mechanisms, naming and lifelong service management, architecture, and deployment methods, to name a few. This paper focuses on the principles underlying the design of services and applications in sensor networks. In particular, since the sensing is inherently distributed, we argue that sensor network applications will themselves be distributed.

Many of the lessons learned from Internet and mobile network design will be applicable to designing sensor network applications. However, this paper hypothesizes that sensor networks have different enough requirements at least initially to provide the overview of the types of applications and services. Specifically, we believe there are significant robustness and scalability advantages to designing applications using *localized algorithms*—where sensors only interact with other sensors in a restricted vicinity, but nevertheless collectively achieve a desired global objective (Section 3). We also describe *distributed diffusion*, a promising model for designing sensor applications (Section 4).

Our research project is starting to investigate the design of localized algorithms using the directed diffusion model. These ideas were developed in the context of a DARPA ISAT study, chaired by one of the authors (Estrin). The

Permission to make digital or hard copies of all or part of this work for personal use or internal institutional use is granted provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. To copy otherwise, to republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee.

Mobicom '99 Seattle Washington USA

An Application-Specific Protocol Architecture for Wireless Microsensor Networks

Wendi B. Heinzelman, Member, IEEE, Anantha P. Chandrakasan, Senior Member, IEEE, and Hari Balakrishnan, Member, IEEE

Abstract

Abstract—Networking together hundreds or thousands of cheap microsensors to perform a task in a harsh environment by intelligently combining the data from many individual nodes. These networks require robust wireless communication protocols that are energy efficient and provide low latency. In this paper, we develop and analyze low-energy adaptive clustering hierarchy (LEACH), a protocol architecture for microsensor networks that uses the idea of self-configuring clusters for data routing and media access together with application-specific data aggregation to achieve good performance in terms of system lifetime, latency, and application-perceived quality. LEACH includes a new, distributed cluster formation technique that enables self-organization of large sensor networks, allowing nodes to self-configure and route to cluster head positions to evenly distribute the energy load among all the nodes, and techniques to enable distributed signal processing to save communication resources. Our results show that LEACH can improve system lifetime by an order of magnitude compared with general-purpose multi-hop approaches.

Index Terms—Data aggregation, protocol architecture, wireless microsensor networks.

I. INTRODUCTION

ADVANCES in sensor technology, low-power electronics, and low-power radio frequency (RF) design have enabled the development of small, relatively inexpensive and low-power sensors, called *microsensors*, that can be connected via a wireless network. These wireless microsensor networks represent a new paradigm for extracting data from the environment and enable the reliable monitoring of a variety of environments for applications that include surveillance, machine failure diagnosis, and chemical/biological detection. An important challenge in the design of these networks is that two key resources—communication bandwidth and energy—are significantly more limited than in a tethered network environment. These constraints require innovative design techniques to use the available bandwidth and energy efficiently.

Manuscript received January 9, 2001; revised July 1, 2001 and August 24, 2001; accepted August 24, 2001. The editor coordinating the review of this paper and approving it for publication was M. Zorzi. The work of W. B. Heinzelman was supported by a Kodak Research Fellow. This work was supported in part by the Defense Advanced Research Project Agency (DARPA) Power Aware Computing Communication Program and the U.S. Air Force Research Laboratory, Air Force Materiel Command, under Agreement F30602-00-2-0551.

W. B. Heinzelman was with the Massachusetts Institute of Technology, Cambridge, MA 02139 USA. She is now with the Department of Electrical and Computer Engineering, University of Rochester, Rochester, NY 14627-0216 USA (e-mail: whinezel@ece.rochester.edu).

A. P. Chandrakasan and H. Balakrishnan are with the Massachusetts Institute of Technology, Cambridge, MA 02139 USA (e-mail: anantha@mit.edu); Digital Object Identifier 10.1109/TWC.2002.804190

The Internet of Things: Roadmap to a Connected World

© 2016 Massachusetts Institute of Technology

2000s: Mesh networks for Internet

Architecture and Evaluation of an Unplanned 802.11b Mesh Network

John Bicket, Daniel Aguayo, Sanjit Biswas, Robert Morris
M.I.T. Computer Science and Artificial Intelligence Laboratory
jbicket, aguayo, biswas, rtm @csail.mit.edu

ABSTRACT

This paper evaluates the ability of a wireless mesh architecture to provide high performance Internet access while demanding little deployment planning or operational management. The architecture considered in this paper has unplanned node placement (rather than planned topology), omni-directional antennas (rather than directional links), and multi-hop routing (rather than single-hop base stations). These design decisions contribute to ease of deployment, an important requirement for community wireless networks. However, this architecture carries the risk that lack of planning might render the network's performance unusually low. For example, it might be necessary to place nodes特意 to ensure connectivity; the omni-directional antennas might provide uselessly short radio ranges; or the inefficiency of multi-hop forwarding might leave some users effectively disconnected.

The paper evaluates this unplanned mesh architecture with a case study of the Roofnet 802.11b mesh network. Roofnet consists of 37 nodes spread over four square kilometers of an urban area. The network provides users with usable performance despite lack of planning: the average inter-node throughput is 1.5 kbytes/second, even though the average node has three hops.

The paper analyzes multiple aspects of the architecture: the effect of node density on connectivity and throughput; the characteristics of the links that the routing protocol elects to use; the usefulness of the highly connected mesh afforded by omni-directional antennas for robustness and throughput; and the potential performance of a single-hop network using the same nodes as Roofnet.

Categories and Subject Descriptors

C.2.1 [Computer-Communication Networks]: Network Architecture and Design—Wireless communication; C.2.2 [Computer-Communication Networks]: Network Protocols—Routing protocols

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. To copy otherwise, to republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee.

MobileCom '05, August 28–September 2, 2005, Cologne, Germany.
Copyright 2005 ACM 1-59593-020-5/05/0008 \$3.00.

General Terms

Design, Experimentation, Measurement, Performance

Keywords

Mesh networks, Multi-hop wireless networks, Ad hoc networks, Wireless routing, Route metrics

1. INTRODUCTION

Community wireless networks typically share a few wired Internet access points that are spread out in an urban area. Two approaches to constructing community networks are common. The first approach is to carefully construct a multi-hop network with nodes in chosen locations and directional antennas aimed to engineer high-quality radio links [31, 8, 29]; these networks require well-coordinated groups with technical expertise, but result in high throughput and good connectivity. The second approach consists of individuals operating “hot-spot” access points to which clients directly connect [5, 4]. These access points often operate independently and are loosely connected, if at all. Access-point networks do not require much coordination to deploy and operate, but usually do not provide much coverage per wired connection as multi-hop networks.

A more ambitious vision for community networks would combine the best characteristics of both network types, operating without extensive planning or central management but still providing wide coverage and acceptable performance. This paper provides an evaluation of such an architecture, consisting of the following design decisions:

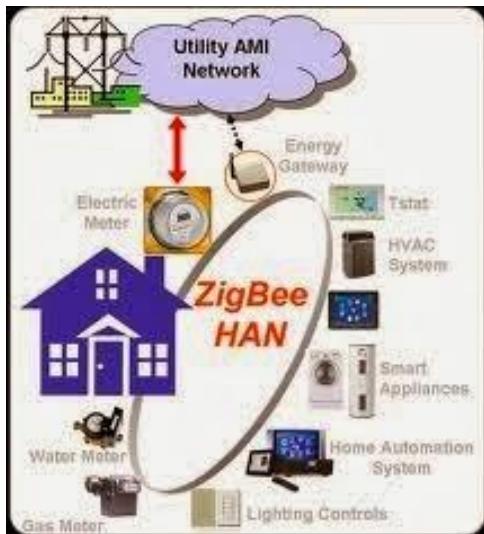
1. Unconstrained node placement, rather than a topology planned for coverage or performance. The network should work well even if the topology is determined solely by where participants happen to live.
2. Omni-directional antennas, rather than directional antennas used to form particular high-quality links. Users should be able to move and antenna without knowing in advance what nodes the antenna might talk to. Nodes should be able to route data through whatever neighbors they happen to find.
3. Multi-hop routing, rather than single-hop base stations or access points. Multi-hop routing can improve coverage and performance despite lack of planning and lack of specifically engineered links.



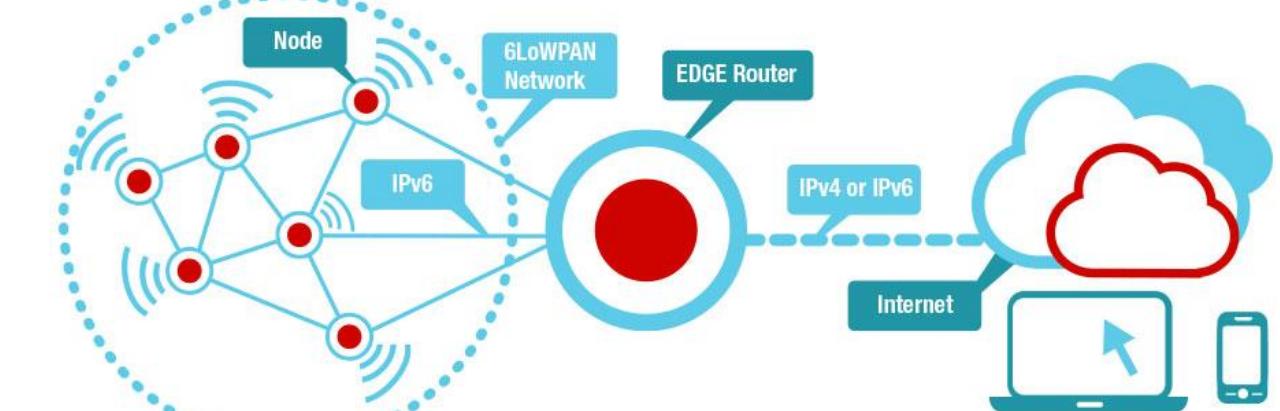
EXTENDING RANGE: MESH NETWORKS

2010s: Mesh networks for IoT

Zigbee



6LoWPAN: IPv6 over low-power wireless personal area networks

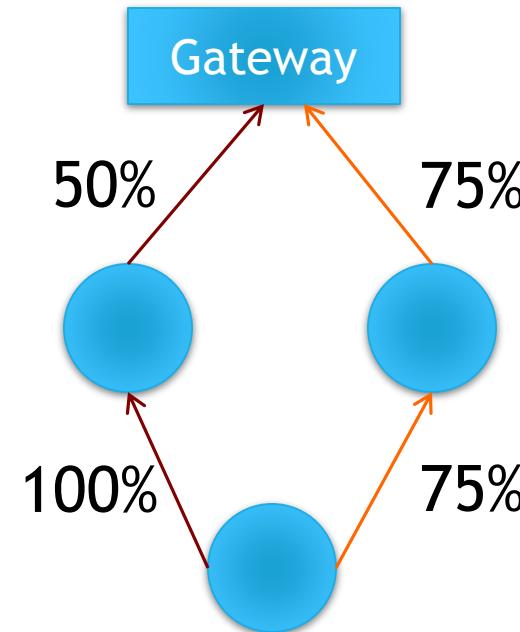
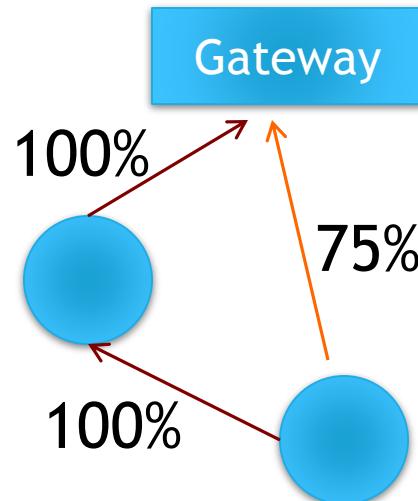


<http://processors.wiki.ti.com/index.php/Contiki-6LOWPAN> (Creative commons)

Both (typically) run over the 802.15.4 MAC standard
Routing protocol with different metrics, such as “expected transmission time”
Use case: devices communicating with gateway across multiple hops
Node duty cycles higher, some nodes do much more work

WIRELESS MESH ROUTING

Each link delivers a packet with some probability
What path should be used for a packet?

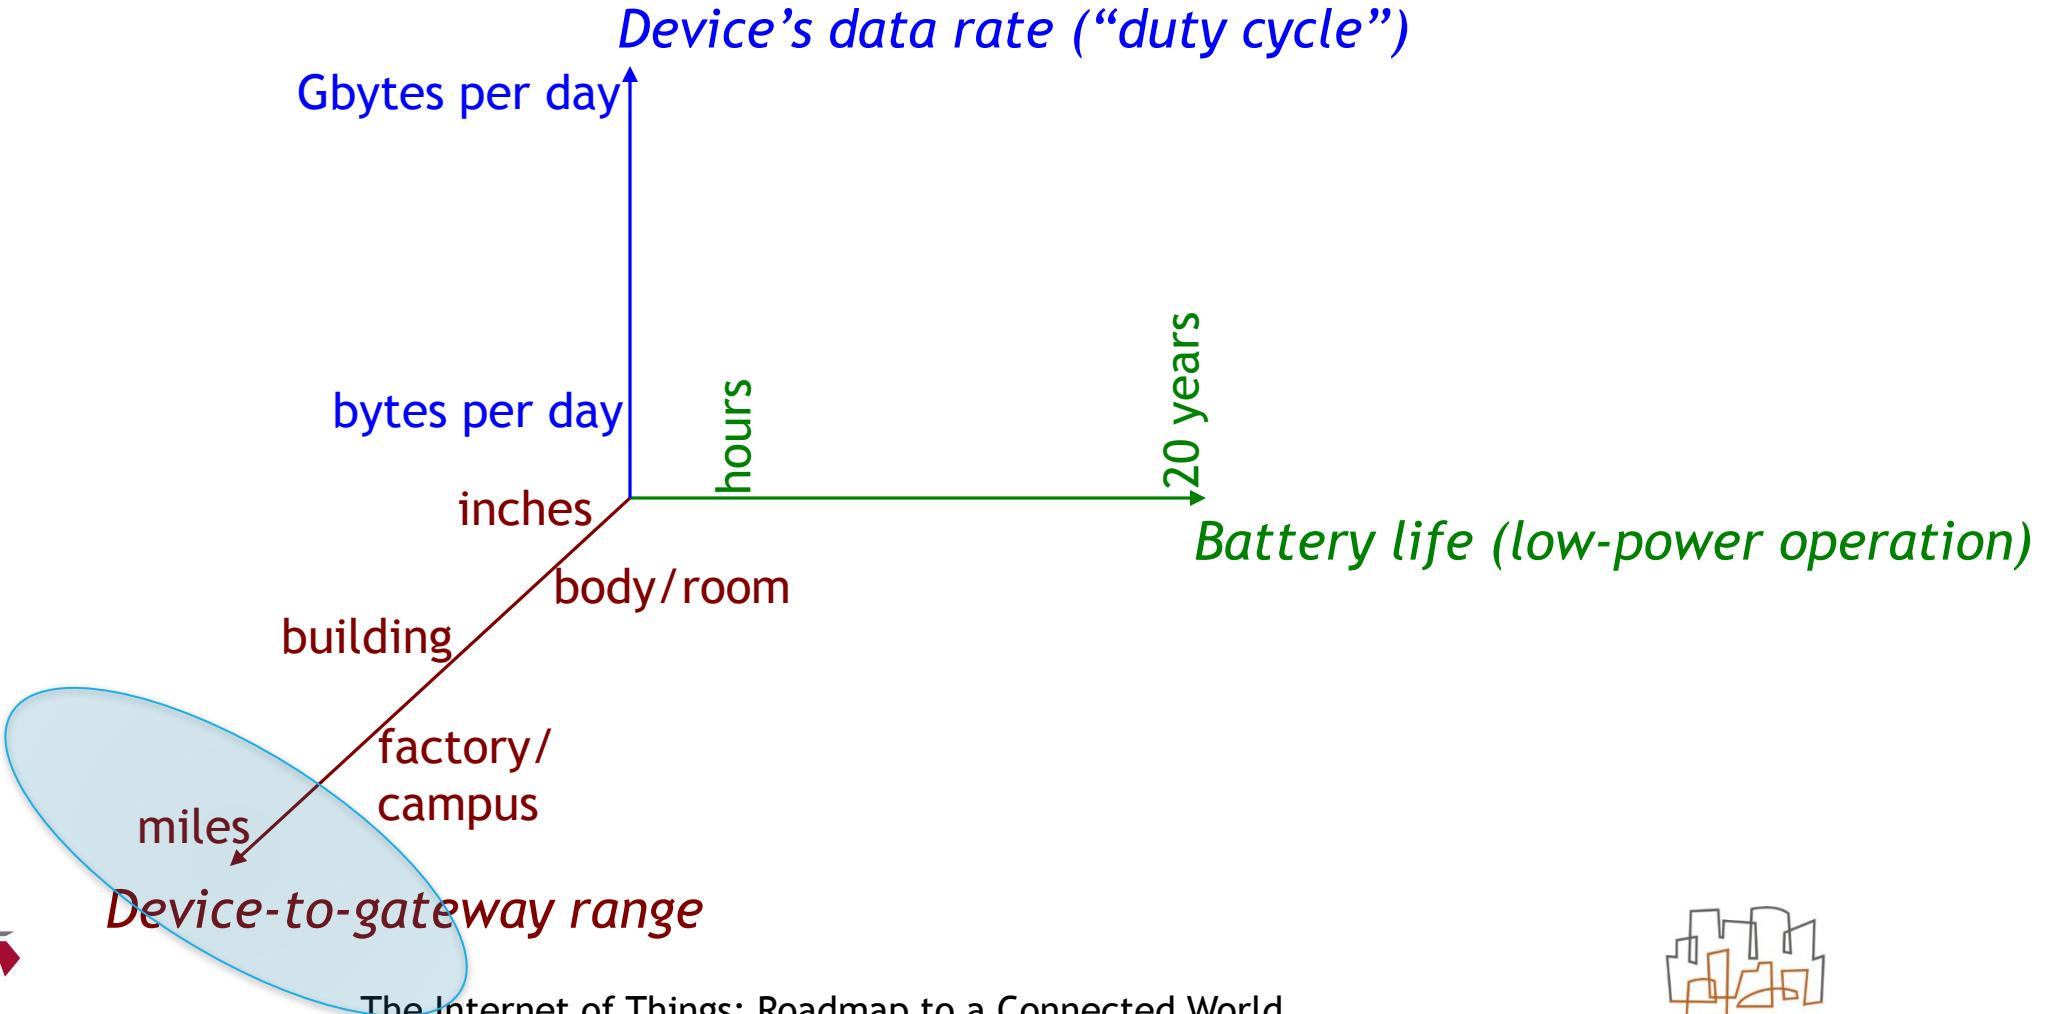


ETX: Expected Transmission Count
(expected tx time if link rate fixed)

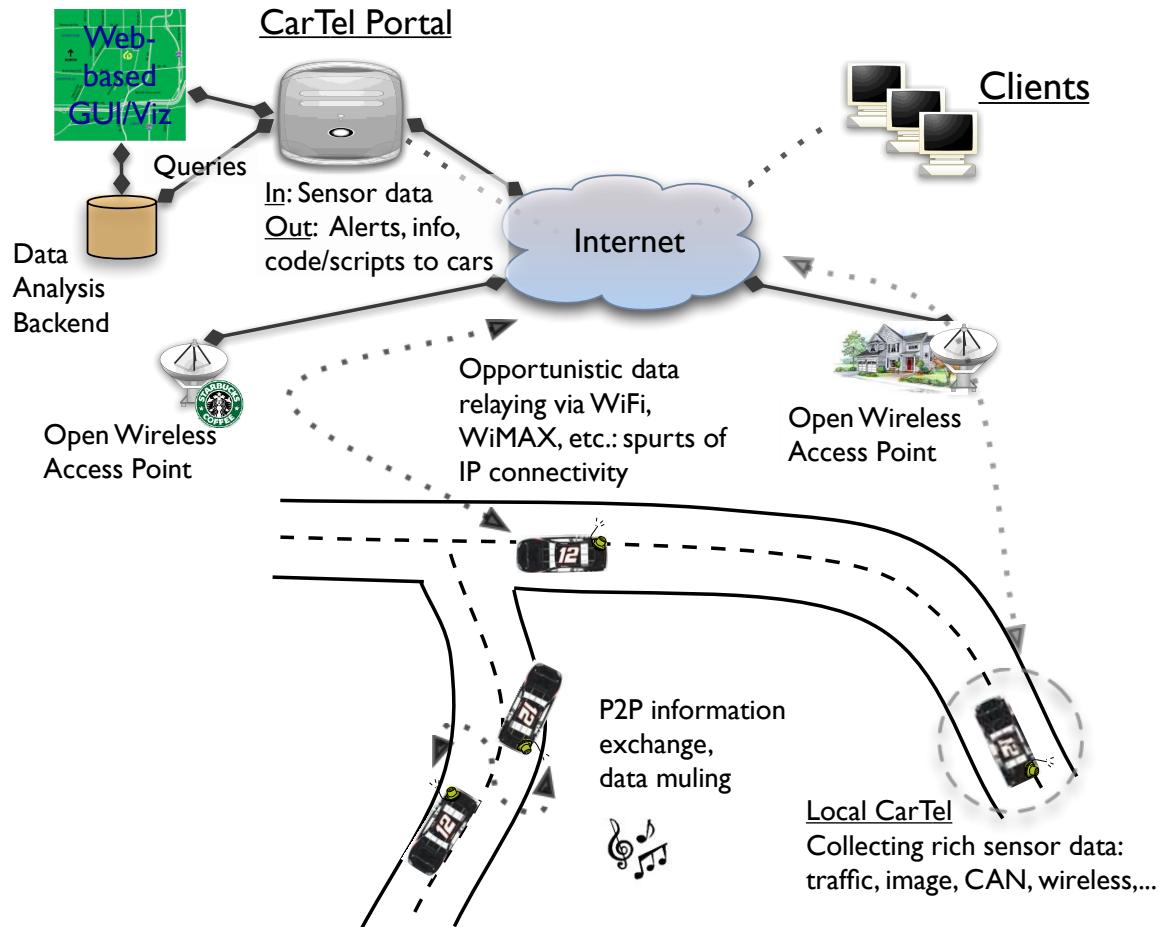
On single link =
 $1 / \text{Probability of packet delivery}$

Add for each link on path and pick
smallest

EVEN LONGER RANGE (CITY-SCALE)



WHEN THE INTERNET IS MILES AWAY



Use mobile devices
as **data mules**
Trade-off: delay
Delay-tolerant network (DTN)



WHAT IF WE WANT LONG RANGE AND LOW DELAY?

“Long-range IoT networks”

Examples: Sigfox, LoRaWAN, cellular IoT proposals
(narrowband LTE, etc.)

Low-power designs (months to years of battery life)

Low or ultra-low throughput (a few bytes per day to achieve long-enough battery life at a rate of a few kbps)

Networks like LoRaWAN also include localization capabilities



WHAT IF WE WANT LONG RANGE AND LOW DELAY

Second choice: Cellular (of course!)

Examples: LTE/4G, 3G, etc.

High-power consumption, so only when power isn't an issue

Delay still a concern for **data-intensive, latency-sensitive applications**



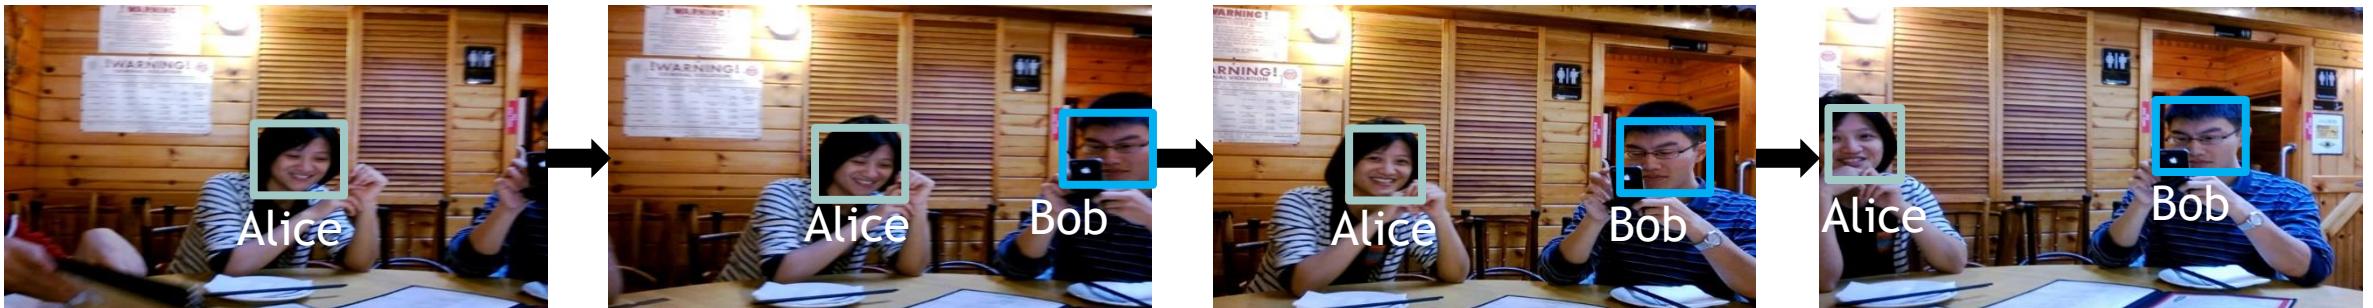
CONTINUOUS RECOGNITION APPS



The Internet of Things: Roadmap to a Connected World

GLIMPSE: CONTINUOUS REAL-TIME RECOGNITION

- Continuous, real-time object recognition on mobile devices in a video stream
- Continuously *identify* and *locate* objects in each frame



T. Chen, L. Ravindranath, S. Deng, P. Bahl, H. Balakrishnan, SenSys 2015

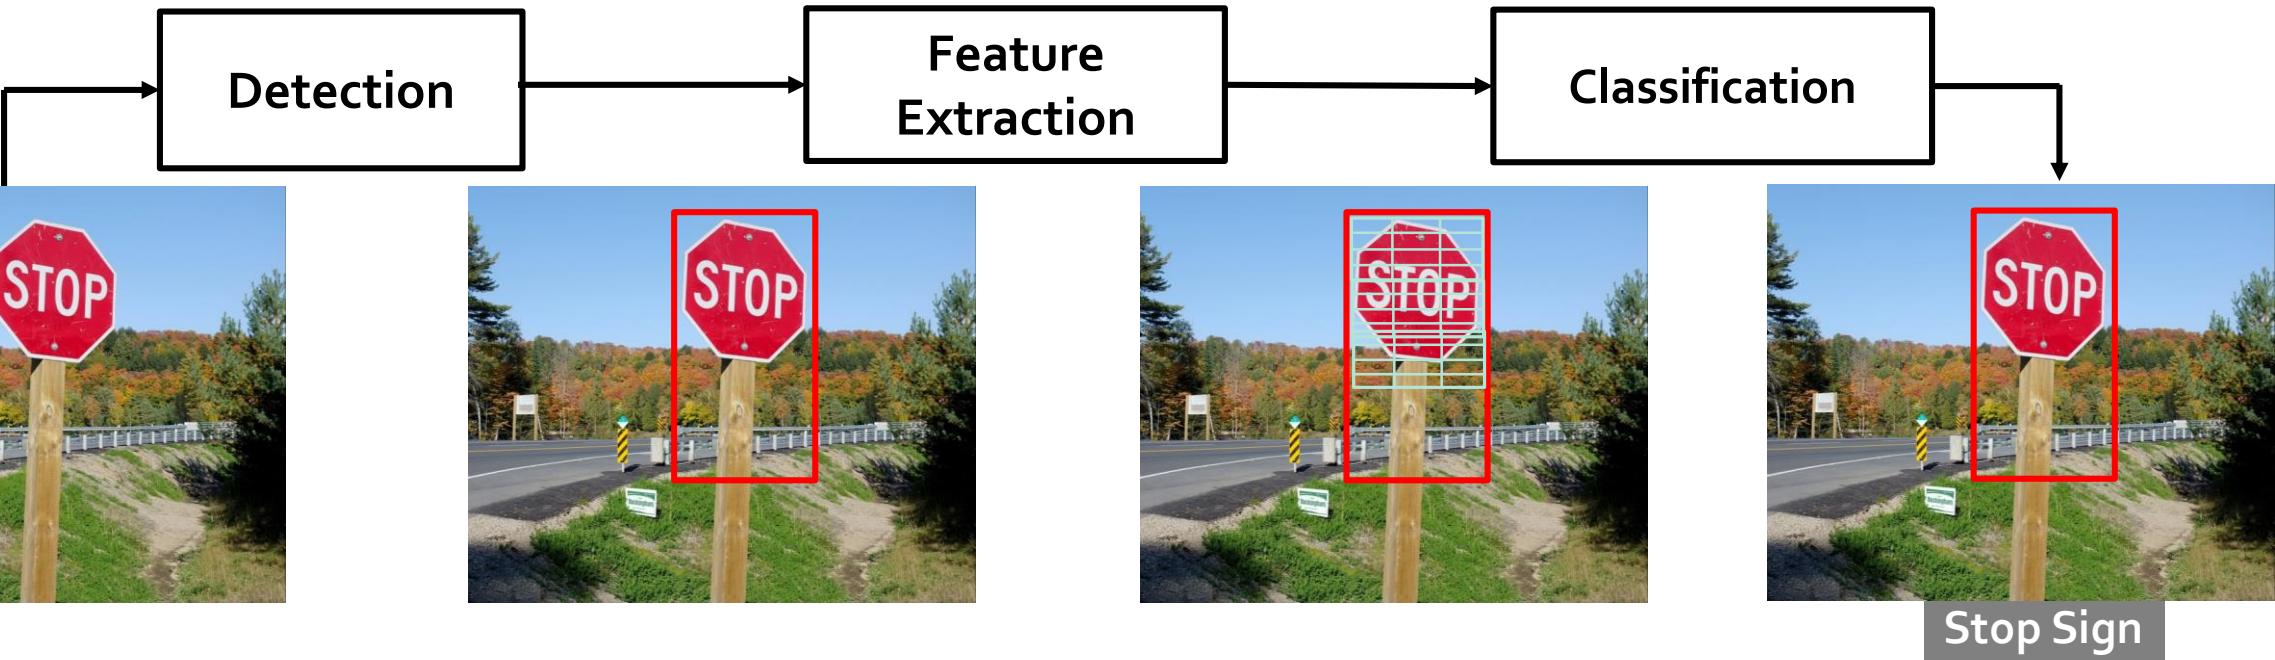


The Internet of Things and the Connected World

© 2016 Massachusetts Institute of Technology



RECOGNITION PIPELINE



T. Chen, L. Ravindranath, S. Deng, P. Bahl, H. Balakrishnan, SenSys 2015

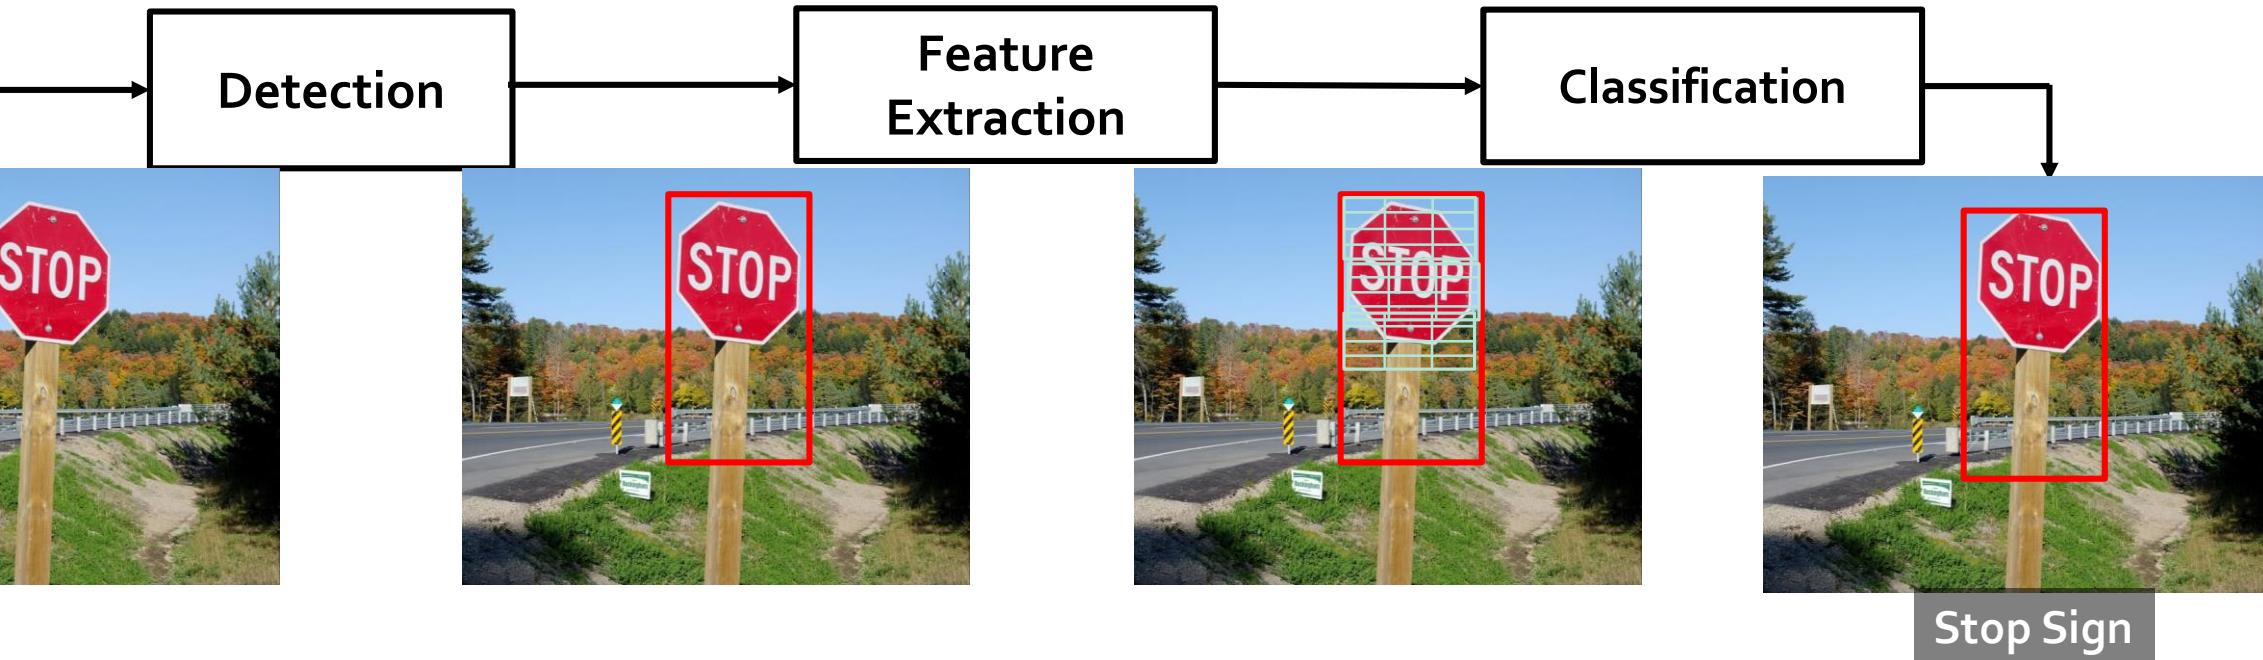


The Internet of Things: Roadmap to a Connected World

© 2016 Massachusetts Institute of Technology



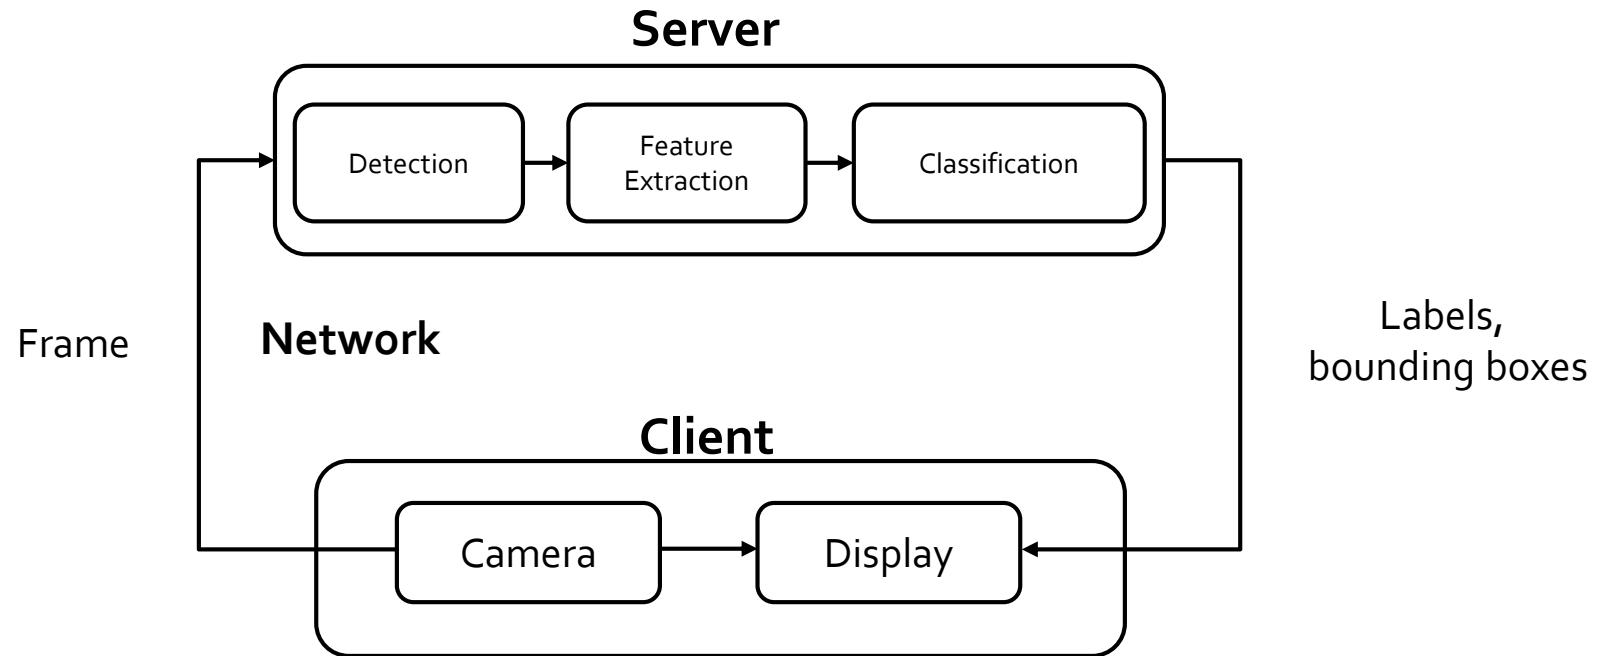
RECOGNITION PIPELINE



Computationally expensive and memory-intensive

- Server is hundreds of times faster than embedded/mobile computer
- Scaling to large number of images requires servers

SOLUTION: SERVER OFFLOAD

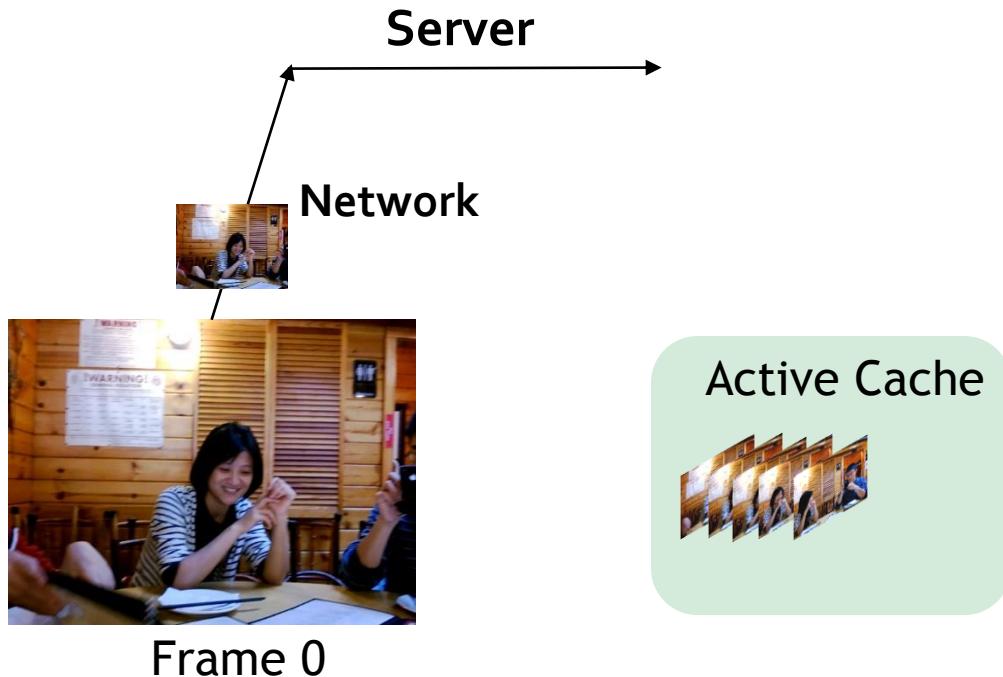


Two big challenges

1. End-to-end latency lowers object recognition accuracy
2. Bandwidth and battery-efficiency

IDEA: DEVICE-SIDE ACTIVE CACHE

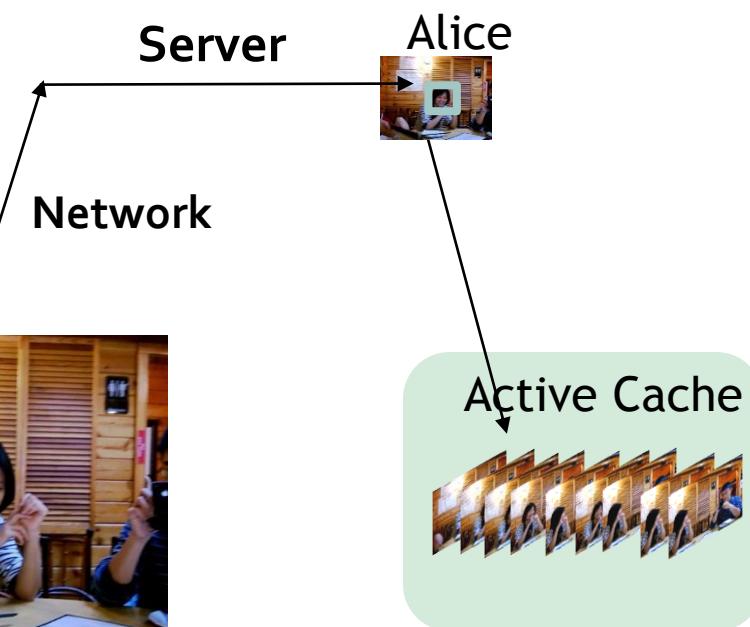
Cache and run tracking through the cached frames



The Internet of Things: Roadmap to a Connected World

© 2016 Massachusetts Institute of Technology





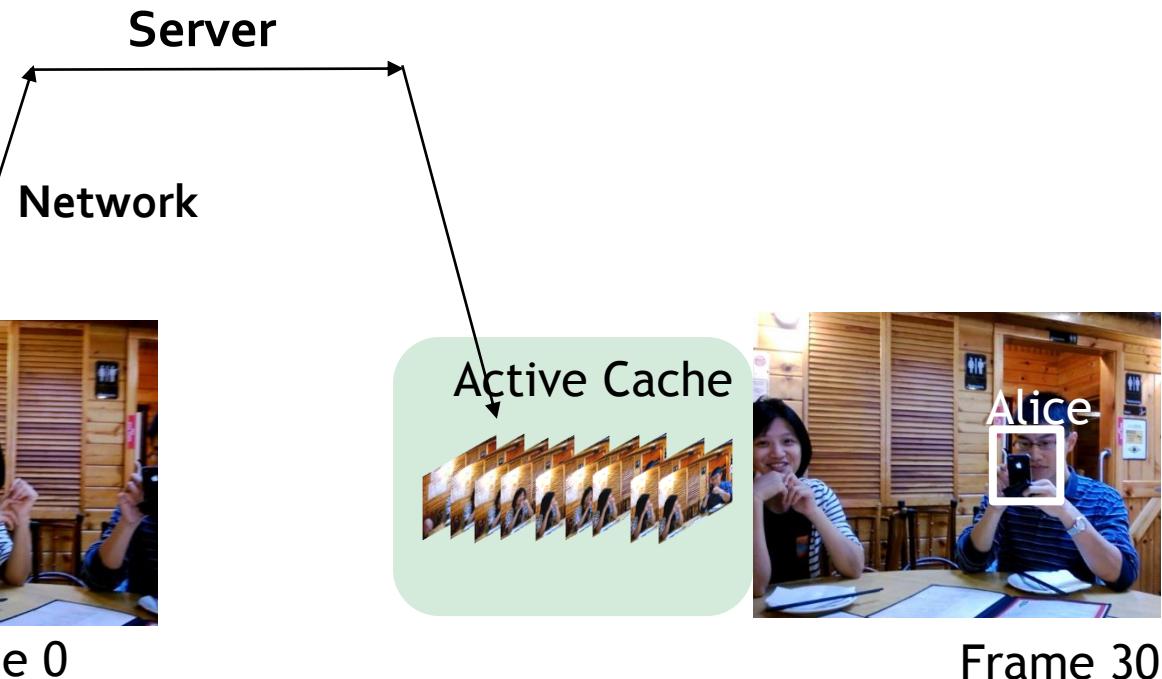
Frame 0



The Internet of Things: Roadmap to a Connected World

© 2016 Massachusetts Institute of Technology





The Internet of Things: Roadmap to a Connected World

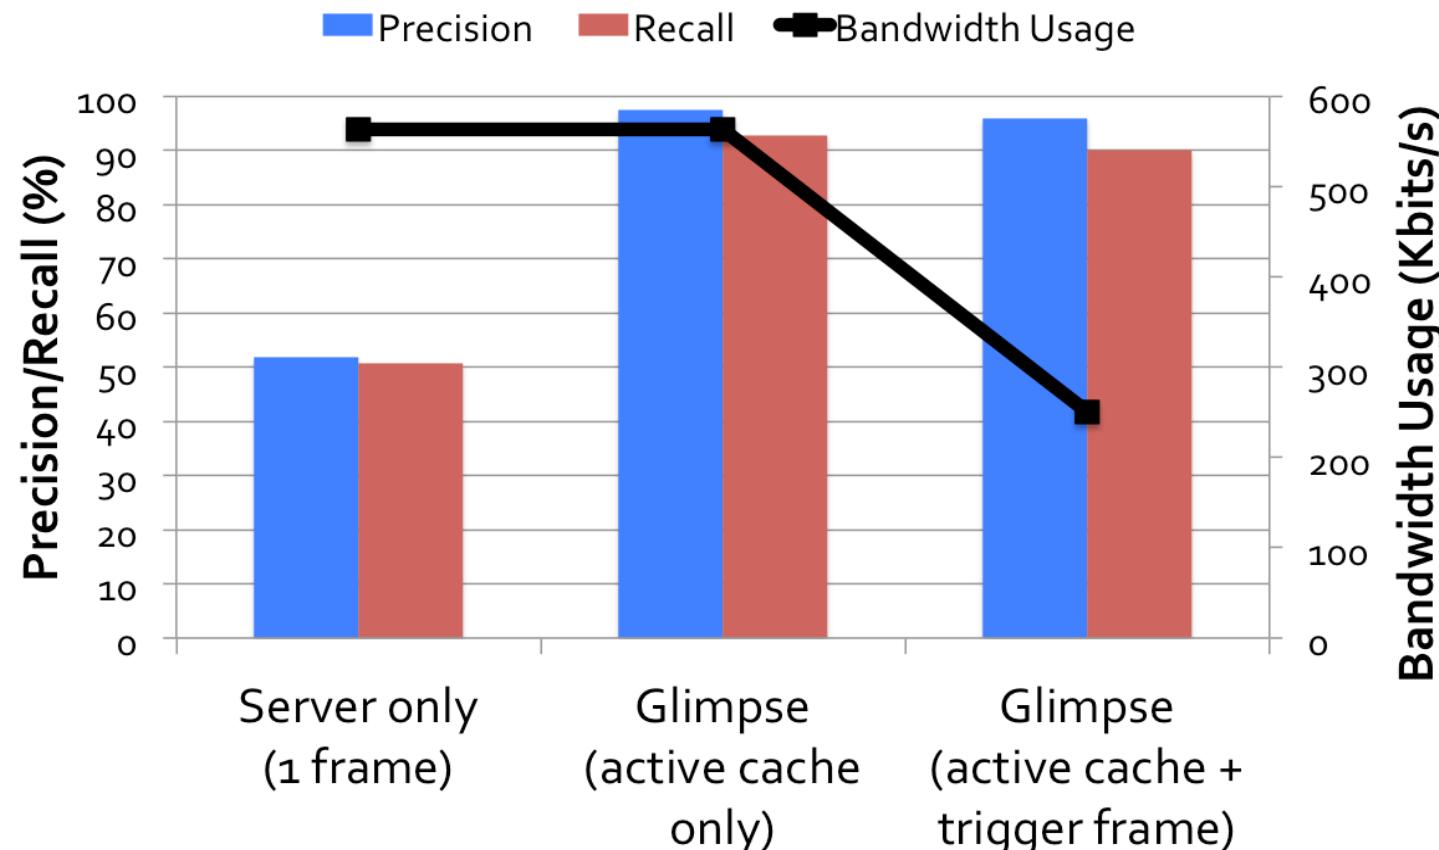
© 2016 Massachusetts Institute of Technology

Run tracking from
Frame 0 to Frame 30



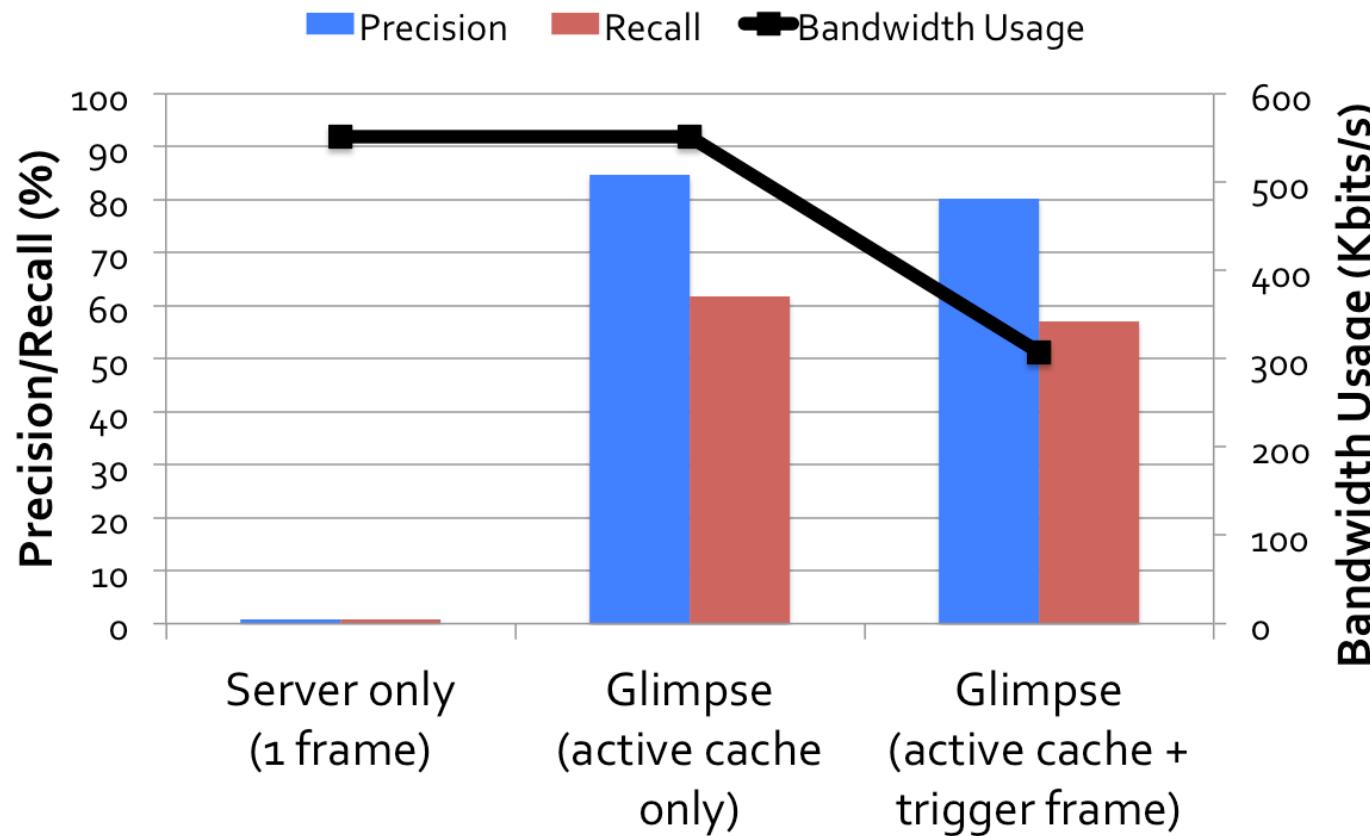
PERFORMANCE ON FACES

Average end-to-end frame delay to server and back: 430 ms



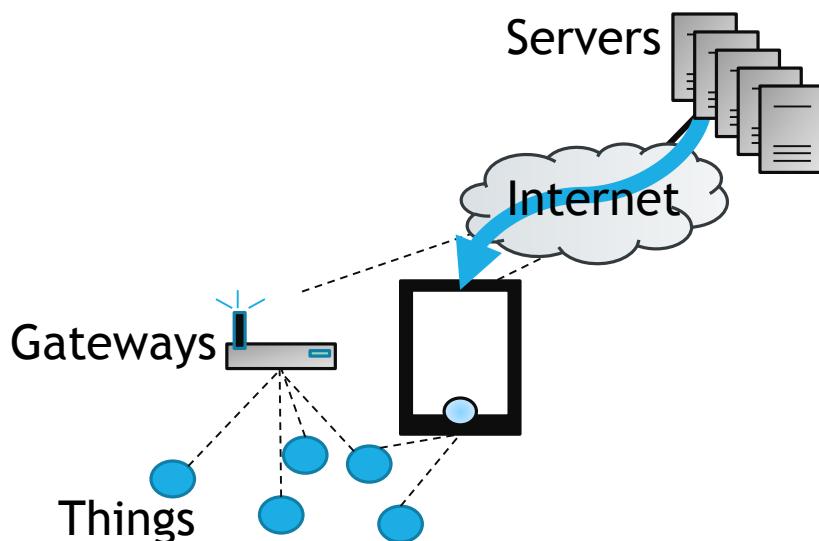
PERFORMANCE ON ROAD SIGNS

Average end-to-end frame delay to server and back: 520 ms



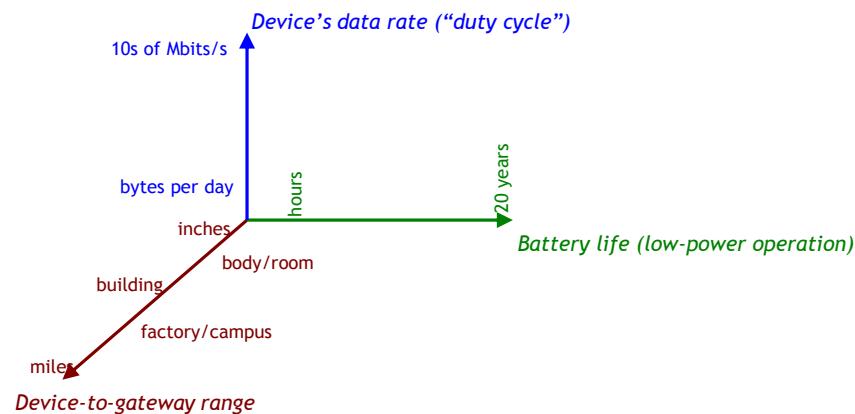
WHAT HAVE WE LEARNED?

Rich design space for things-gateway communication



Think along three dimensions:

1. data rate/duty cycle
2. battery
3. range



Three case studies

1. Low-power design (Bluetooth LE): advertisement, time-scheduled MAC
2. Range extension techniques: muling & meshing (Zigbee, 6LoWPAN)
3. Data-intensive IoT: continuous recognition

OPEN QUESTIONS AND FUTURE WORK

We are not even at the end of the beginning of IoT!
(Who wants to turn on a light bulb by going to an app?!)

What if you want city-scale, high-rate, low-power sensing?
(e.g., high-fidelity vibration, weather, image sensors)

Current systems gated by standby power (microWatts)
Recent advances have shown nanoWatt standby power
How will this change IoT networks?

Current IoT apps are “siloed” from each other
How to integrate them?



The Internet of Things: Roadmap to a Connected World

© 2016 Massachusetts Institute of Technology



DE-SILOING

Today: build IoT devices/sensors, build an app, build a cloud service

Vertically-integrated: hard to integrate and slows innovation

Gateway functions are repeatedly invented

The issue: real value is in the data, not in the devices!

Possible (non-exclusive) approaches

1. Coordinate access to data via server-side APIs in the cloud
2. Provide access to data in smartphone apps via “kits”
(HomeKit, Healthkit, Google Fit, ...)
3. Develop a generic gateway (multiple technologies)



PREDICTIONS

1. Shake-up in standards: multiple winners, but they will divide up the “three-dimensional space”
2. Ultra-low power IoT systems and networks
3. Compute-intensive (data-intensive) IoT systems and networks
4. De-siloed architectures
5. Smartphone-centric v. hidden (ubiquitous) computing



The Internet of Things: Roadmap to a Connected World

© 2016 Massachusetts Institute of Technology



The Internet of Things: Roadmap to a Connected World

THANK YOU!

HARI BALAKRISHNAN

Fujitsu Professor, EECS Department

Computer Science and Artificial Intelligence Laboratory (CSAIL)
Massachusetts Institute of Technology



The Internet of Things: Roadmap to a Connected World

© 2016 Massachusetts Institute of Technology





HARI BALAKRISHNAN

Fujitsu Professor of EECS



The Internet of Things: Roadmap to a Connected World

© 2016 Massachusetts Institute of Technology

