

# CYBERSECURITY

## Systems Security



PROFESSIONAL  
EDUCATION



MIT COMPUTER SCIENCE AND ARTIFICIAL INTELLIGENCE LABORATORY



## Hardware for Security

Howard Shrobe

Principal Research Scientist, Director Security@CSAIL

Computer Science and Artificial Intelligence Laboratory (CSAIL)

Massachusetts Institute of Technology



# Overview

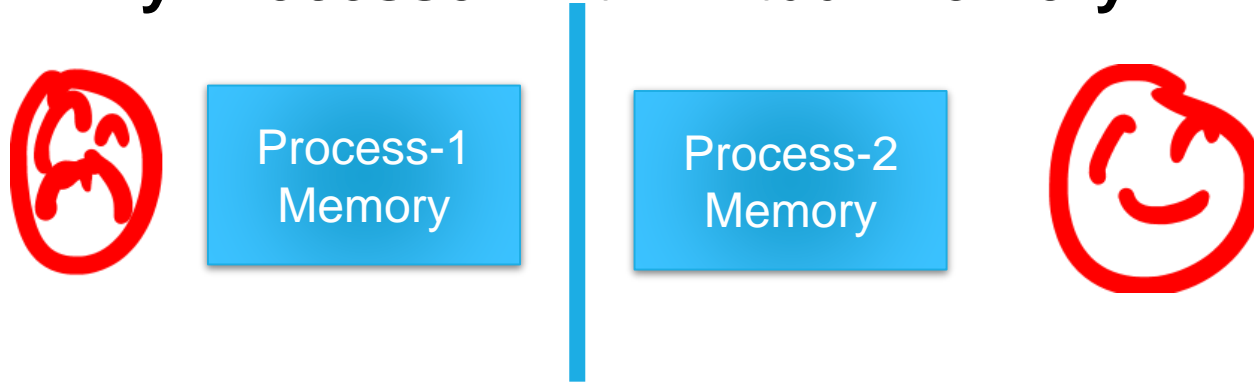
# Why should we think about hardware-based protection?

- Current computers already have hardware-based protections
- New features are in the pipeline
- There are important properties that can be enforced more systematically in hardware
- There are important properties that can only be done efficiently in hardware
- Hardware can enforce the intended semantics of your software


# Examples of hardware-based protection

## Process Isolation


Any Processor with Virtual Memory



## Data Protection - Ascend




## Enforcement of Access Policies



**Policy:**

- High processes can access anything
- Low processes can only access low data

## Memory Safety Intel-MPX



# Agenda

1. Learning from the past: Multics
2. Examples of what can go wrong
3. Tagged architectures
  1. Memory safety
  2. Type safety
  3. Information flow
  4. “Zero Kernel”
4. Capability architectures

# Historically Interesting Machines

- Multics: 1964 (last machine decommissioned in 2000)
  - Segments, Rings
- Burroughs B5500/6500: 1967
  - Data type tags
- Cambridge CAP Machine: 1970 – 1977
  - Capability machine
- MIT Lisp Machine: 1973 (last machine still running)
  - Data type tags, bounds checking
- IBM System 38: 1978 (now IBM System I)
  - Capability Machine – current product

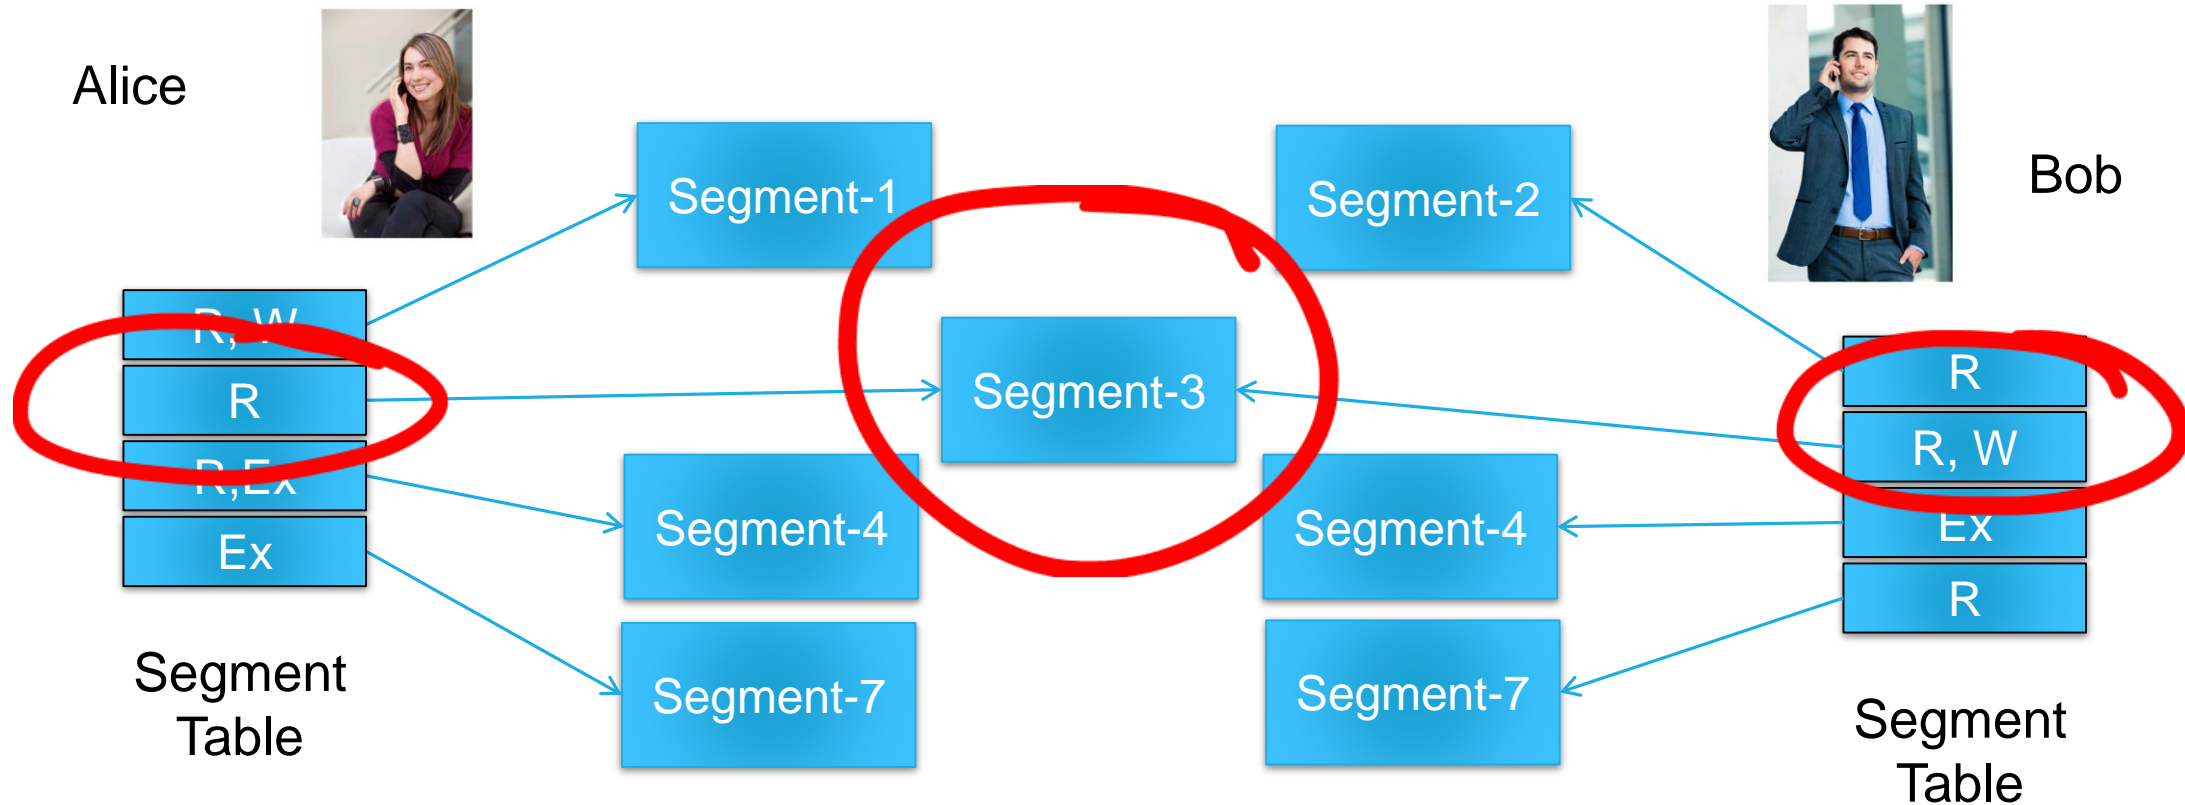
# Multics: An early secure hardware base



# The Multics Model: Controlled Sharing

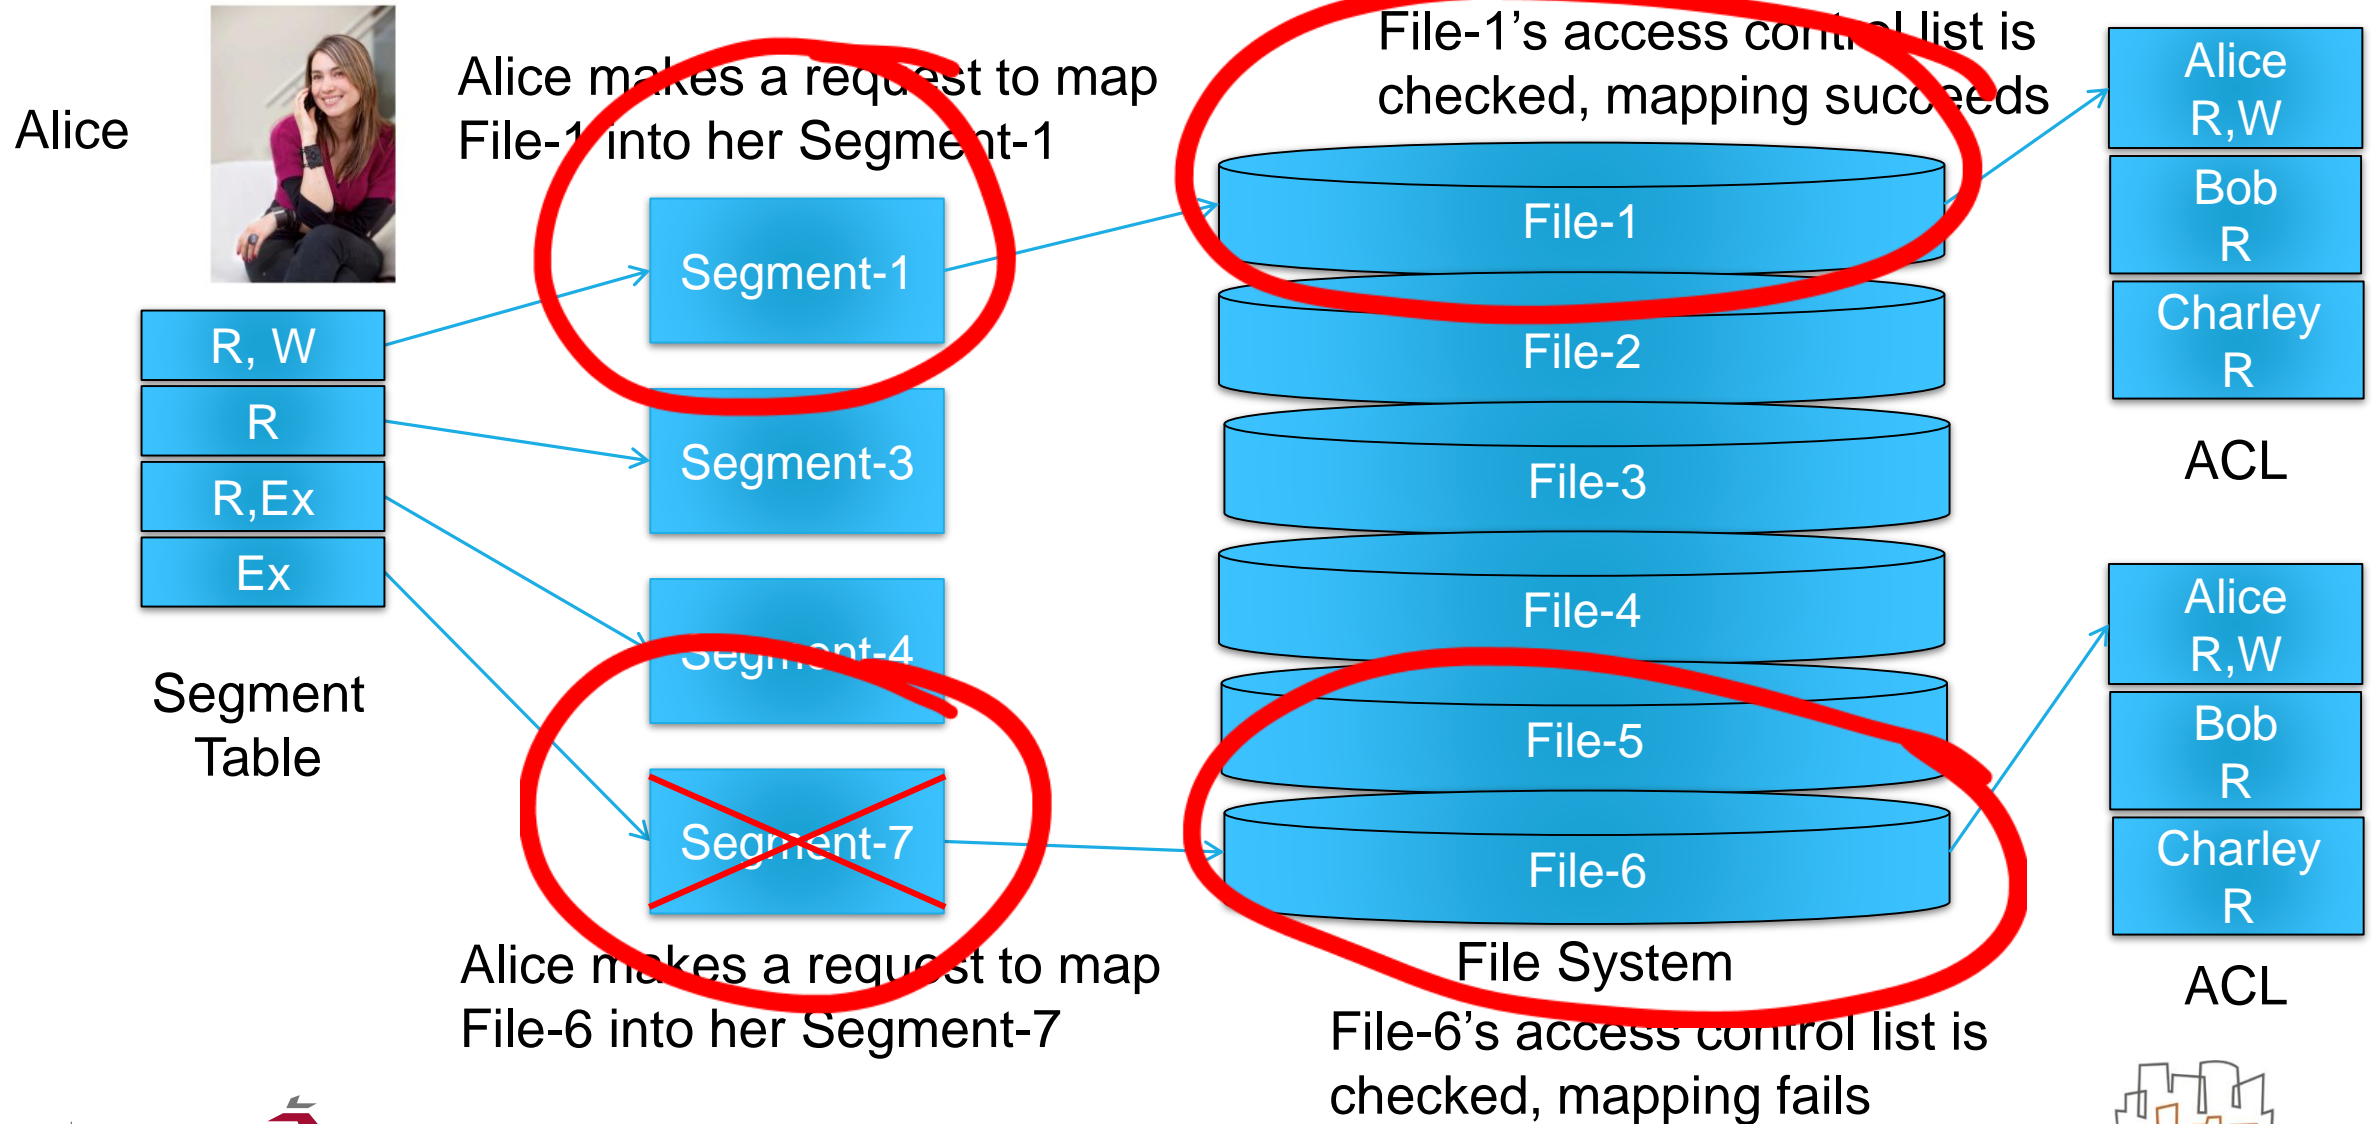
- **Multics** was an early MIT computer architecture project (started in 1964)
- Multics was conceived as a *computer utility*, supporting many users
- Users needed to share resources to support collaboration
- Users also required protection of their private data

# Multics: Segmented Memory & Permissions




- Memory addresses are two dimensional: Segment #, Address within segment
- Segments are composed of virtual memory pages
- Each user has a segment table, which points to a page table

# Multics: Segments, Files, Access Lists



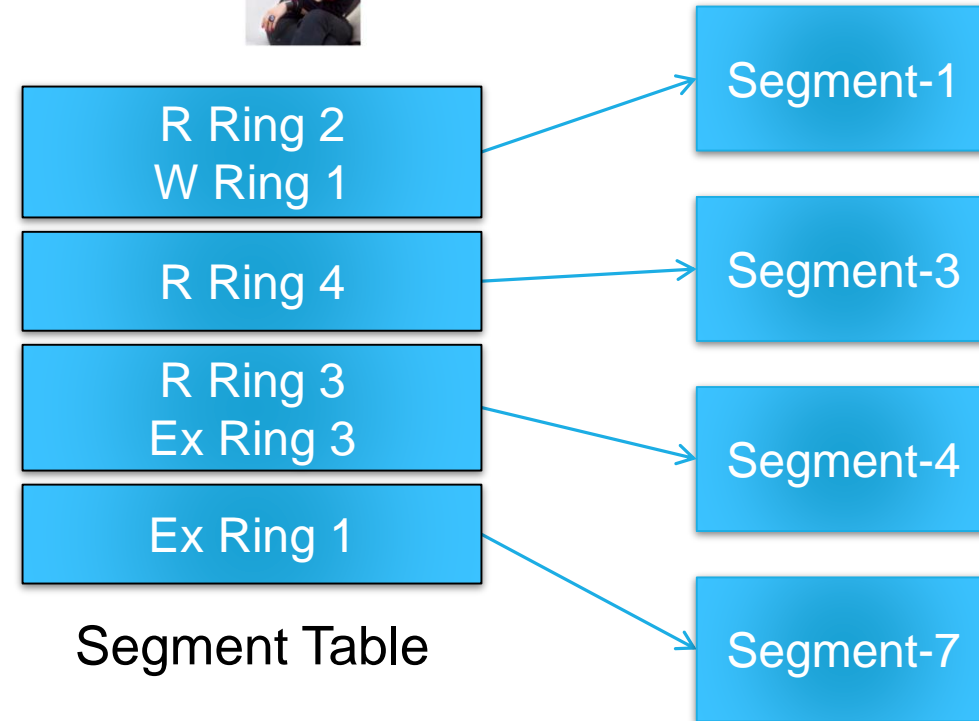
# Multics: Rings of Protection

- Each ring is strictly more privileged than the next
- Ring-0 is the most privileged



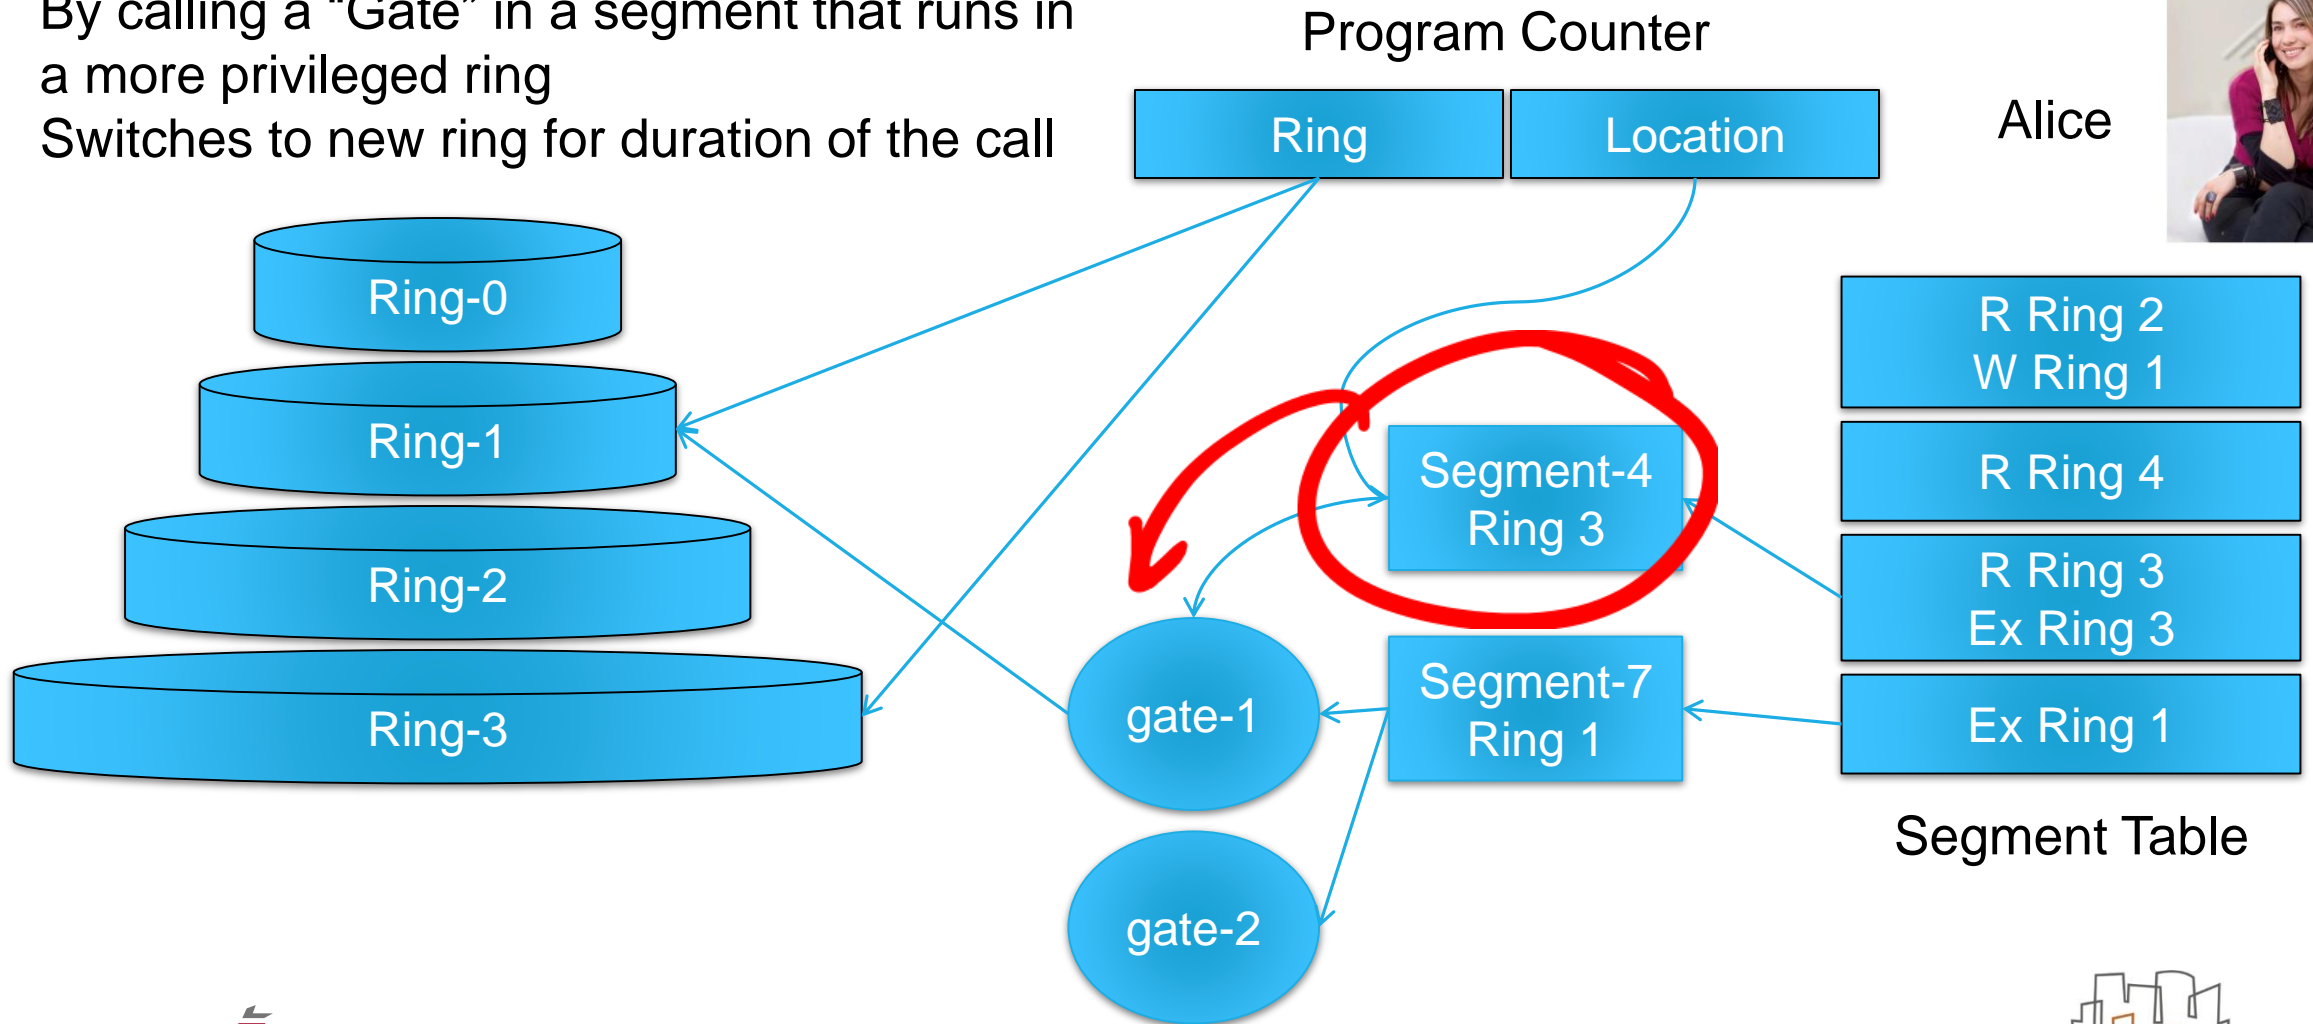
- Alice can read Segment-1 if she's executing in Ring 0,1 or 2
- Alice can write Segment-1 but only if she's executing in Ring 0 or 1

Alice



# Multics: Gates for changing privilege levels

- How do you transition between rings?
- By calling a “Gate” in a segment that runs in a more privileged ring
- Switches to new ring for duration of the call





# Key Principles

1. Complete Mediation
2. Separation of Privilege
3. Least Privilege
4. Economy of Mechanism
5. Acceptability & Productivity


Saltzer & Schroeder: The Protection of Information in Computer Systems

Revised version in *Communications of the ACM* 17, 7 (July 1974)

# The Multics Legacy


- For its day the Multics hardware was too expensive & used too many resources
- Unix was created as cheaper, simpler system, but with fewer guarantees
  - No segments, permission on pages
  - User & Kernel modes; mode switch is expensive
  - Kernel mode is totally privileged
  - C programming language for efficiency
- We've continued that trend ever since

# What do attackers do?



# What can go wrong. Examples of attacks:

- Memory corruption
- Data disclosures
- Code injection
- Control flow diversion
- Return-oriented programming



# Memory corruption exploits


Suppose your program is copying data from the network to an array:

Computer, here is A. Copy A to C



No problem.

Computer, here is B. Copy B to C



- If bounds aren't checked, there is a big problem, D is over-written
- Bounds that aren't represented can't be checked

Aleph One. Smashing the Stack for Fun and Profit. Phrack, 7(49), November 1996.



# Heart bleed memory disclosure attack

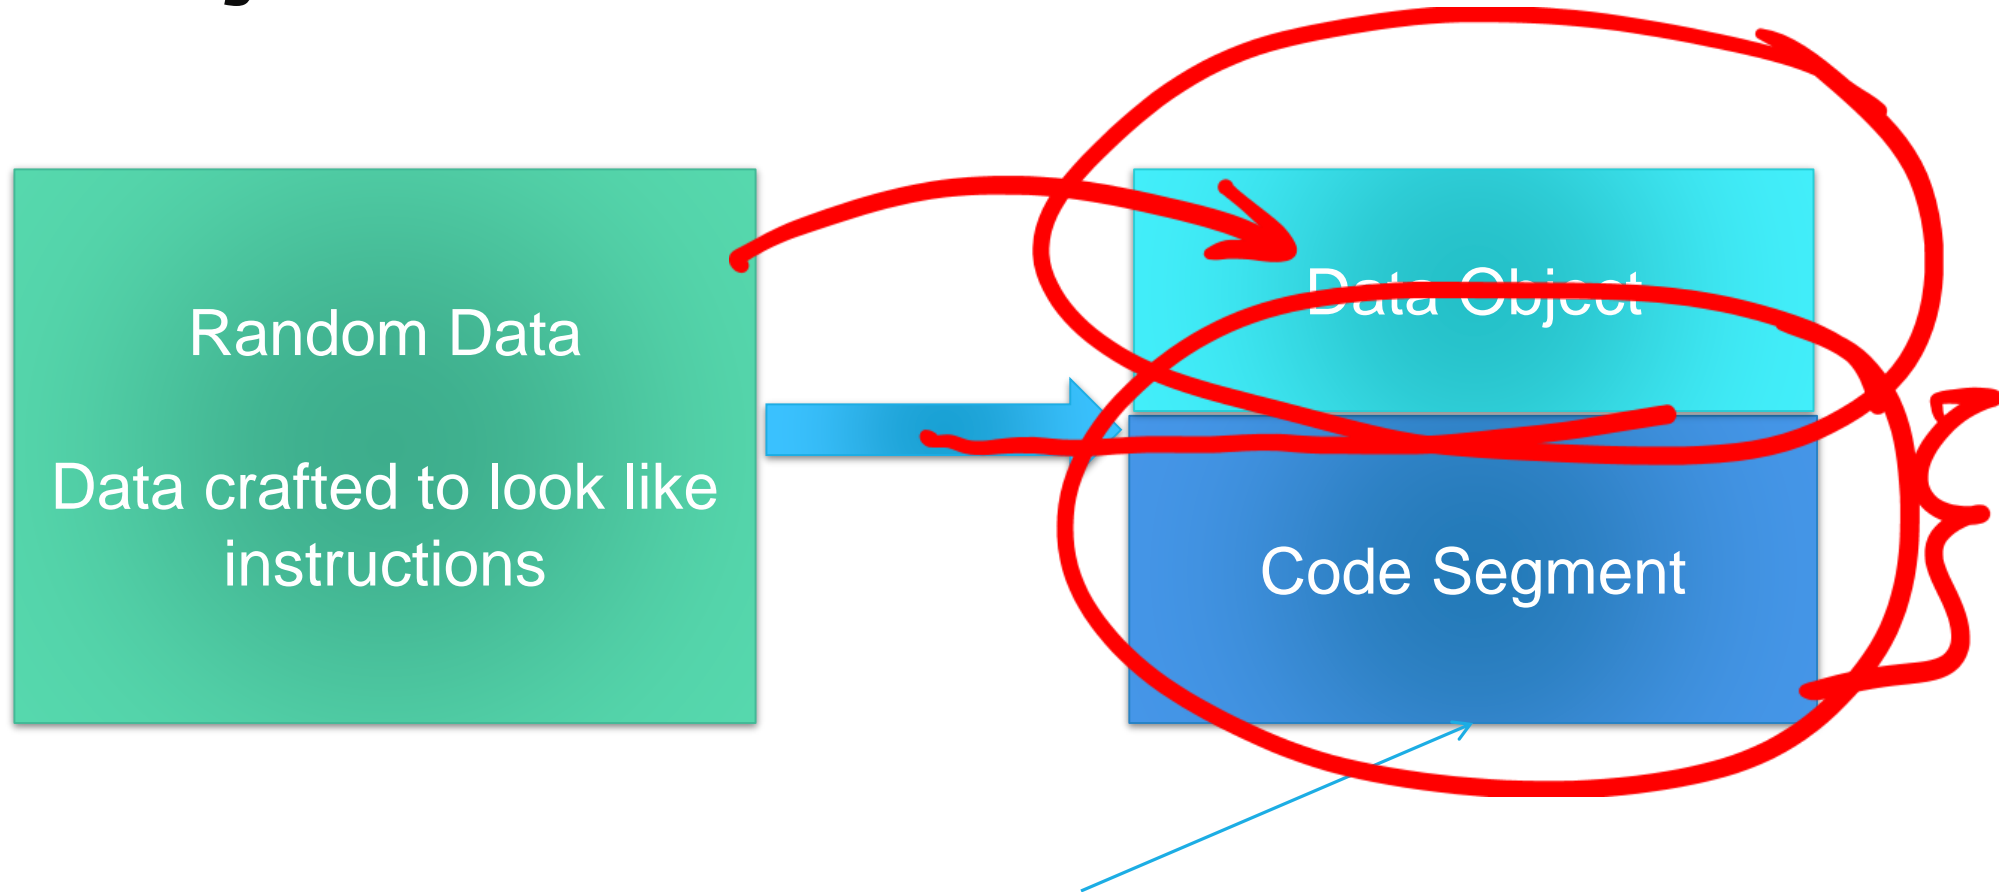
Network protocol request:

Echo 100,000 characters of this string




When you honor the request,  
you disclose your data

# Code injection



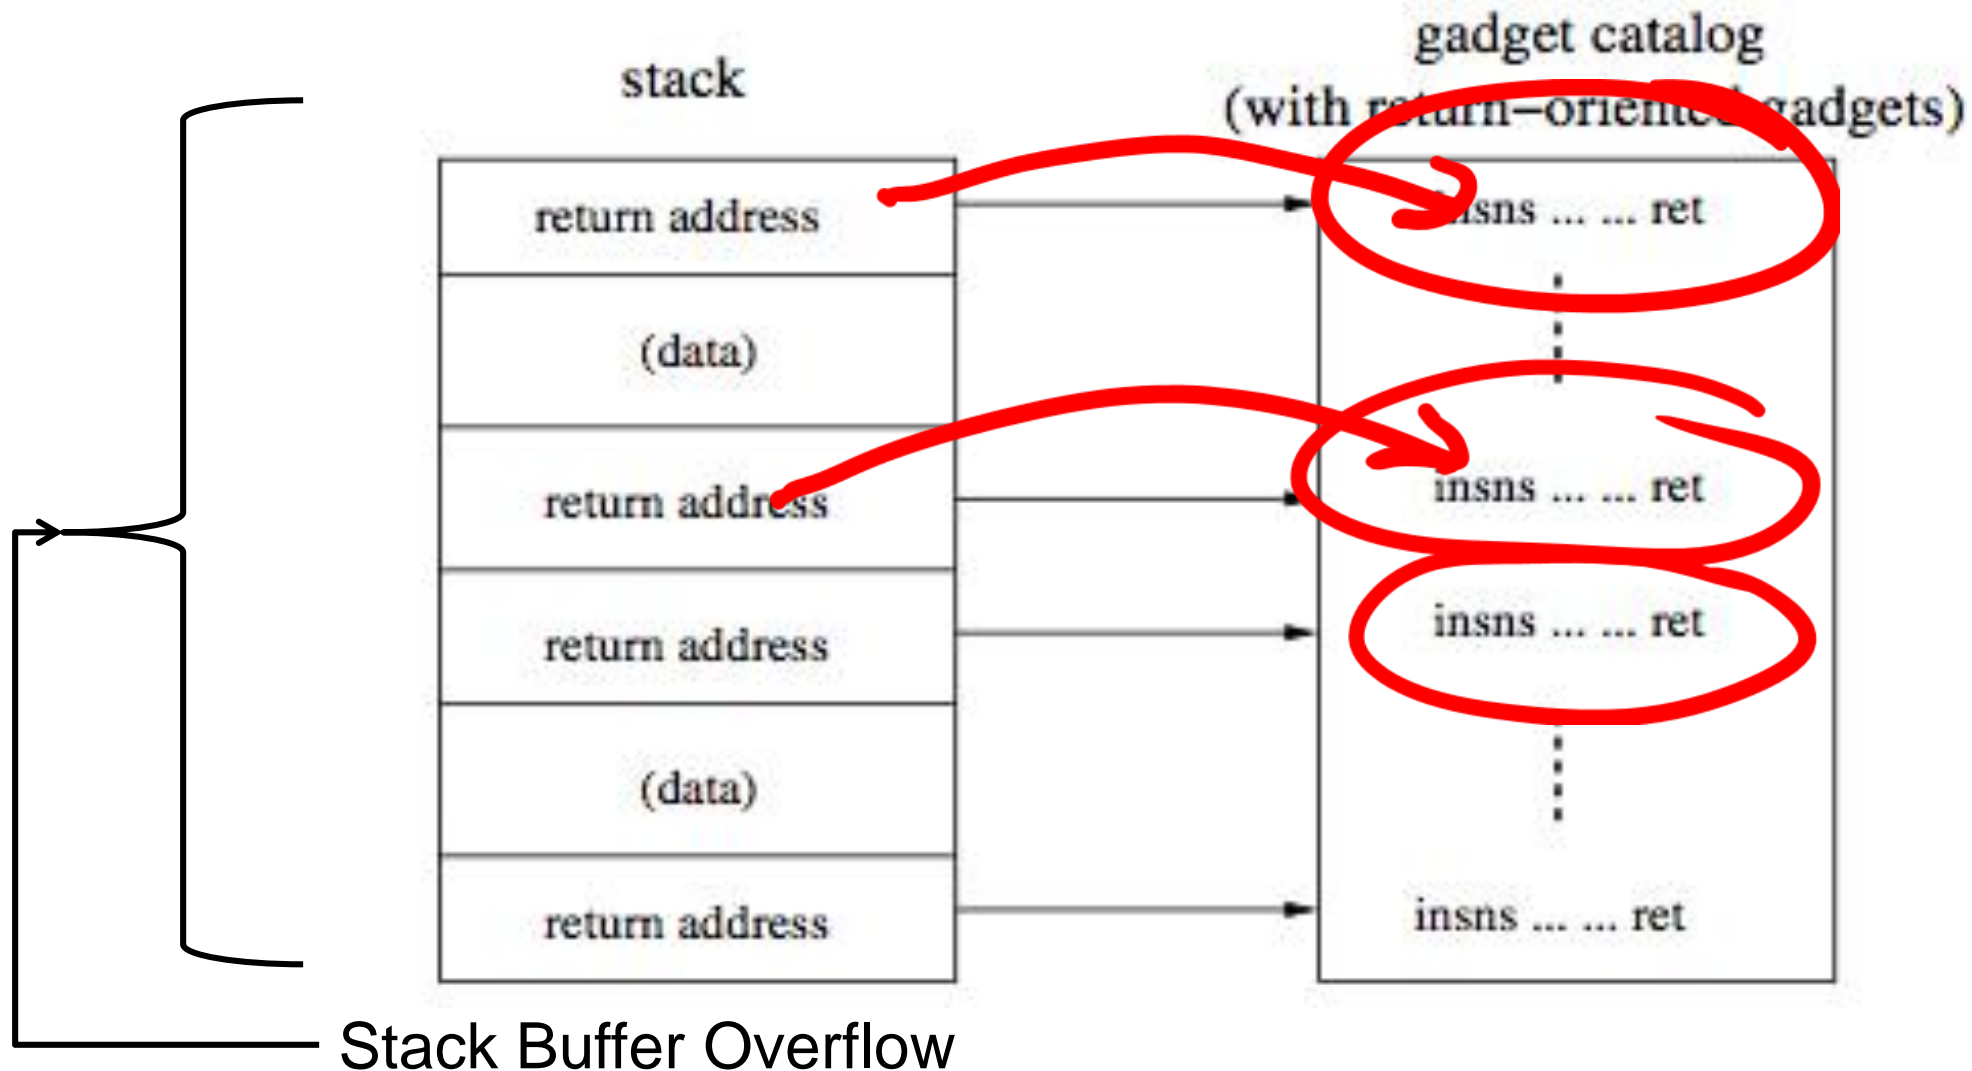
When someone calls this routine, the attacker's instructions get executed

# Pointer over-write



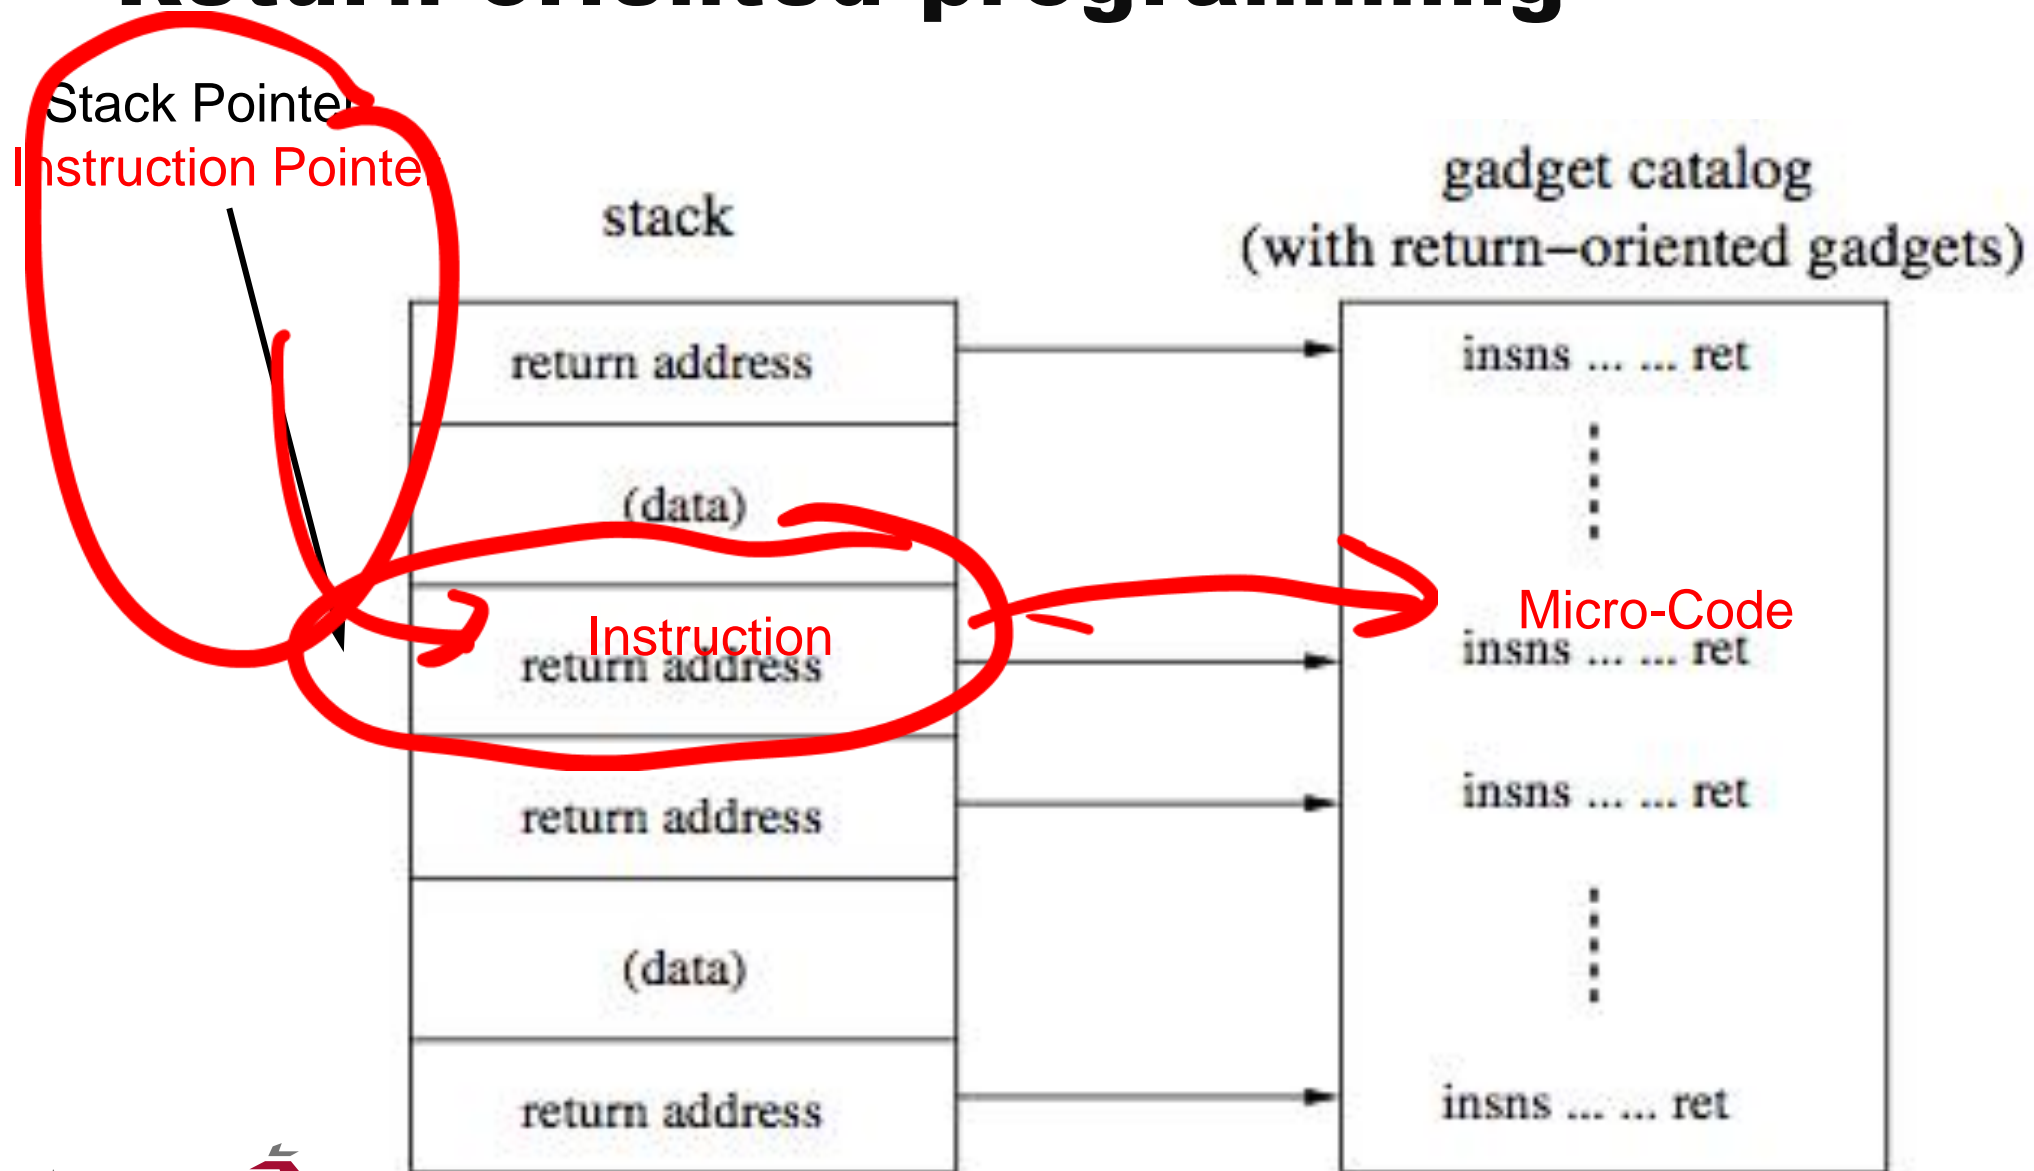
Over-writing the return address  
diverts execution to a place of the  
attacker's choosing

# Return-oriented programming



Stack Buffer Overflow

# Return-oriented programming





# The Underlying Problem:

- Computers don't make important semantic distinctions
- At runtime, there's nothing but "Raw Seething Bits"

10010011100101110

Instruction?

01110010111010010

Data?

01011101001001110

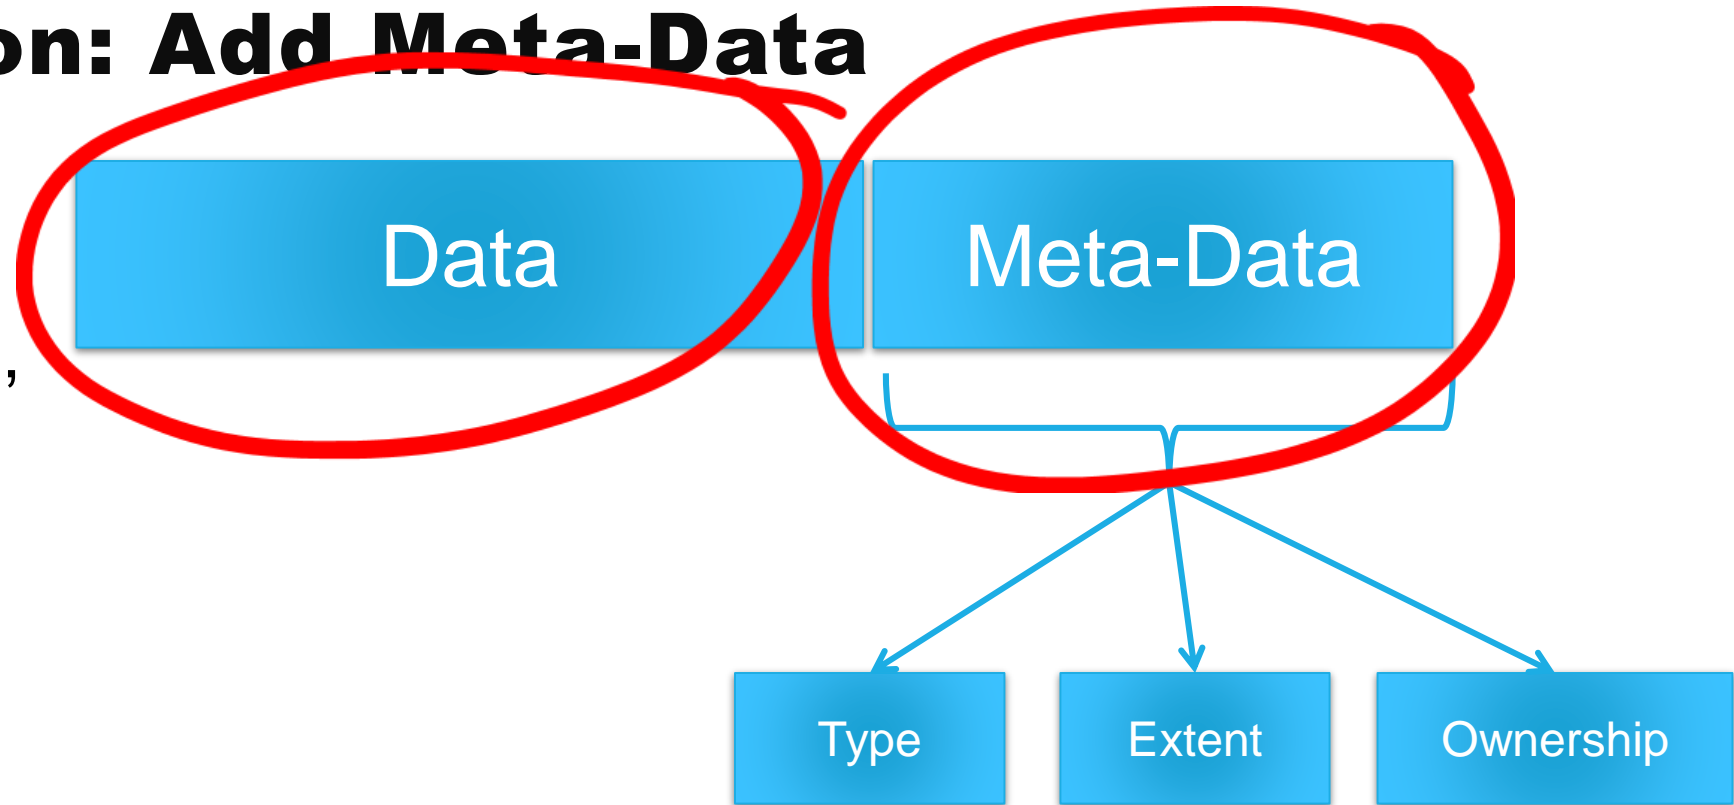
Pointer?

10010011100101110.....

Where's the end?

# The Solution: Add Meta-Data

- Type
  - Instruction
  - Immediate Data,
  - Pointer
- Extent
  - Base
  - Bounds
- Ownership
  - Compartment
  - Access rights



# Type Tagged Architectures



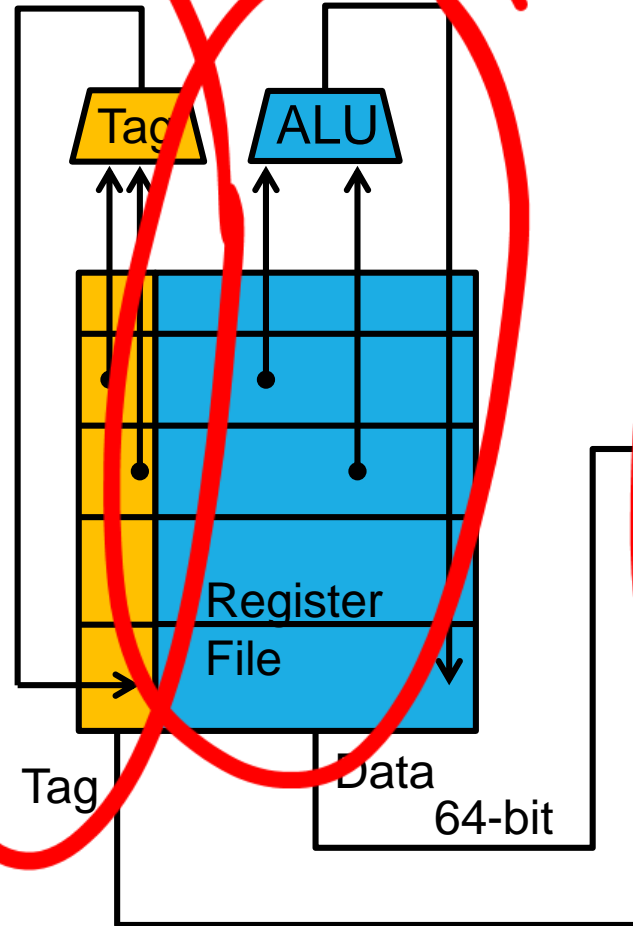
# The Solution: Enforce policies

- Type
  - Only instructions can be executed
  - Only pointers can be followed
  - The argument to a call must be a function pointer
- Extent
  - You can't increment a pointer outside its bounds
  - You can't reference outside a pointer's bounds
- Ownership
  - When you combine data the result maintains compartment
  - Only designated “principles” can access data


# Adding tags for meta-data

Tag processor checks the meta-data tag:

- Does it make sense? Am I trying to add 2 strings?
- Is it allowed? Can I access data of this type?
- What's the tag of the result?

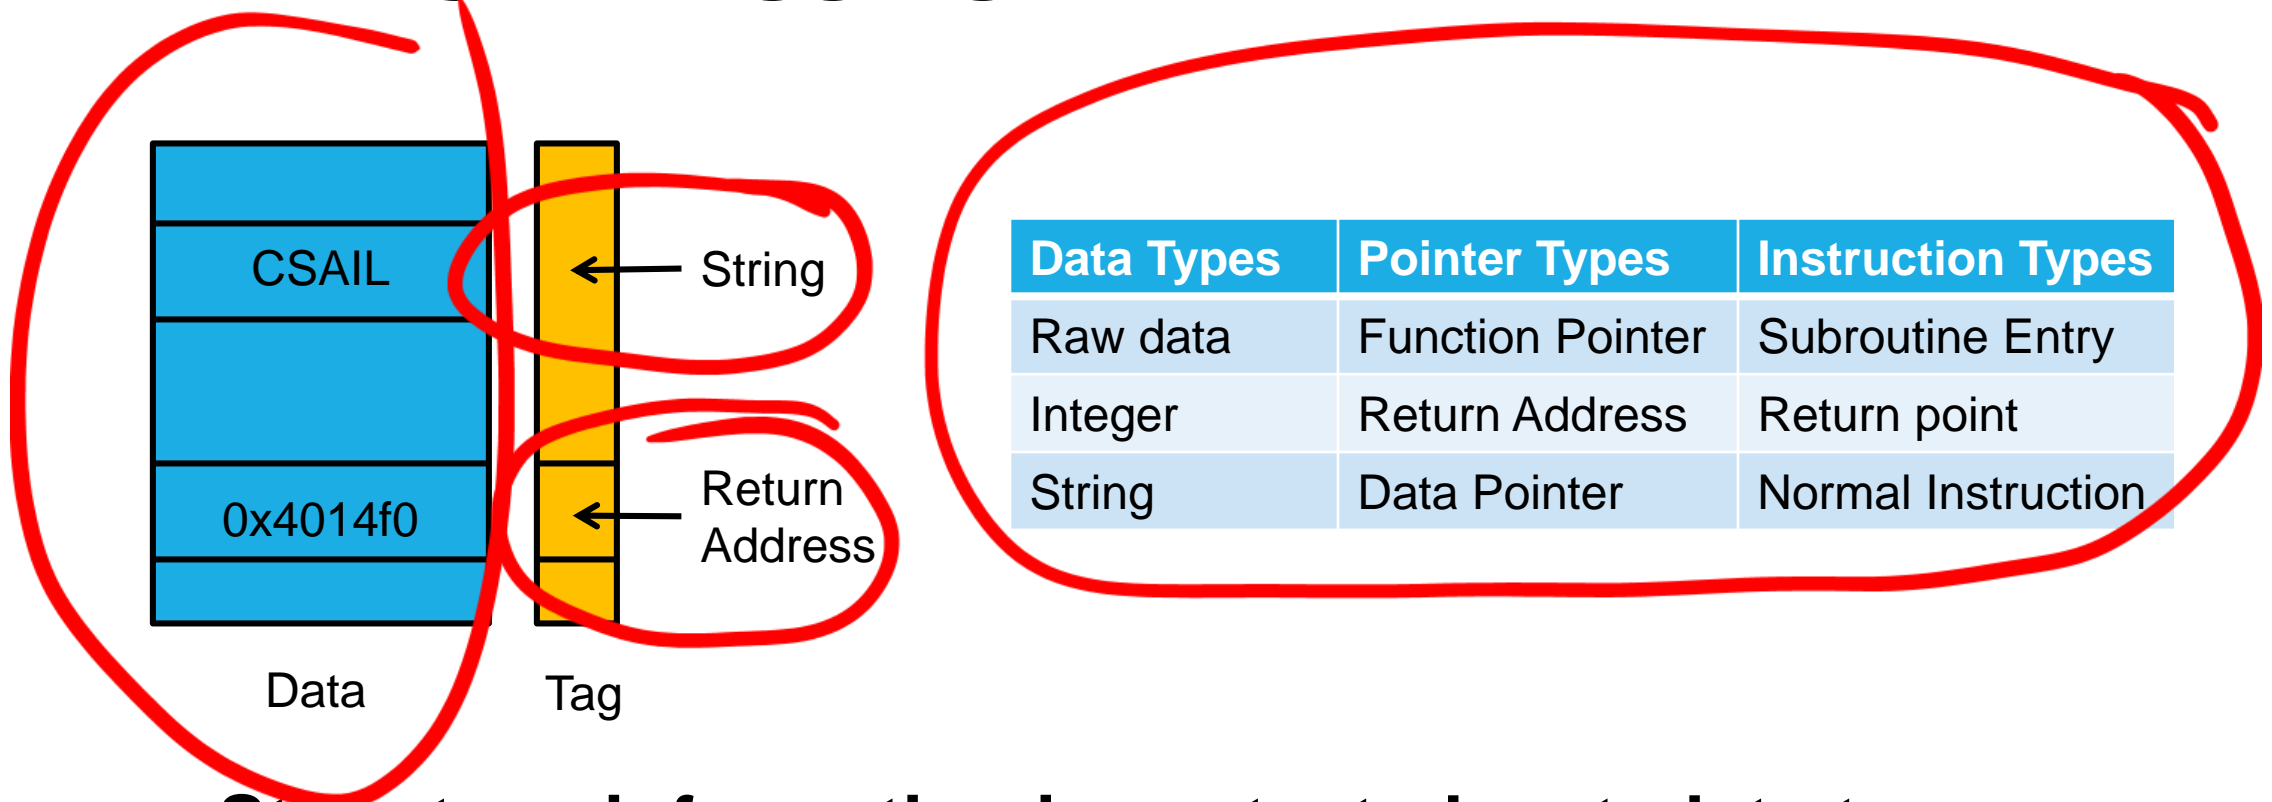


Blue: Conventional architecture  
Yellow: Tag processing additions



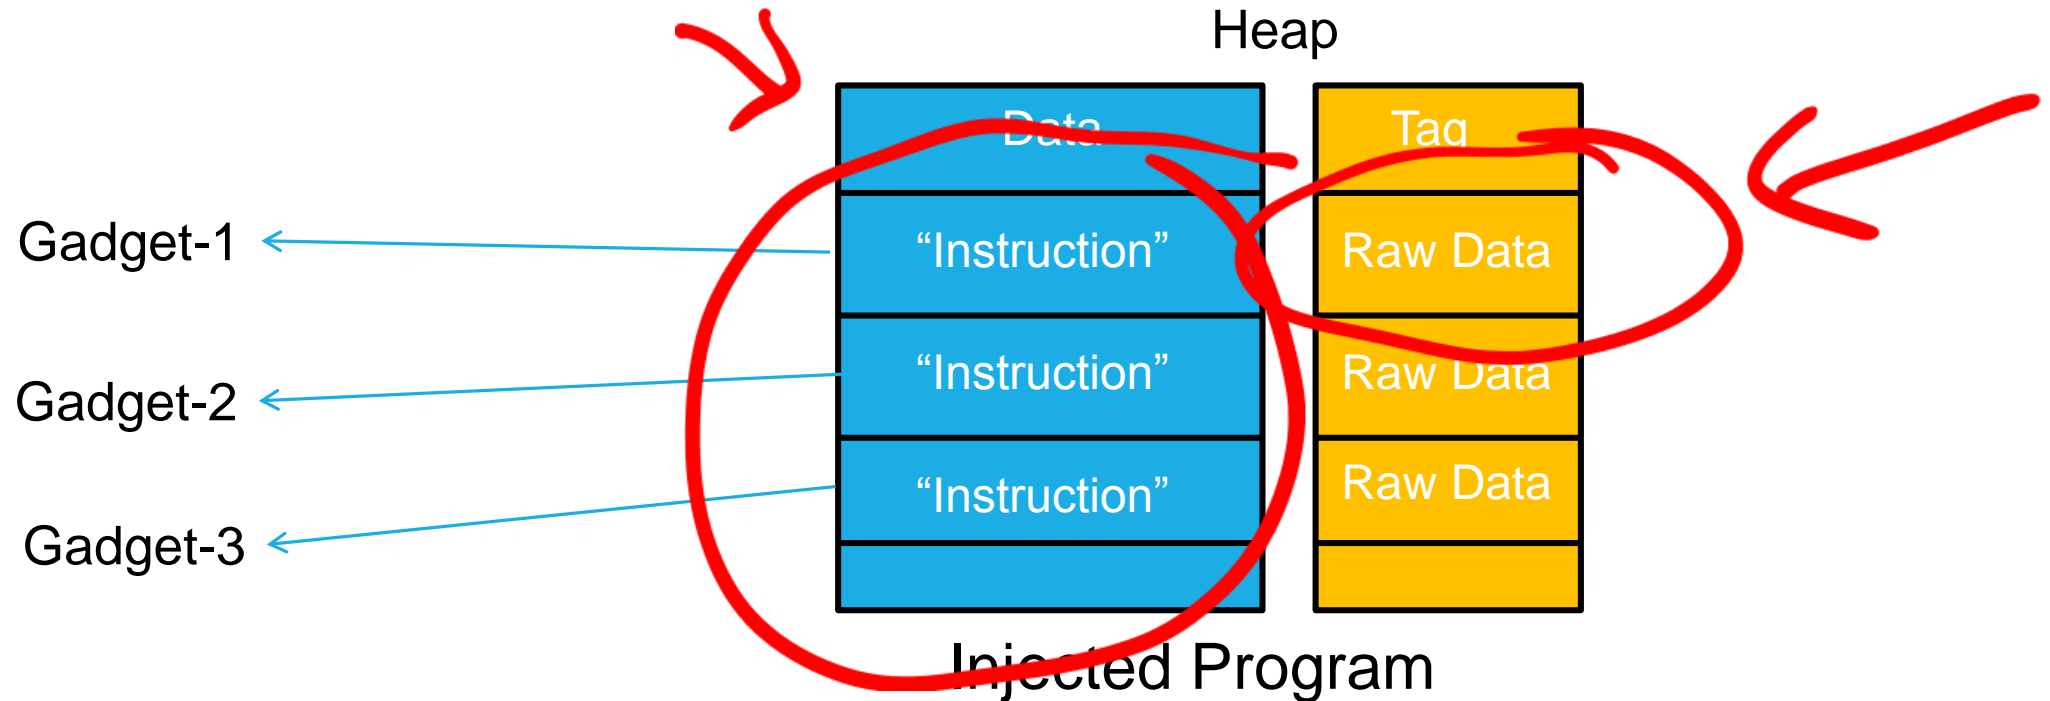


# Data type tagging



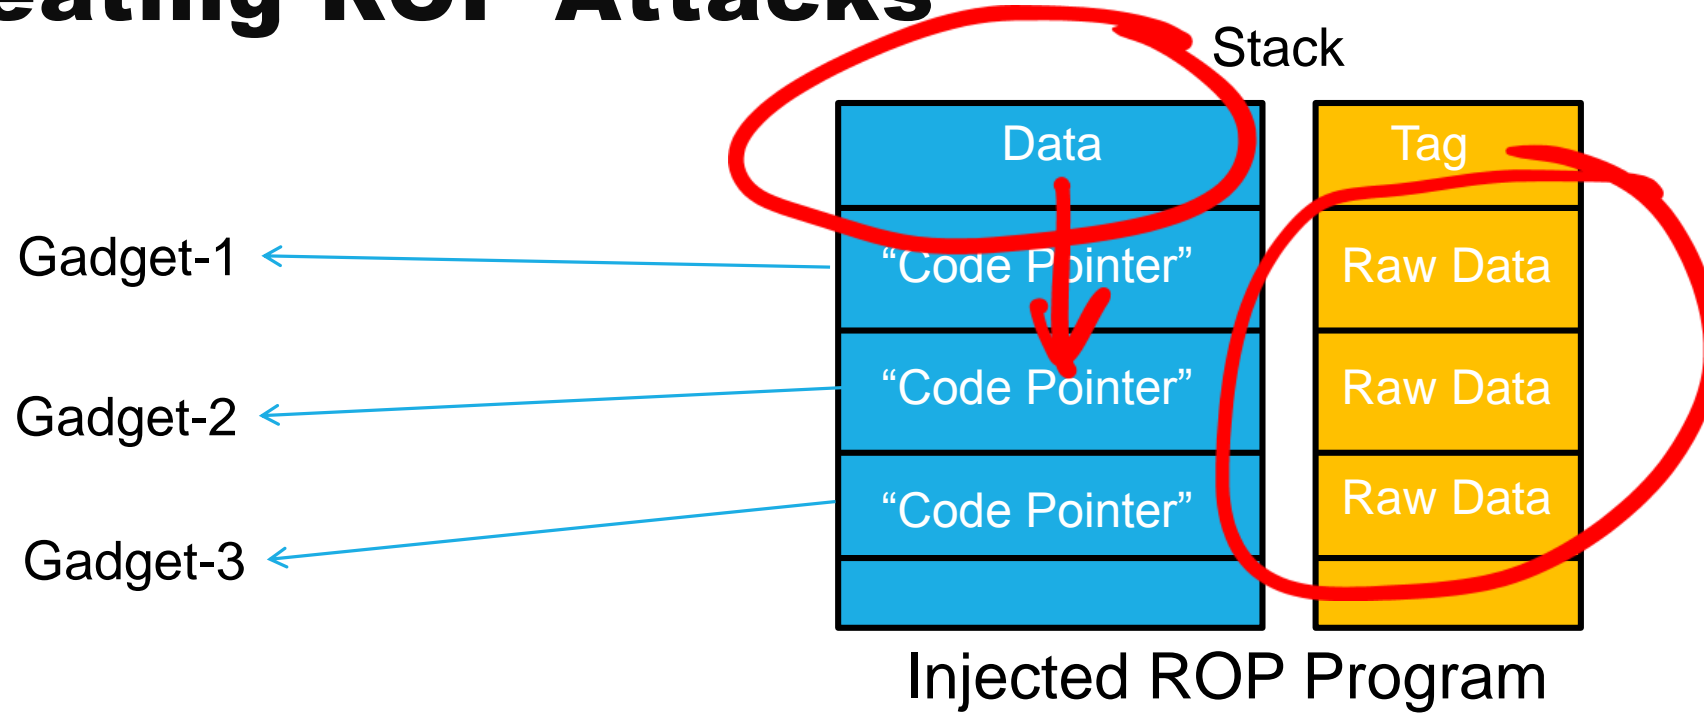
- **Store type information in protected metadata tag**
- **Enforce data-type oriented policies on each cycle**

# Defeating Code Injection Attacks



- Although purported "instructions" can be pushed into the heap, they will have the tag of "Raw Data"
- **Attacker cannot forge instructions and hijack control flow**

# Defeating ROP Attacks



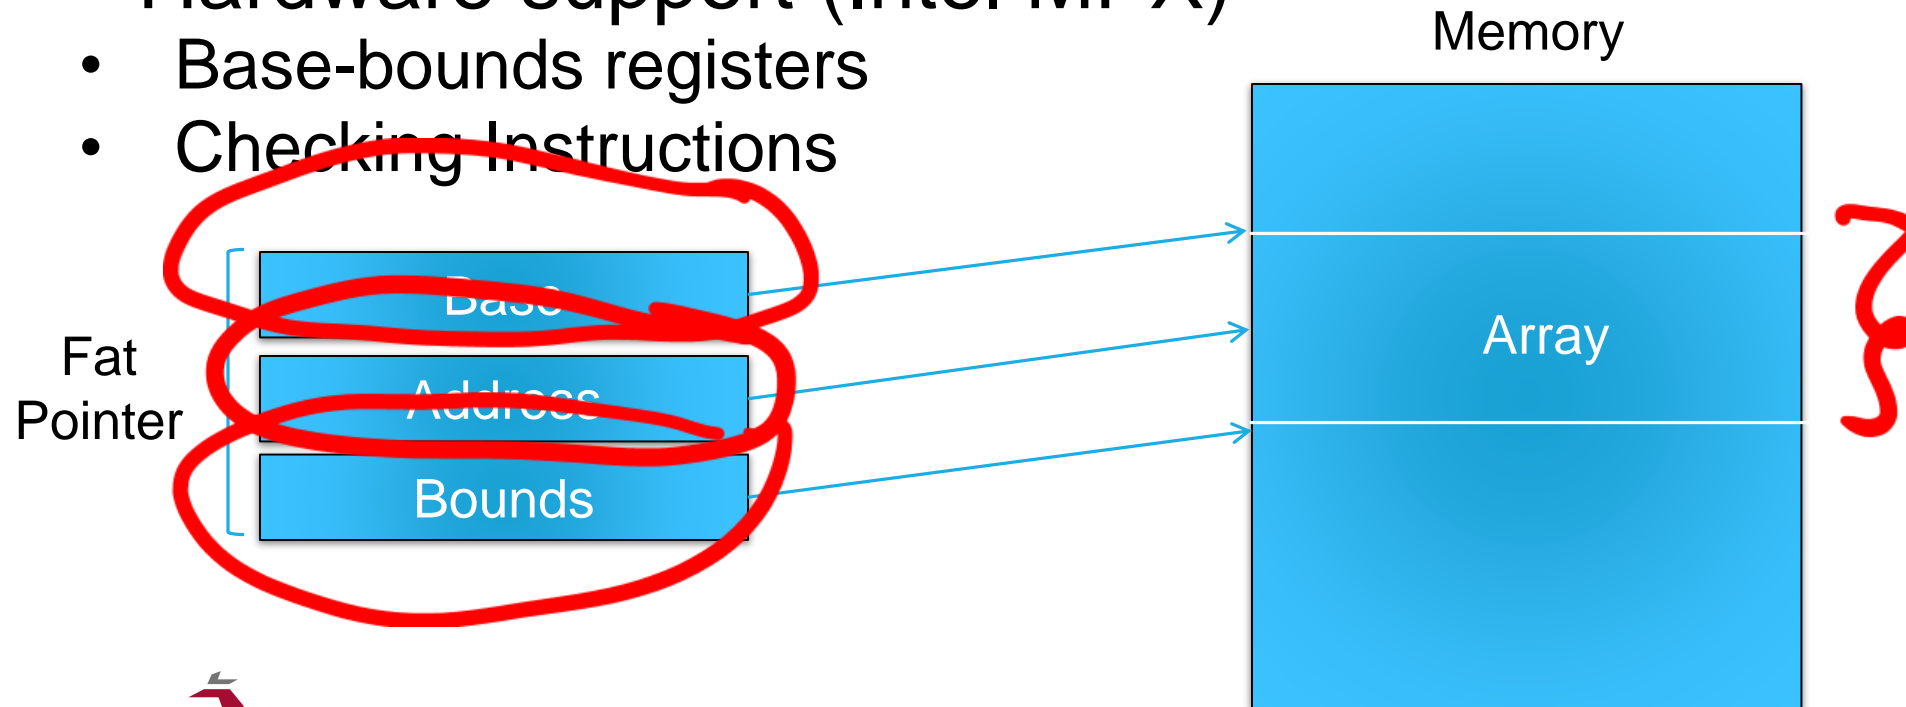
- Although purported “code pointers” can be pushed onto the stack, they will have the tag of “Raw Data”
- **Attacker cannot forge code pointers and hijack control flow**

# Memory Safe Architectures



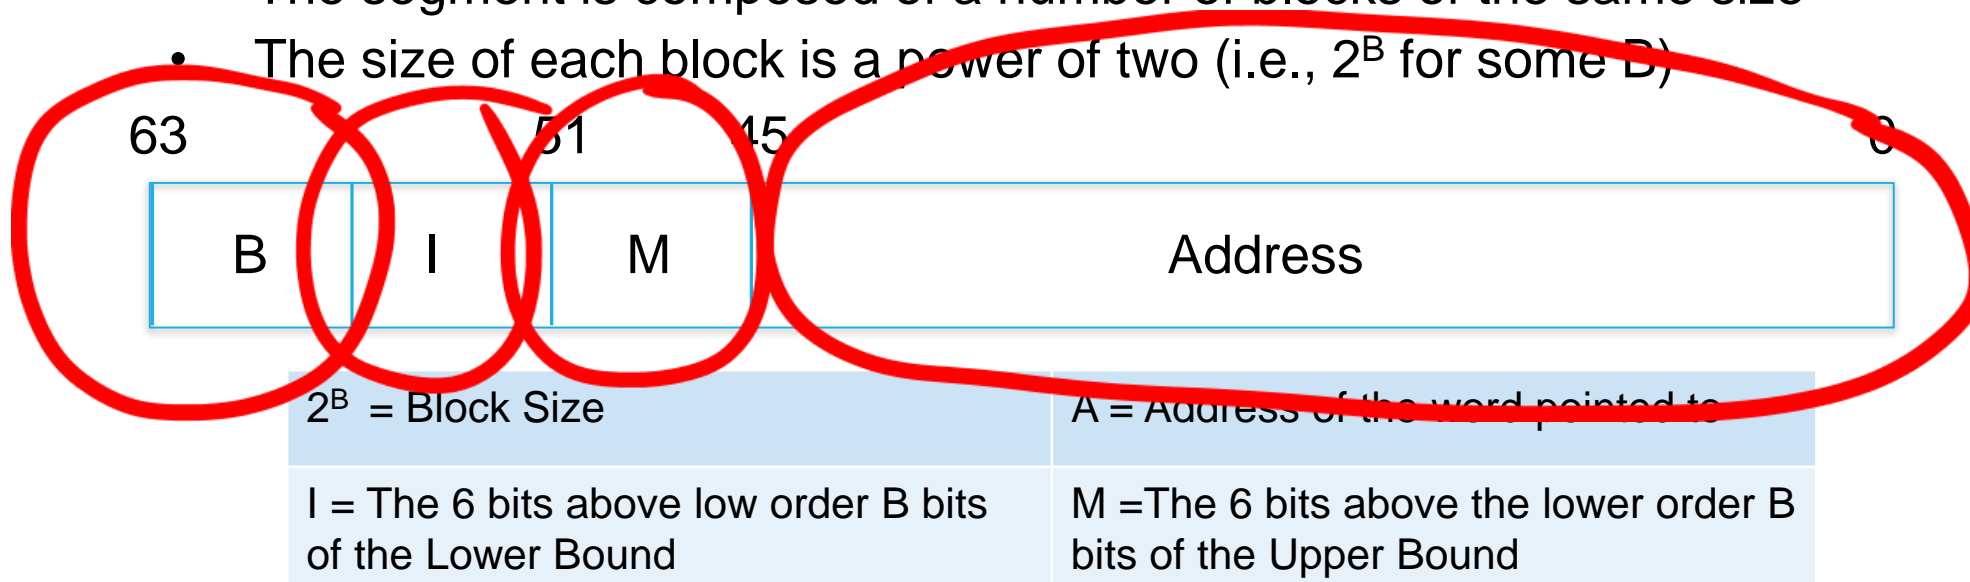
# “Fat Pointers” for memory safety

- Every pointer has 2 other words associated with it for base and bounds
  - 100% runtime overhead when done in software
  - Lots of extra loads, stores and checks
- Hardware support (Intel MPX)
  - Base-bounds registers
  - Checking Instructions



# “Low Fat Pointers”

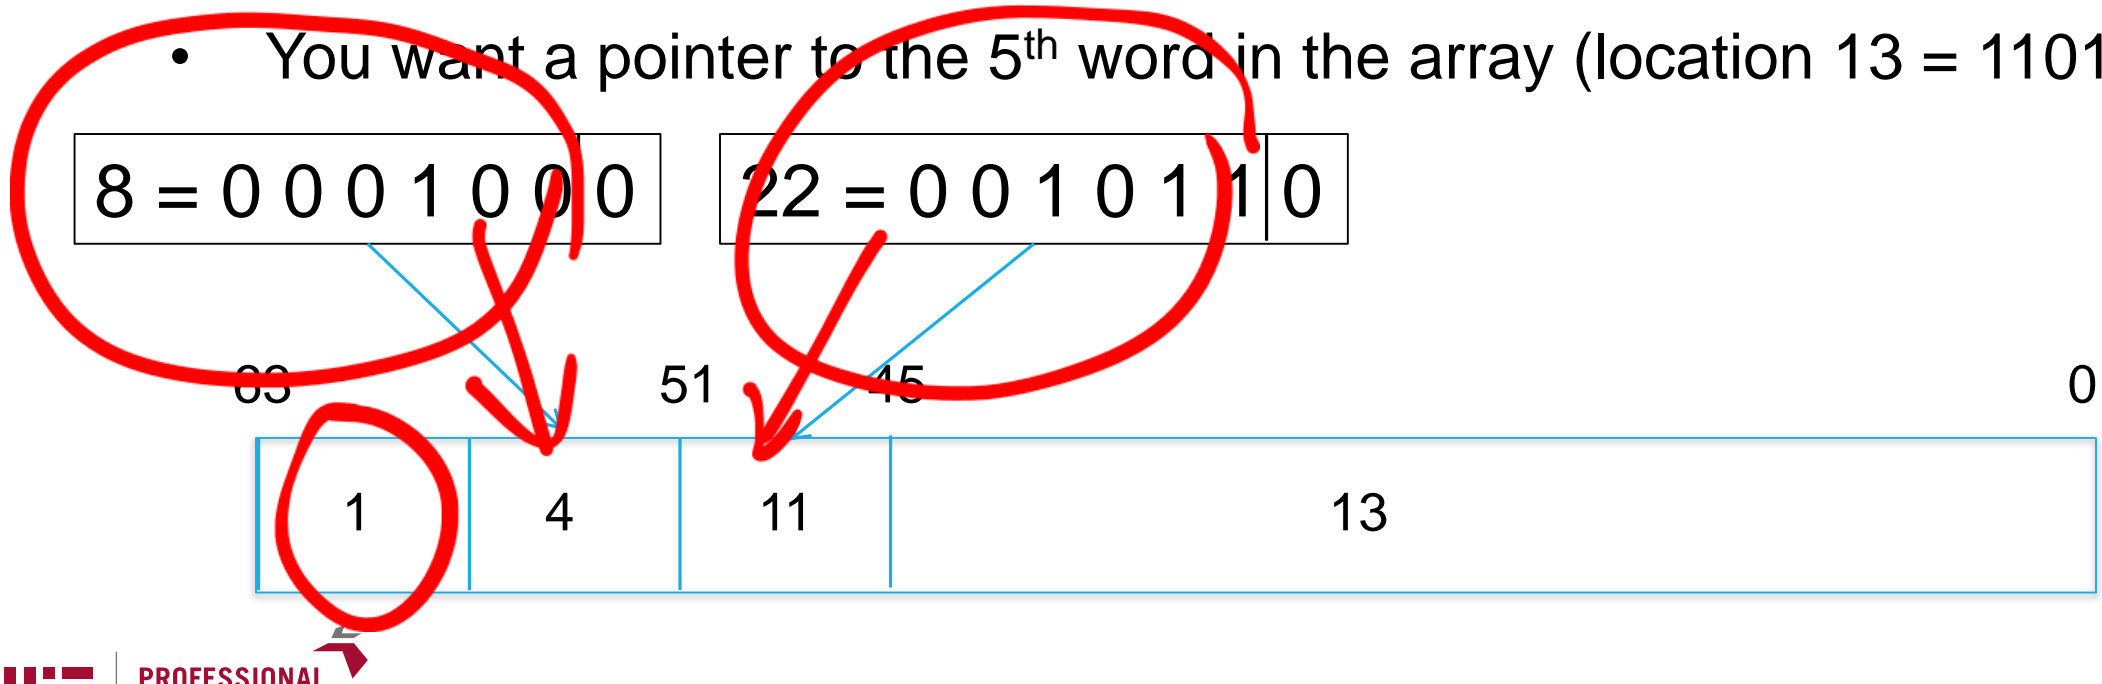
- “SAFE” processor in DARPA CRASH program
- Trades allocation constraints for compactness
  - The segment is composed of a number of blocks of the same size
  - The size of each block is a power of two (i.e.,  $2^B$  for some  $B$ )



**Low-Fat Pointers: Compact Encoding and Efficient Gate-Level Implementation of Fat Pointers ...**  
[20th ACM Conference on Computer and Communications Security, November 6, 2013, Berlin, Germany.](#)

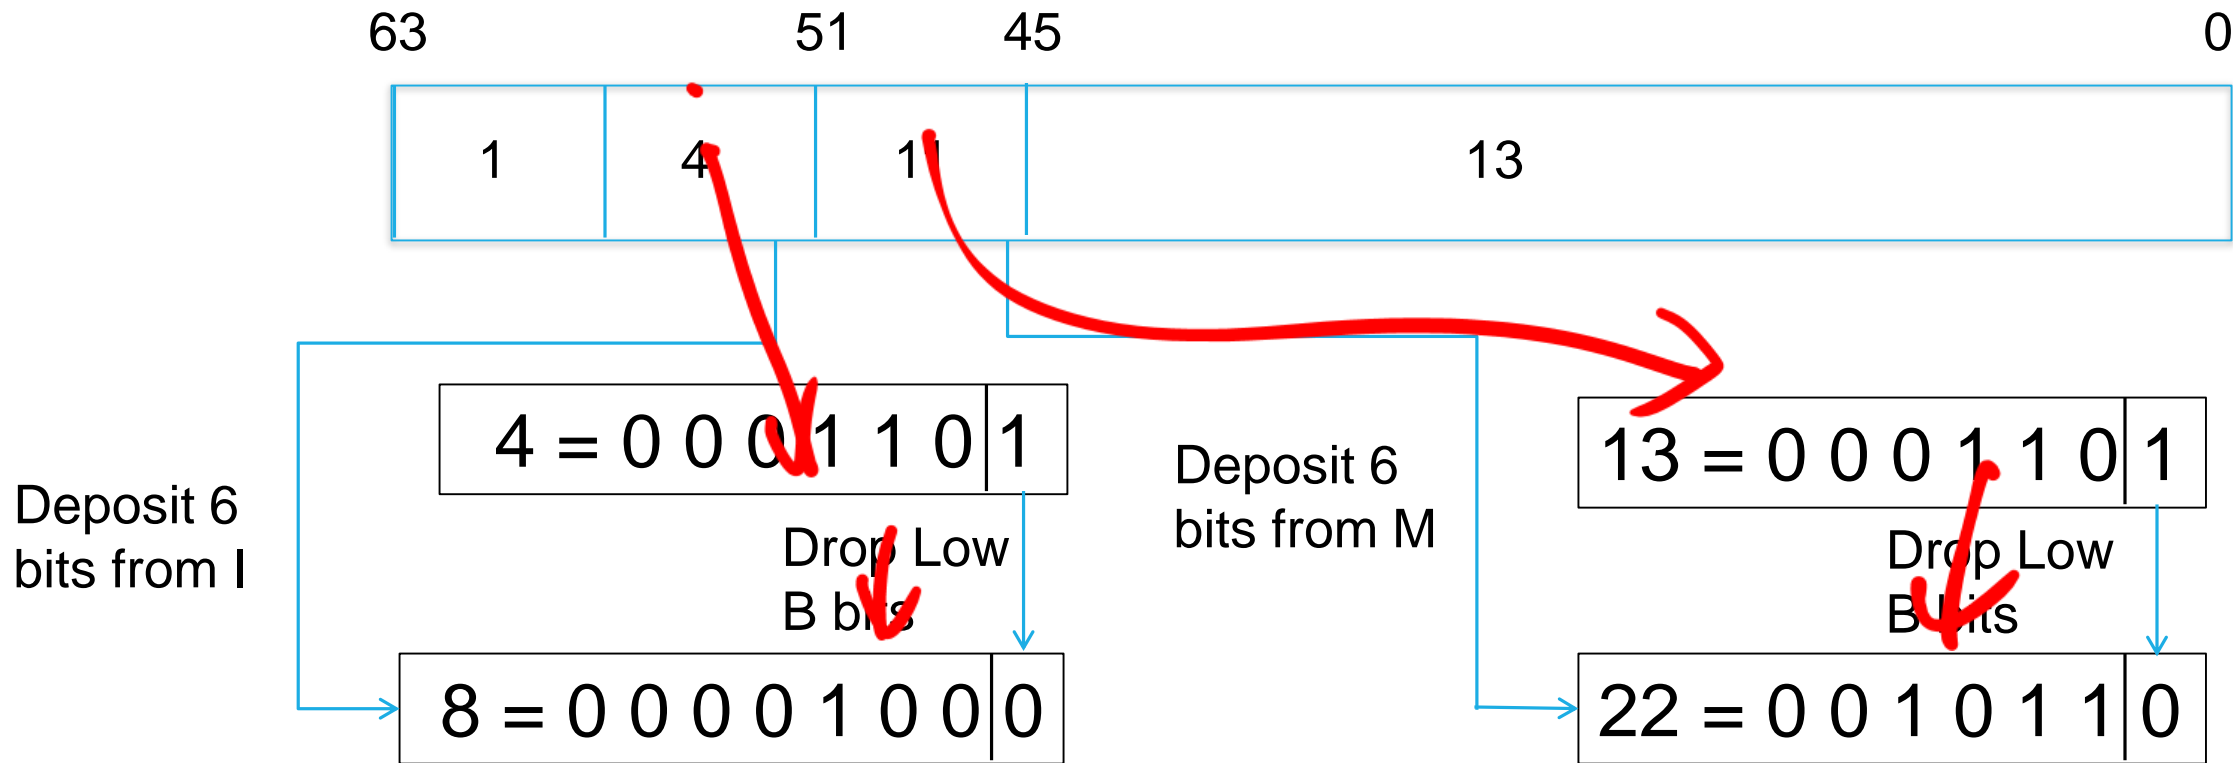
# Low Fat Pointer Example

- Suppose you want an array of 13 words
- The next available space is at location 7
- Base = 8 (aligned to 2)                      01000 in binary
- You need 14 words (7 blocks of 2 words),
- Bound = 22                                      10110 in binary
- You want a pointer to the 5<sup>th</sup> word in the array (location 13 = 1101)





# Recovering base and bound



\* It's actually a little more complicated in some cases, but the hardware can handle it

# “Information Flow”



# Information flow

- Keeping data from getting to where it shouldn't:
- Every word has a “label”
  - A set of “Principals”, the owners of the data
  - The “compartment” the data is classified in
- As data is combined by the ALU, labels are combined by the Tag Management Unit
- The processor is always running on behalf of some specific principal
- Rules specify which principals can perform what operations on the data, given the labels

# Example information flow policy

- Howie has a shared data compartment
  - Any friend of Howie can read this data
  - John is a friend of Howie
- John has a shared data compartment
  - John's friends can read data in that compartment
- John combines data from his shared data compartment with data from Howie's shared data compartment
  - This data can only be read by someone who is both a friend of Howie and a Friend of John

# Information flow policy as rules

If the Principal is Howie

the read-data is in Howie's private compartment

the operation is Read

Then OK with Tag Howie's private compartment

If the Principal is John

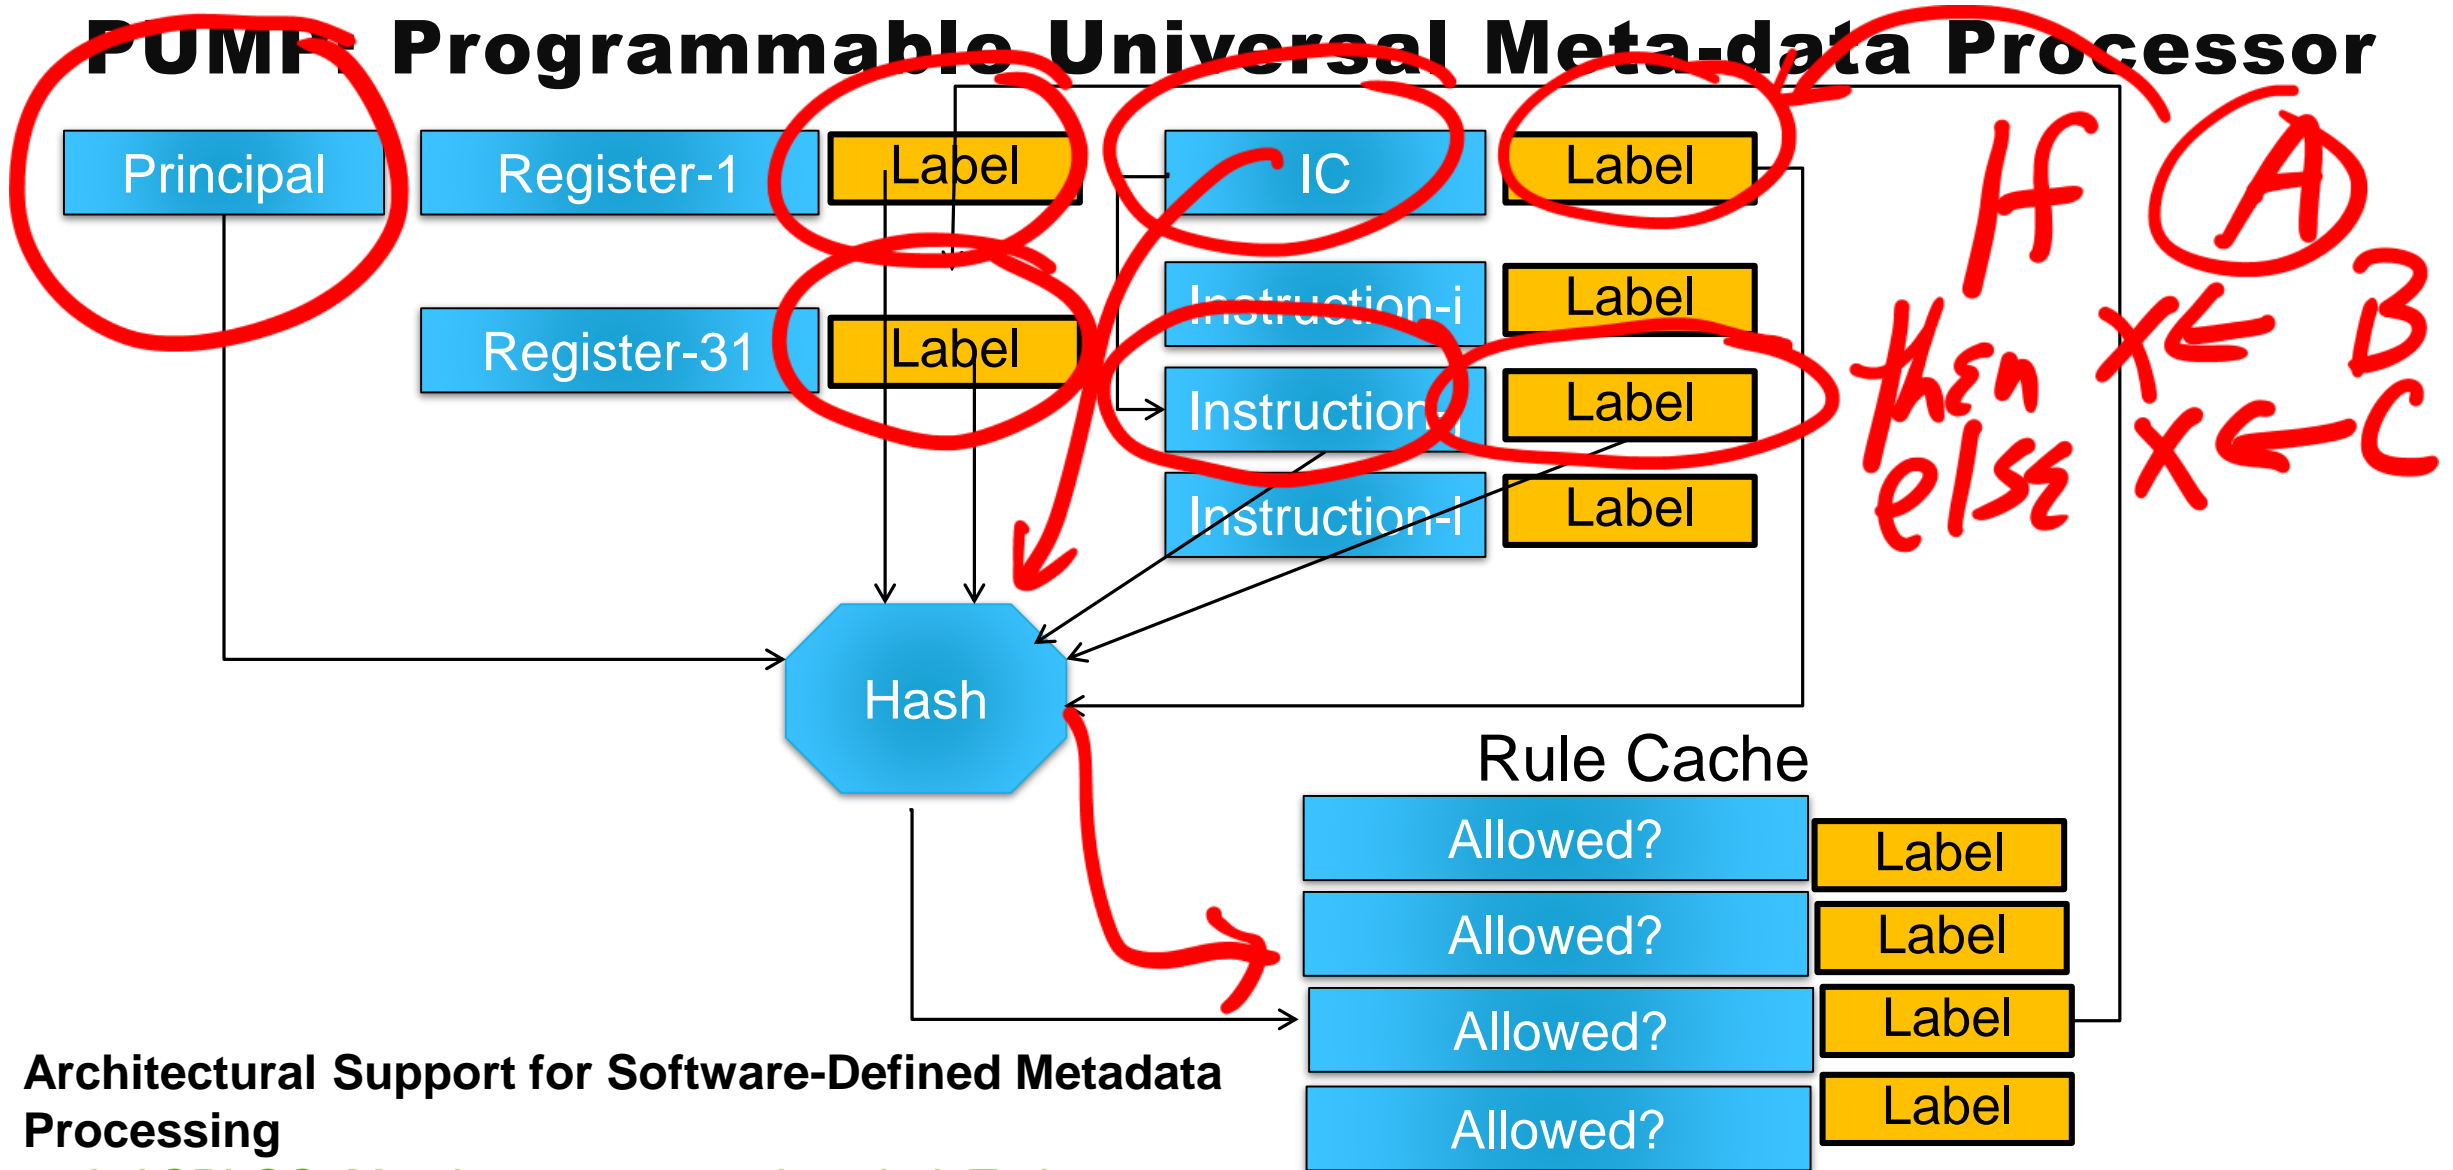
the first operand is in Howie's shared compartment

the second operand is in John's shared compartment

the operation is Add

Then OK with Tag John & Howie shared compartment

# PUMP: Programmable Universal Meta-data Processor



Architectural Support for Software-Defined Metadata Processing

20th ASPLOS, March 14-18, 2015, Istanbul, Turkey.


# Capability architectures





# The problem with current architectures

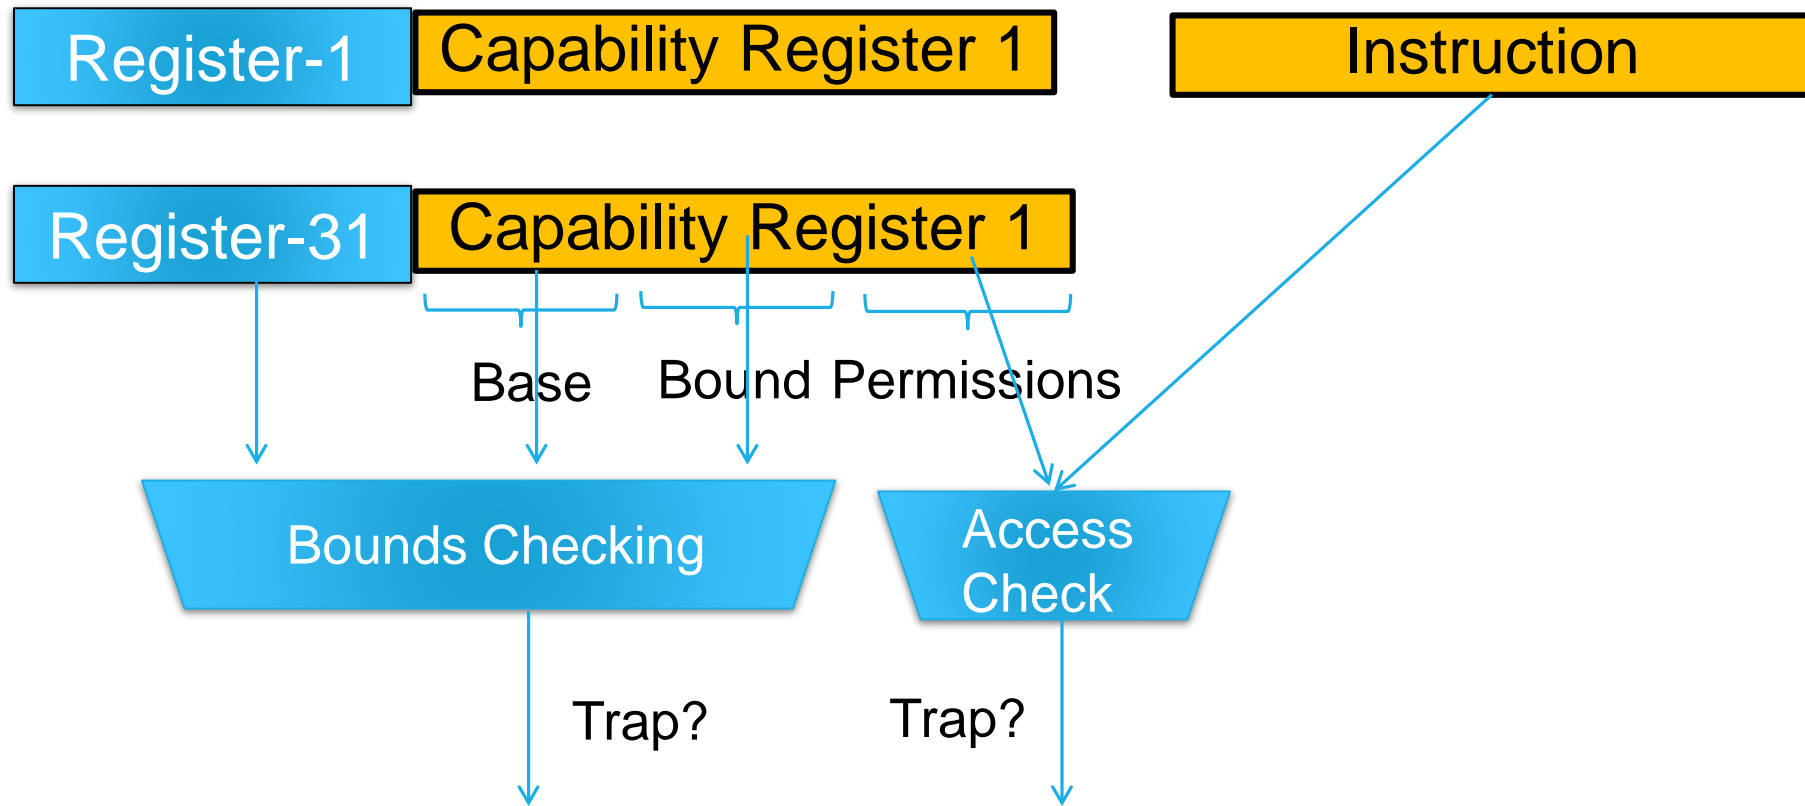
- Address space and protection are identified
  - Context switching is very expensive
- We have, at best, a ring-like system in which the innermost ring (the kernel) is all privileged.
- This violates “least privilege”
  - Why should the scheduler be able to read my personal data?



# Capability Architectures

- A flat memory space
  - Process isolation through virtual memory isn't necessary
- No “raw pointers”
- A Capability is:
  - A pointer
  - Base and bounds
  - Access rights to the thing pointed at
  - Typing information
- You can only access things that you have a capability for

# Capability hardware



# Capability Architectures: Issues


- Passing a capability is equivalent to granting privilege
  - Can lead to “confused deputy” problems
  - Revocation can be difficult
- What you want is to pass a “less capable” capability
  - Tighter bounds
  - More limited access rights
  - More specific type

# A “Zero-Kernel” System

- Merge capability and information flow ideas
- Principals, Tagged Objects, Access Rules
- Gate calls:
  - Invoke a new routine with a changed principal
  - Principals are not strictly hierarchical
- The Kernel is broken up into many pieces
  - Each with its own principal and compartment
  - Controlled sharing within the kernel
  - Least privilege with respect to user data

# Zero Kernel System: Building Blocks

Gate:  
A Procedure Bound  
with a New Principal



Compartment:  
A collection of Data with  
common access rights



Access Rule:  
Which principal can do what  
operation on what data


P1	C1	C2	Add	OK	$C1 \wedge C2$
P2	C1	C2	Add	Trap	
P1	C3	C4	Add	Trap	
P2	C3	C4	Add	OK	C3

# Zero Kernel System: Least Privilege

I'm the Principal  
accessing data in my  
Compartment



I make a Gate Call:  
The new Principal is the  
Scheduler acting for me



Only the necessary  
data is copied into a  
shared compartment

This Principal can only  
access data in the  
shared compartment.

Scheduler  
Acting  
For Howie




This Principal can make  
Gate Calls to the core  
of the Scheduler, but it  
can't pass my data to it.



The Scheduler Principal  
can only access data in  
its own compartment.

Core  
Scheduler





# Summary

- Novel hardware can work with the OS and language runtime to enforce the intended semantics
- Memory & Type Safety & Information Flow
- Complete Mediation
- Least Privilege

# THANK YOU

Howard Shrobe

Principal Research Scientist, Director Security@CSAIL

