

CYBERSECURITY

Case Studies



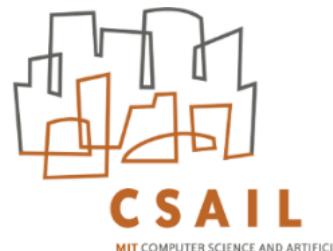
Disk Encryption: BitLocker

Nickolai Zeldovich

Associate Professor

Computer Science and Artificial Intelligence Laboratory (CSAIL)

Massachusetts Institute of Technology



Segment 1: Introduction

- Goals and threat model
- Approaches
- Challenges

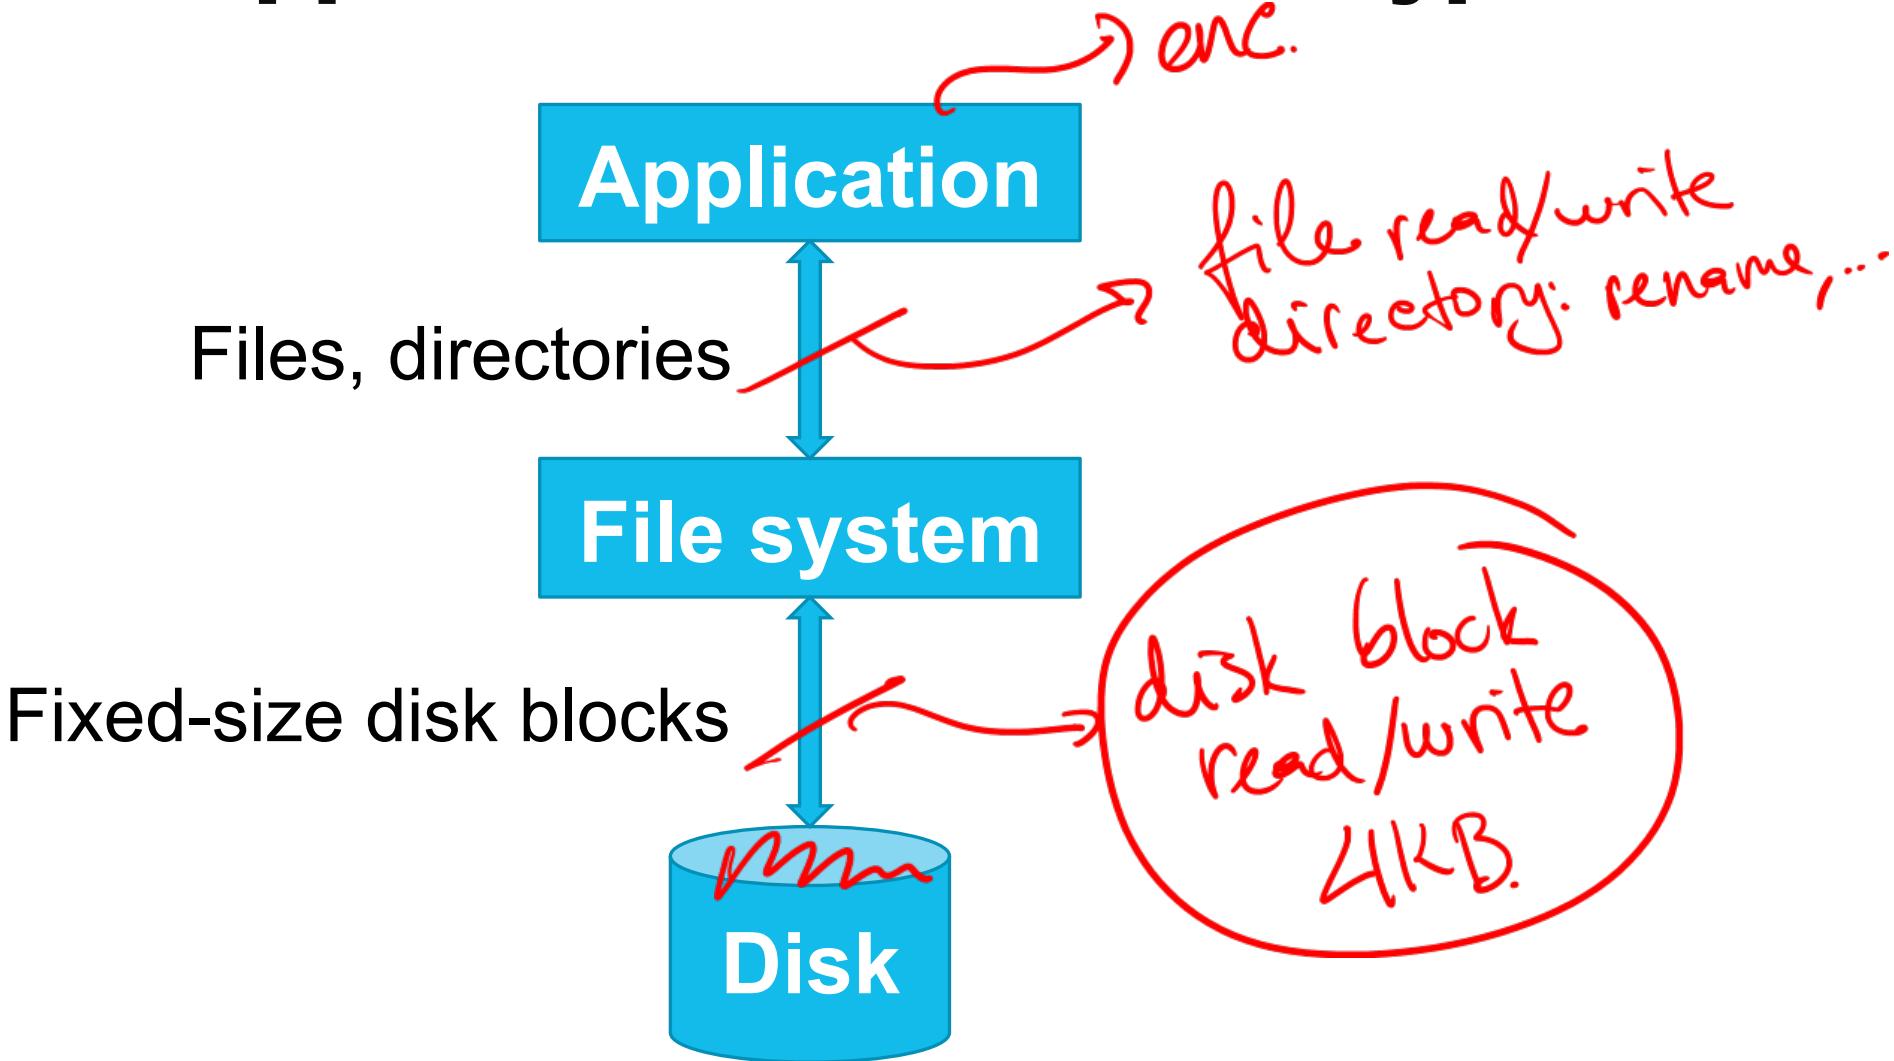
Disk encryption goals

- Scenario: adversary steals a laptop
- Primary goal: *secrecy*
 - Ensure confidential data not leaked to adversary
- Related goal: *integrity*
 - Adversary cannot tamper with data on disk
 - Will be indirectly important later on
- Others: *freshness, deniability, ...*

Threat model

- What can the adversary do?
 - Boot up the laptop
 - Boot from a CD or USB drive, change password
 - Remove disk, plug into another computer
 - Make copies of the disk, so as to try different things
- What is the adversary *unable* to do?
 - Break cryptography, use lots of compute power
 - Break tamper-resistant hardware
 - Know the user's exact password, PIN, ...

Approaches to disk encryption



Challenges in disk encryption

- Where does the key come from?
 - Password, USB device, trusted hardware (TPM), ...
- How is the data encrypted on disk?
 - Using block ciphers, randomness, authentication, ...

Segment 2: Keys from passwords

- The trouble with passwords: entropy
- Key-derivation functions
- Salting passwords

Turning a password into a key

- Password: variable-length string
- Key: fixed number of bits
- Naïve approach: use a hash function

$$\text{key} = \text{SHA1}(\text{password})$$

Problem: passwords have low entropy

- Users choose predictable passwords
 - Top 5,000 passwords account for 20% of users (out of 32M)
- Our threat model: adversary can guess
 - Guess requires hashing password, seeing if decryption works
 - Decryption must be designed to be fast
 - Hash functions also designed to be fast
 - Adversary can make over 1M guesses/second

Idea 1: expensive key derivation

- Plan: make the hash function expensive

$$\text{key} = \text{SHA1}(\text{SHA1}(\dots\text{SHA1}(\text{password})))$$

- Slows down each guess by adversary
- Not a big problem for user: just one time

Problem: pre-computed tables

- Adversary can pre-compute table of hashes
 - Time-consuming, but worth it to break passwords quickly!

Password	Hash
aaa	7e240de7...
acorn	f5d193d9...
alice	522b276a...
...	...

Idea 2: salting the passwords

- Introduce random “salt” into key generation

$$\text{key} = H(\text{password}, \text{salt})$$

- To compute key, need password *and* salt
 - Salt stored in plaintext on disk, but crucially, has high entropy
 - Adversary can get the salt value after stealing the laptop
 - However, now cannot pre-compute table of possible hashes!

Real-world key derivation functions

- PBKDF2
- bcrypt
- scrypt

Limitations of using passwords

- Susceptible to guessing, even with salting
 - Ideally, want to allow just a handful of guess attempts!
- Usability challenges
 - Disk encryption key needed as soon as computer boots up
 - User must enter password into BIOS, before OS running
 - OS may require separate authentication: password, fingerprint, ...
- Nonetheless, can be useful for recovery

Segment 3: Password alternatives

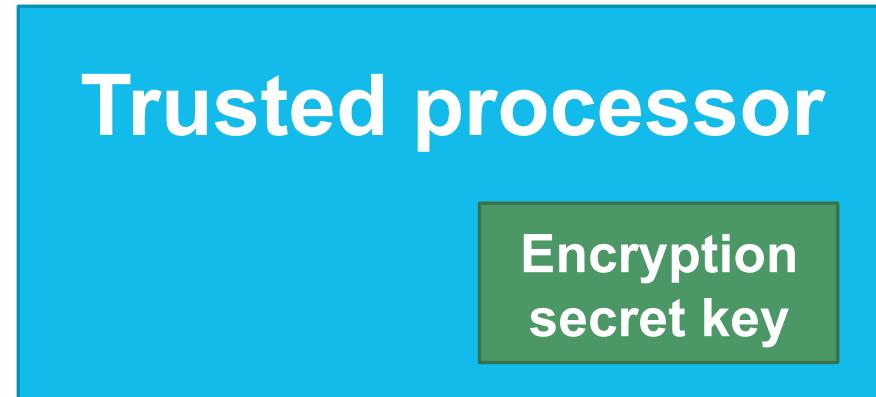
- Removable devices
- Trusted hardware
- Signed boot

Removable devices

- Idea: store key on removable device
 - USB drive, or other devices (e.g., Yubikey)
- Good: strong, hard-to-guess key
 - Can generate high-entropy keys, no need to rely on user
- Bad: user has to carry separate device
- Bad: device could well be in stolen laptop
 - Device has no idea who is using it (adversary or user)

Trusted hardware

- Idea: store key in a tamper-proof processor
 - E.g., system CPU or a nearby chip
 - Ensure that physically “capping” the chip is costly
- What software is allowed to access the key?



Signed boot

- System requires all software to be signed
 - Vendor's public key hard-wired into the trusted processor
 - Vendor signs only approved software releases
 - E.g., used by some game consoles, cell phones, etc



Disk encryption with signed boot

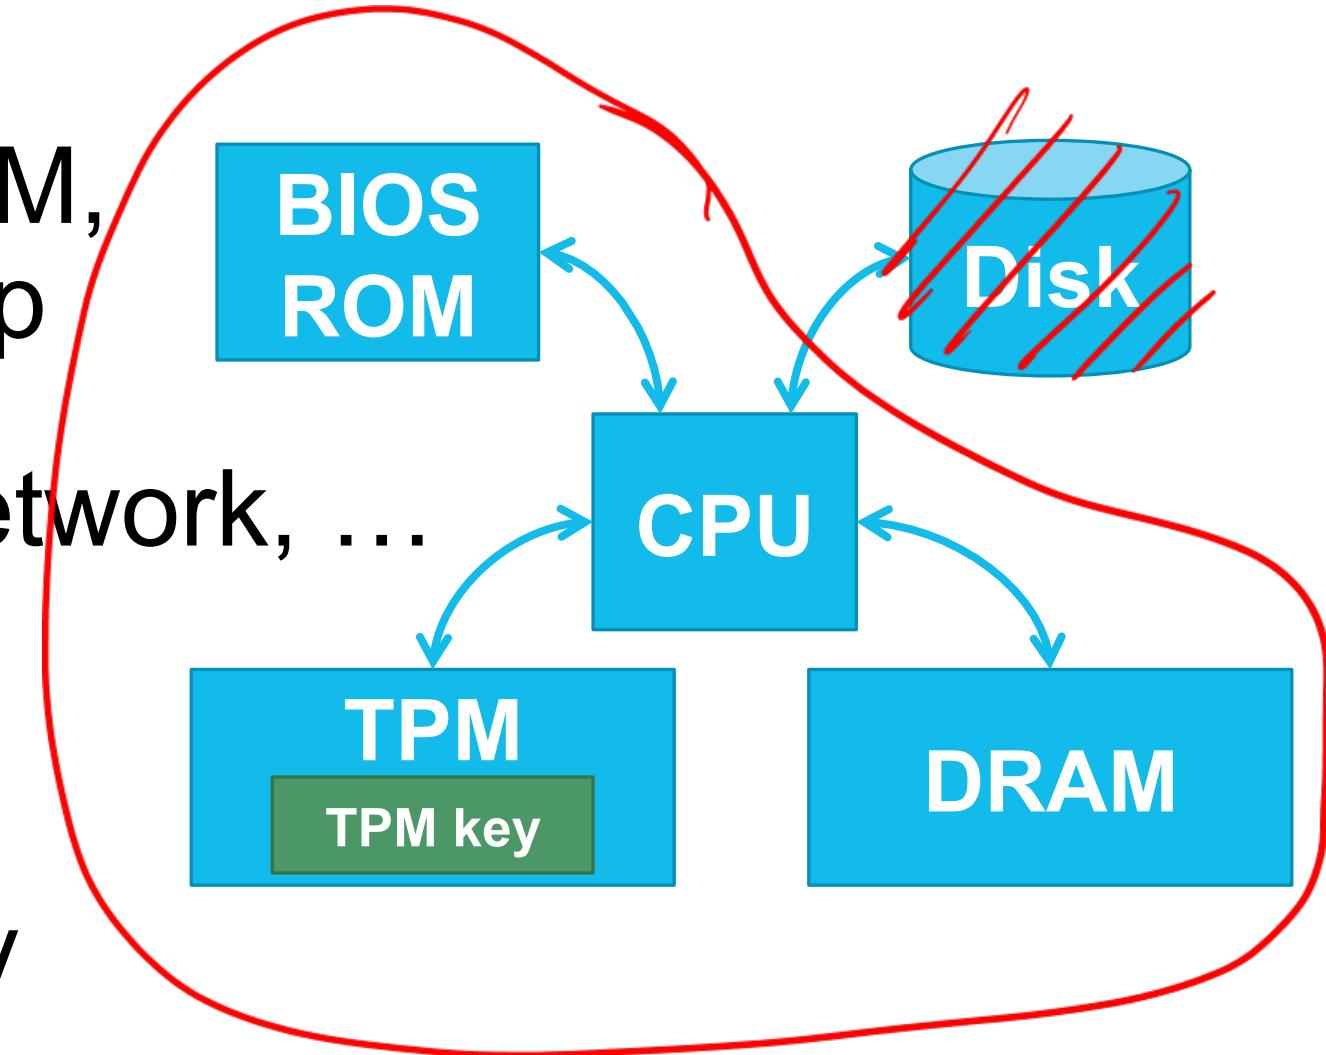
- Use key in trusted processor for disk enc.
- Signed OS implements user authentication
 - Passwords, fingerprints, face recognition, ...
 - No need to generate key from authentication mechanism
 - Easy to enforce strong limits on guessing in software
- Downside: too restrictive in many cases
 - All OS code must be signed by a single vendor

Segment 4: Trusted Platform Module (TPM)

- System model
- TPM interface
- Measured boot

TPM system model

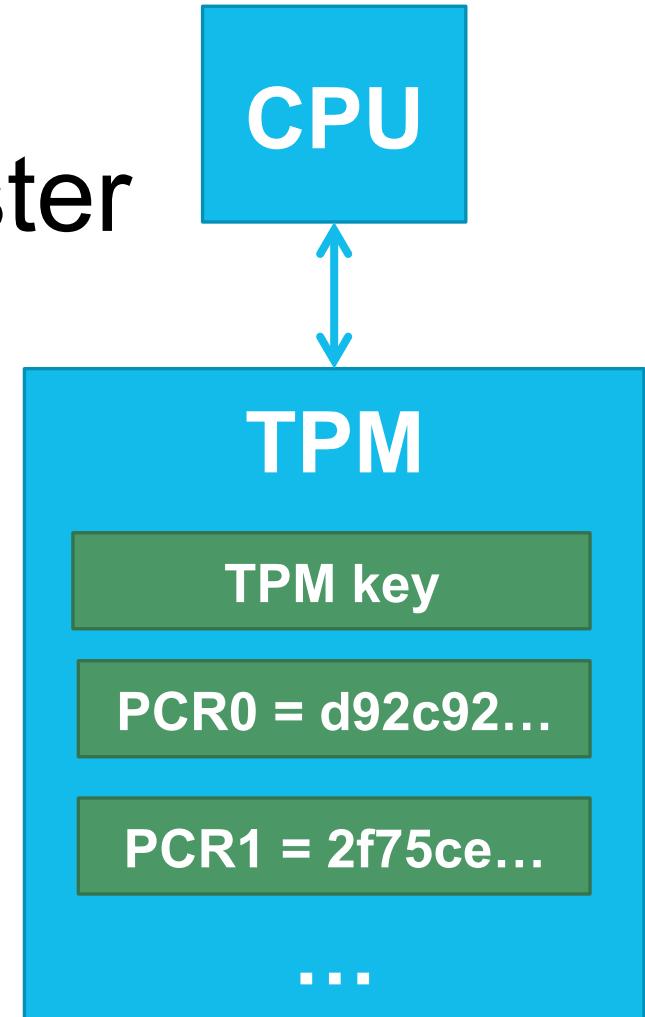
- Trusted: CPU, DRAM, BIOS, and TPM chip
- Not trusted: disk, network, ...
- TPM keeps track of what software runs, selectively uses key



TPM interface (partial)

- PCR: platform configuration register
- $\text{TPM_extend}(n, m)$:
$$\text{PCR}_n \leftarrow H(\text{PCR}_n \parallel m)$$
- $\text{TPM_seal}(n, v, \underline{\text{msg}})$:
return ciphertext
- $\text{TPM_unseal}(\text{ciphertext})$:

return msg if $\text{PCR}_n = v$



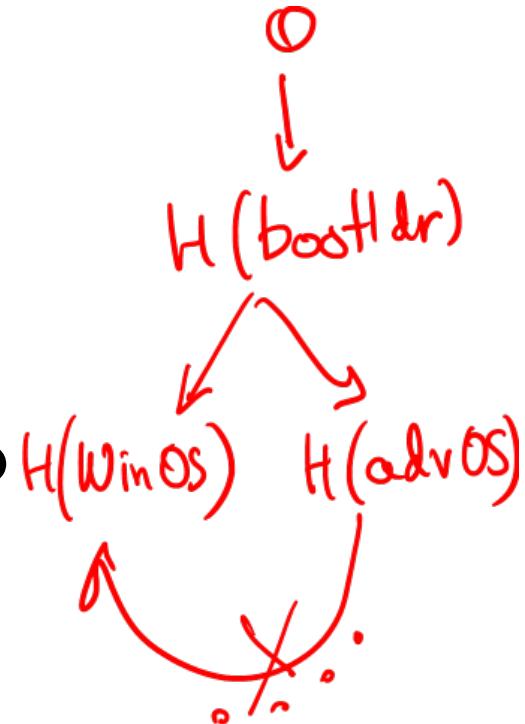
Measured boot

- PCR values reset to zero on boot
- CPU starts executing the BIOS
- Each boot step calls **TPM_extend**
 - BIOS: $\text{TPM_extend}(H(\text{bootloader}))$
 - Bootloader: $\text{TPM_extend}(H(\text{OS kernel}))$
 - OS kernel: $\text{TPM_extend}(H(\text{drivers}))$
 - ...

$$\begin{aligned} \text{PCR}_i &= 0 \\ H(0) \parallel H(\text{bootldr}) \\ H(=) \parallel H(\text{OS kern}) \\ \dots \end{aligned}$$

Sealing data using the TPM

- **TPM_seal(i , pcr , msg)**
 - Can be unsealed to recover msg only if $\text{PCR}_i = pcr$
 - Should mean running software corresponding to pcr
- What if adversary boots another OS?
 - BIOS calls `TPM_extend()` on adversary's bootloader
 - PCR value now differs from the expected sequence
 - Secure hash function: infeasible to find another way to get the same PCR value pcr via `TPM_extend()`
 - Will not be able to call `TPM_unseal()`



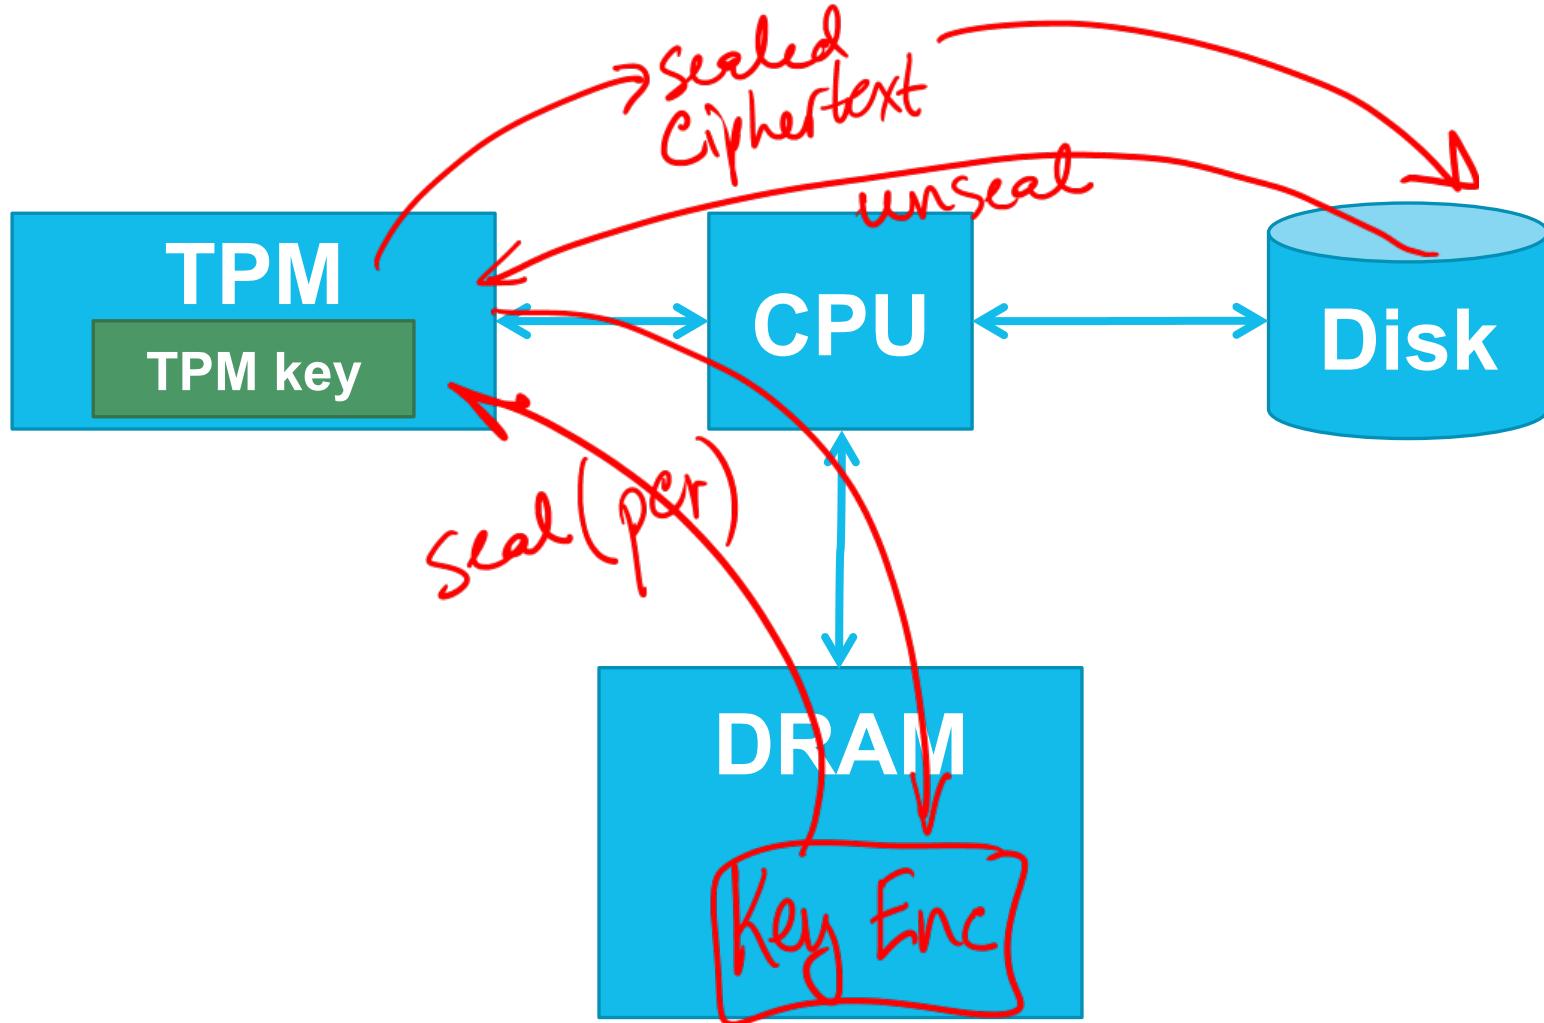
Measured boot flexibility

- Measured boot process updates PCRs to track what software was loaded
- TPM allows encrypting data that will be decrypted only under a specific PCR value
- Many benefits (and issues) of signed boot, but without hard-coding the signing key

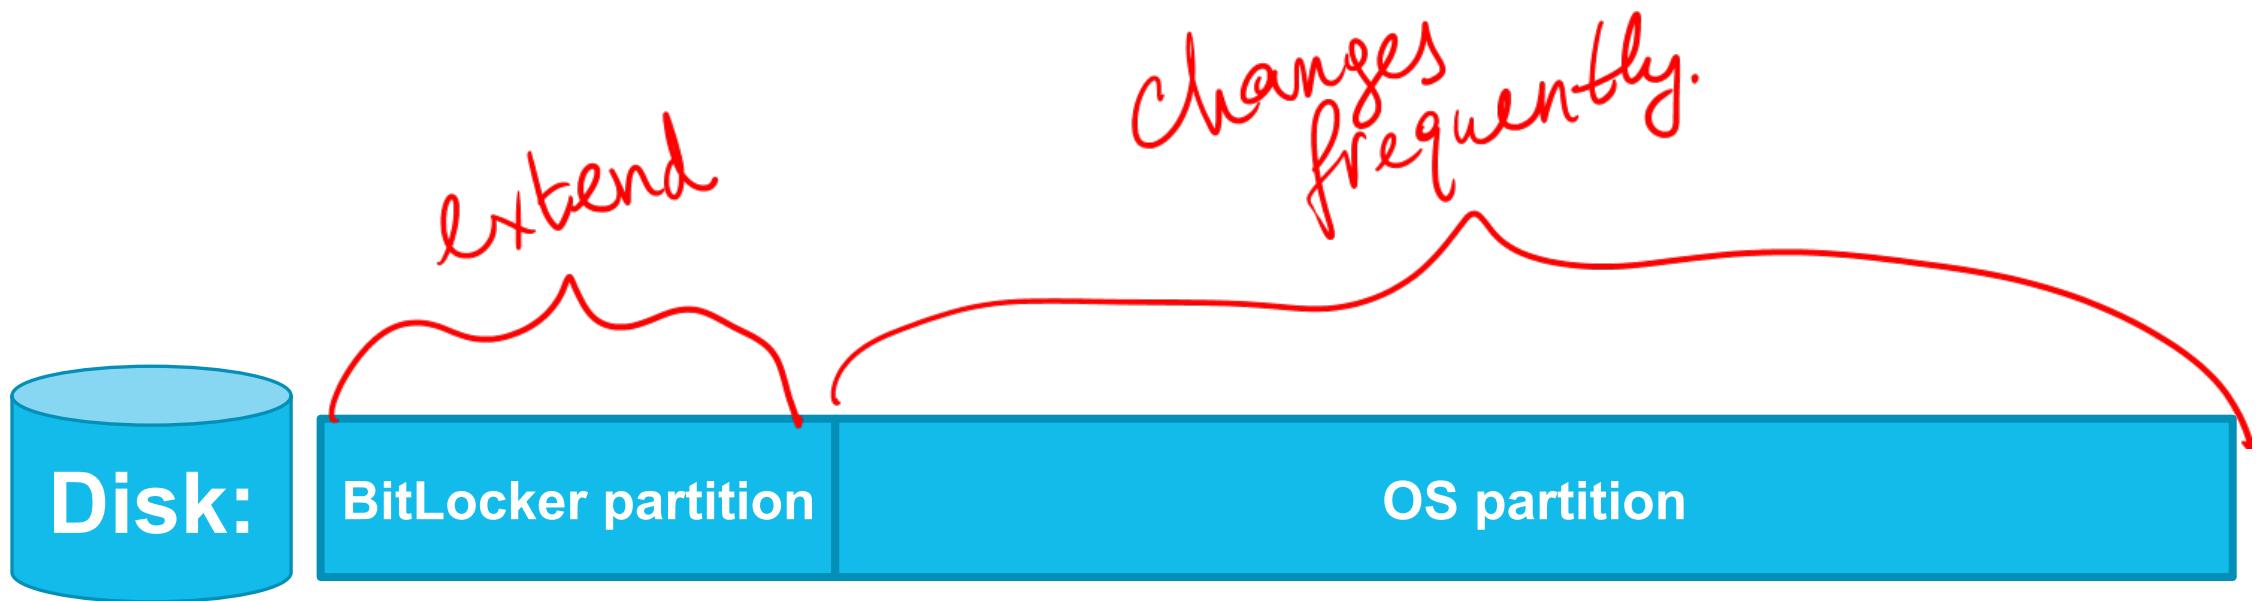
Segment 5: BitLocker keys using the TPM

- Deciding what to seal
- Deciding what to measure
- Upgrades and recovery

BitLocker seals disk encryption key

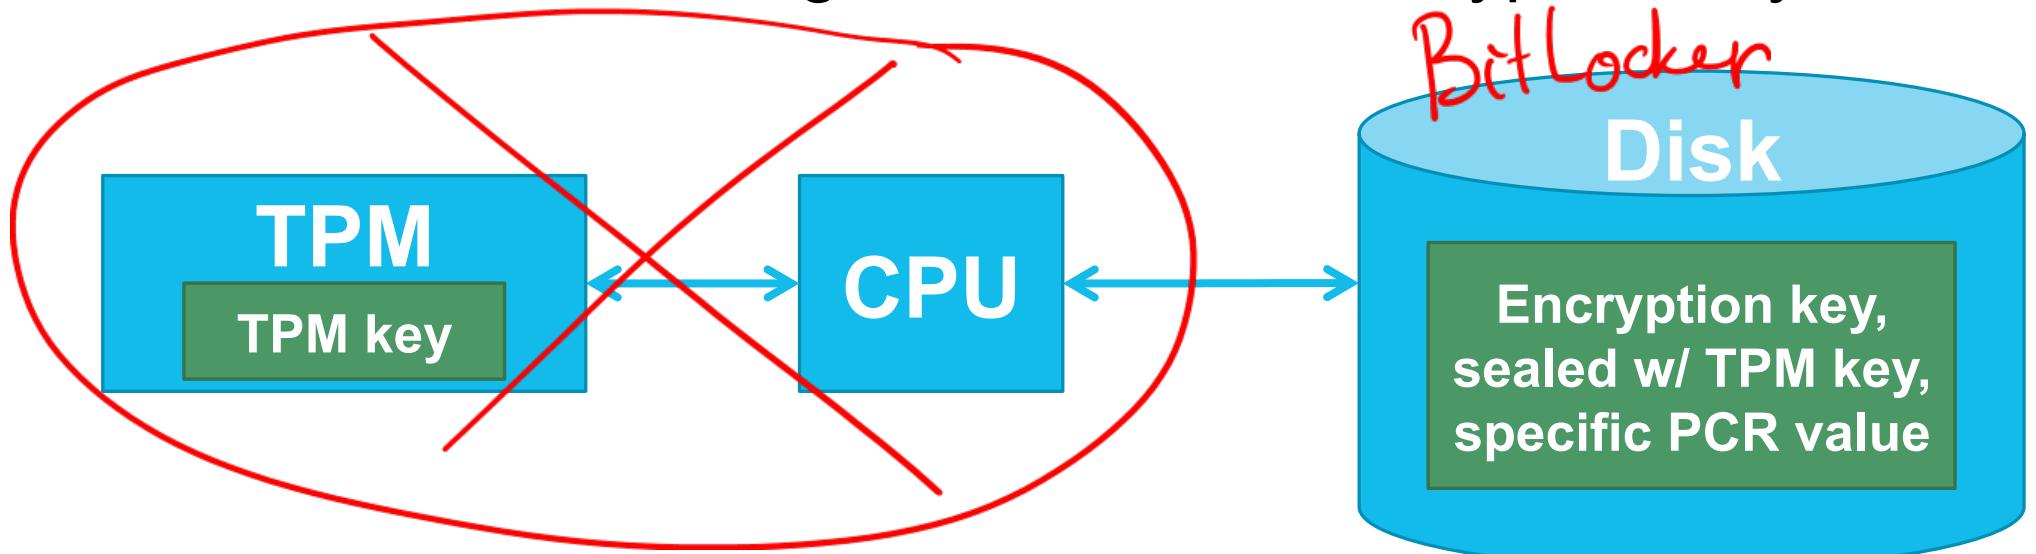


BitLocker measures small boot partition



Recovery and upgrade challenge

- Key sealed on specific TPM, for specific pcr
 - Upgrade BitLocker partition: PCR value changes
 - Upgrade or replace computer: TPM key changes
 - In both cases, can no longer unseal disk encryption key!



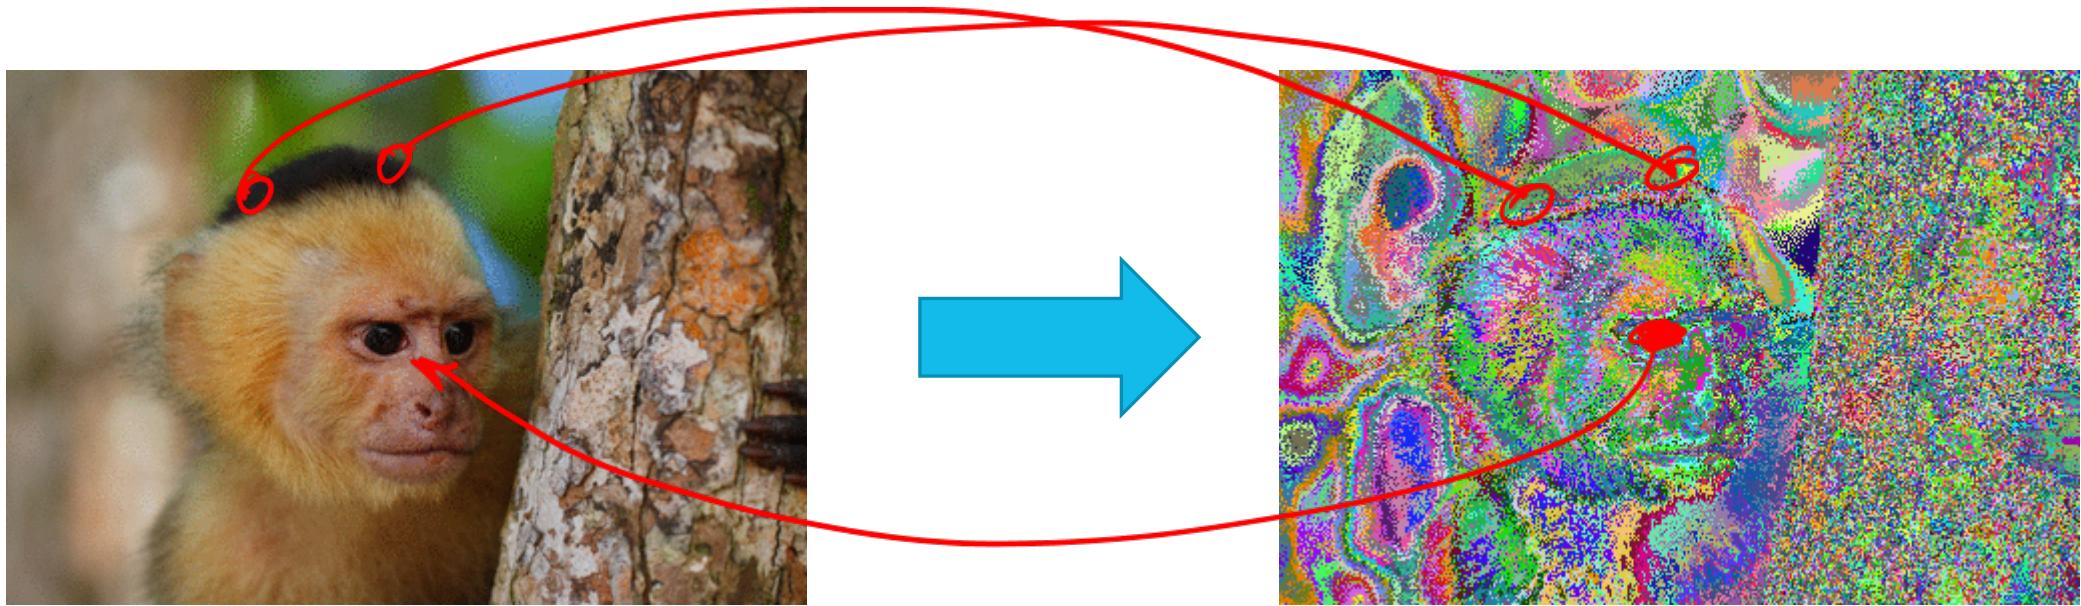
Recovery and upgrades in BitLocker

- Recovery: backup copies of disk enc. key
 - Offline USB device
 - Central server (e.g., Windows active directory)
- Upgrades: temporarily store unsealed key

Segment 6: Data encryption

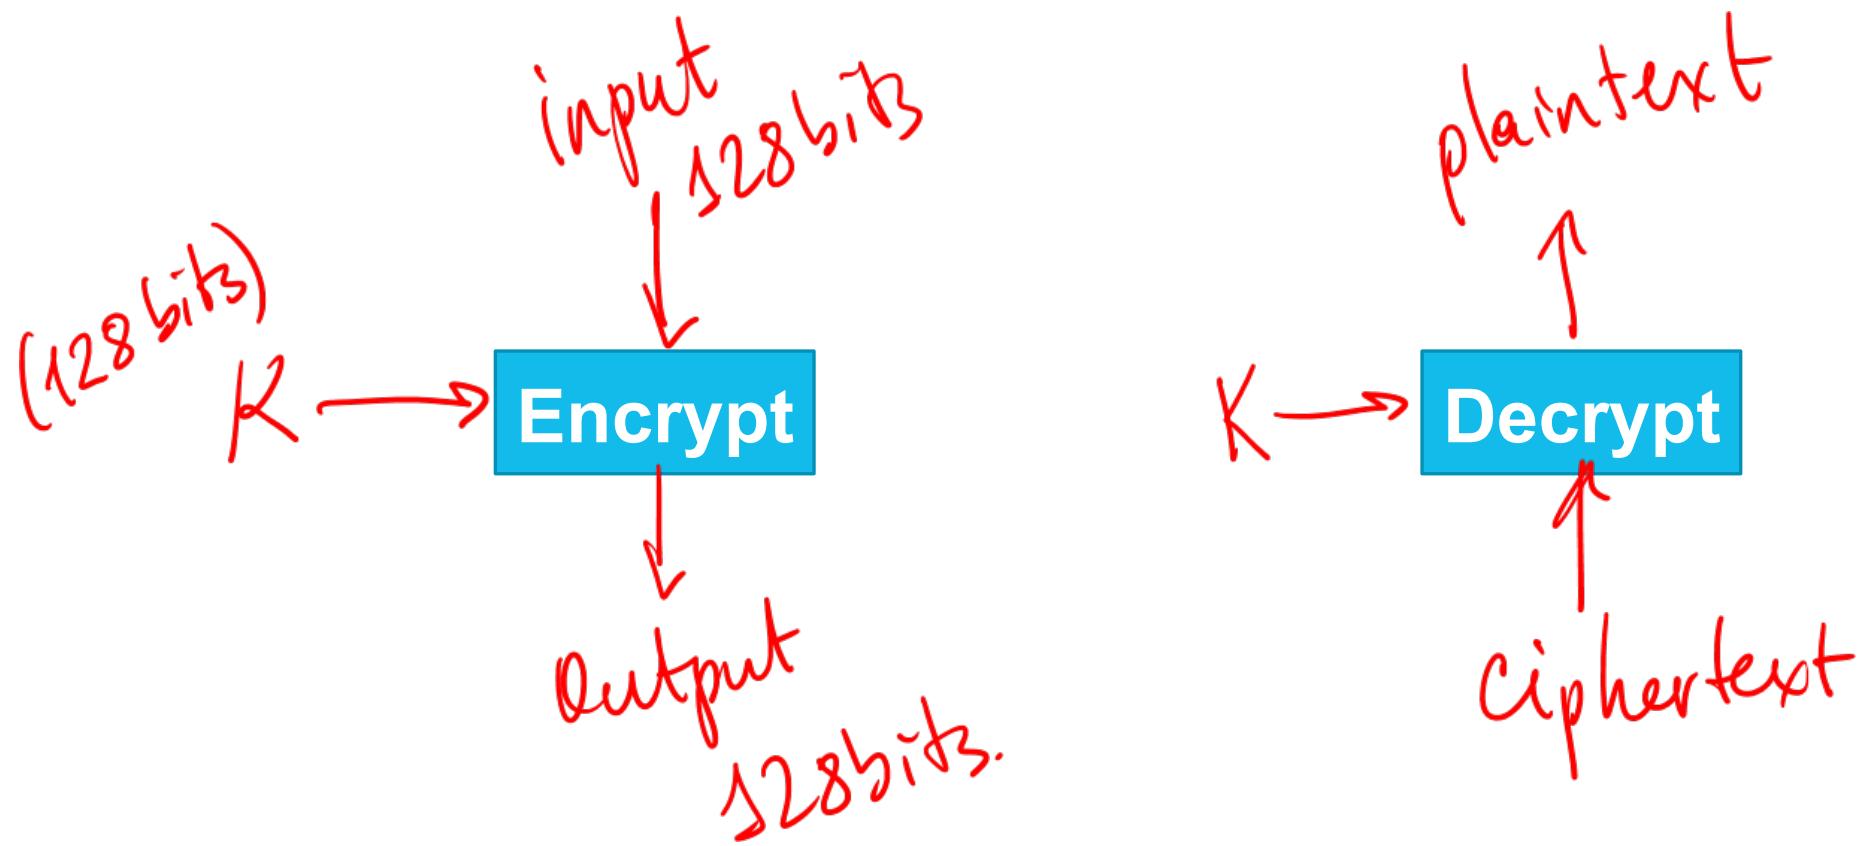
- Challenges in encrypting data
- Block ciphers
- Authentication
- Constraints of disk encryption

Problems with naïve encryption

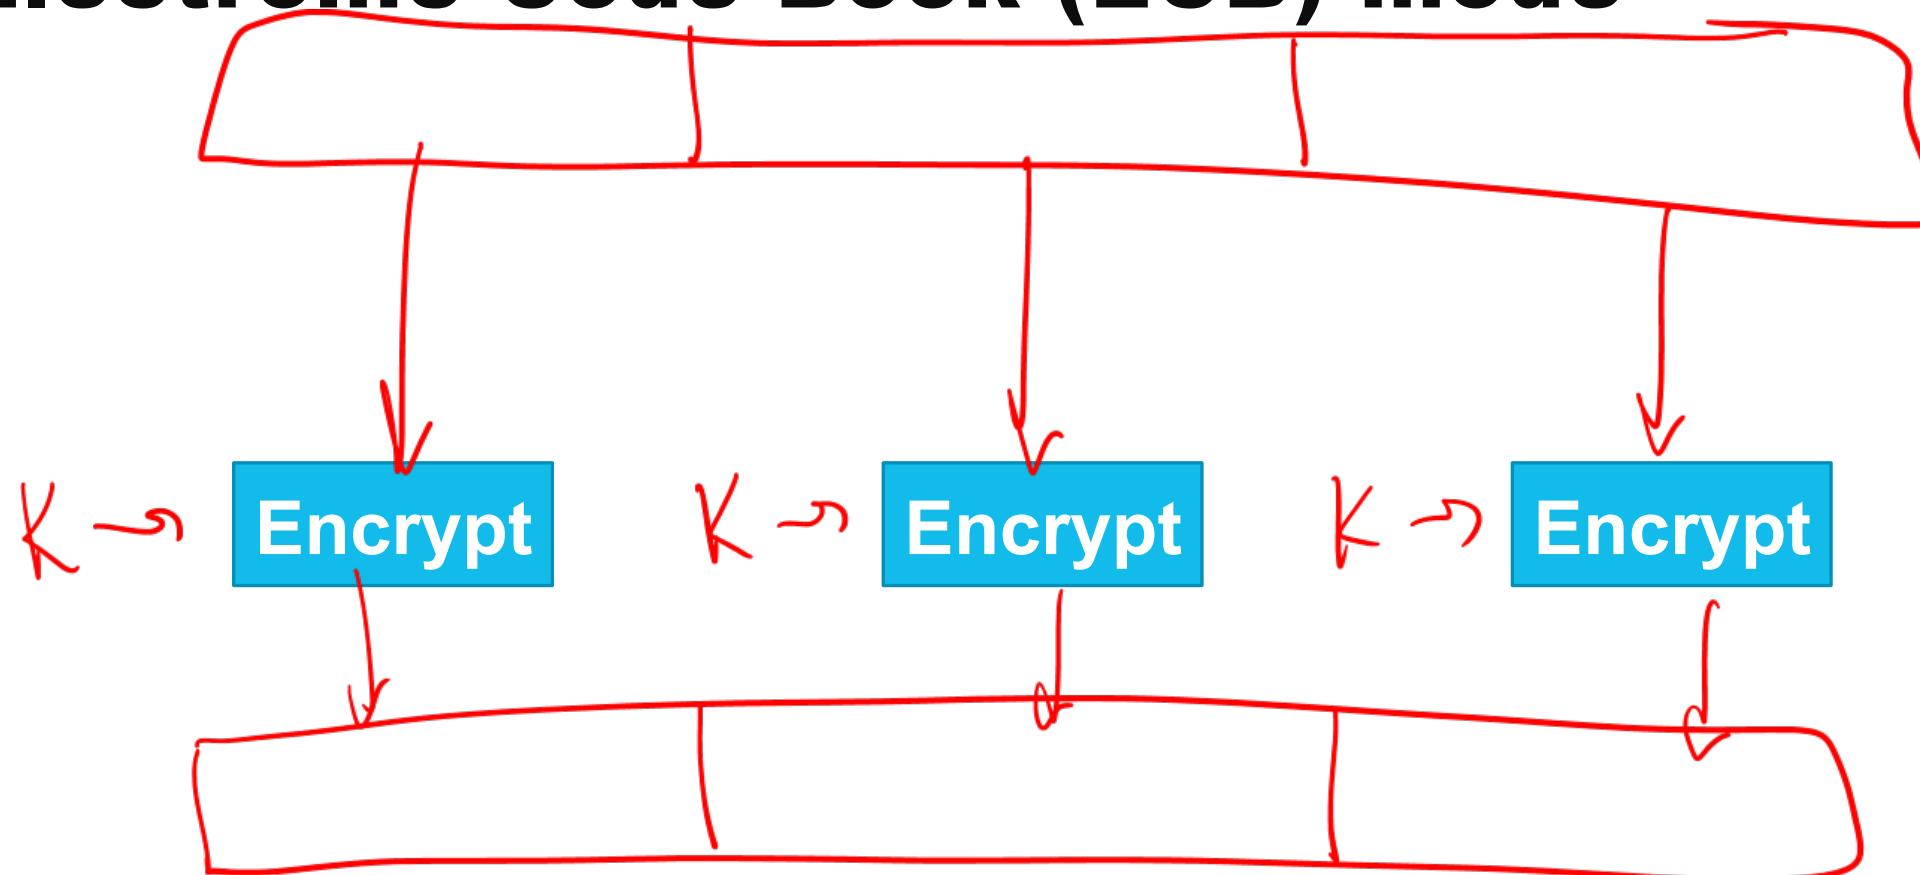


1. Deterministic
2. Malleability.

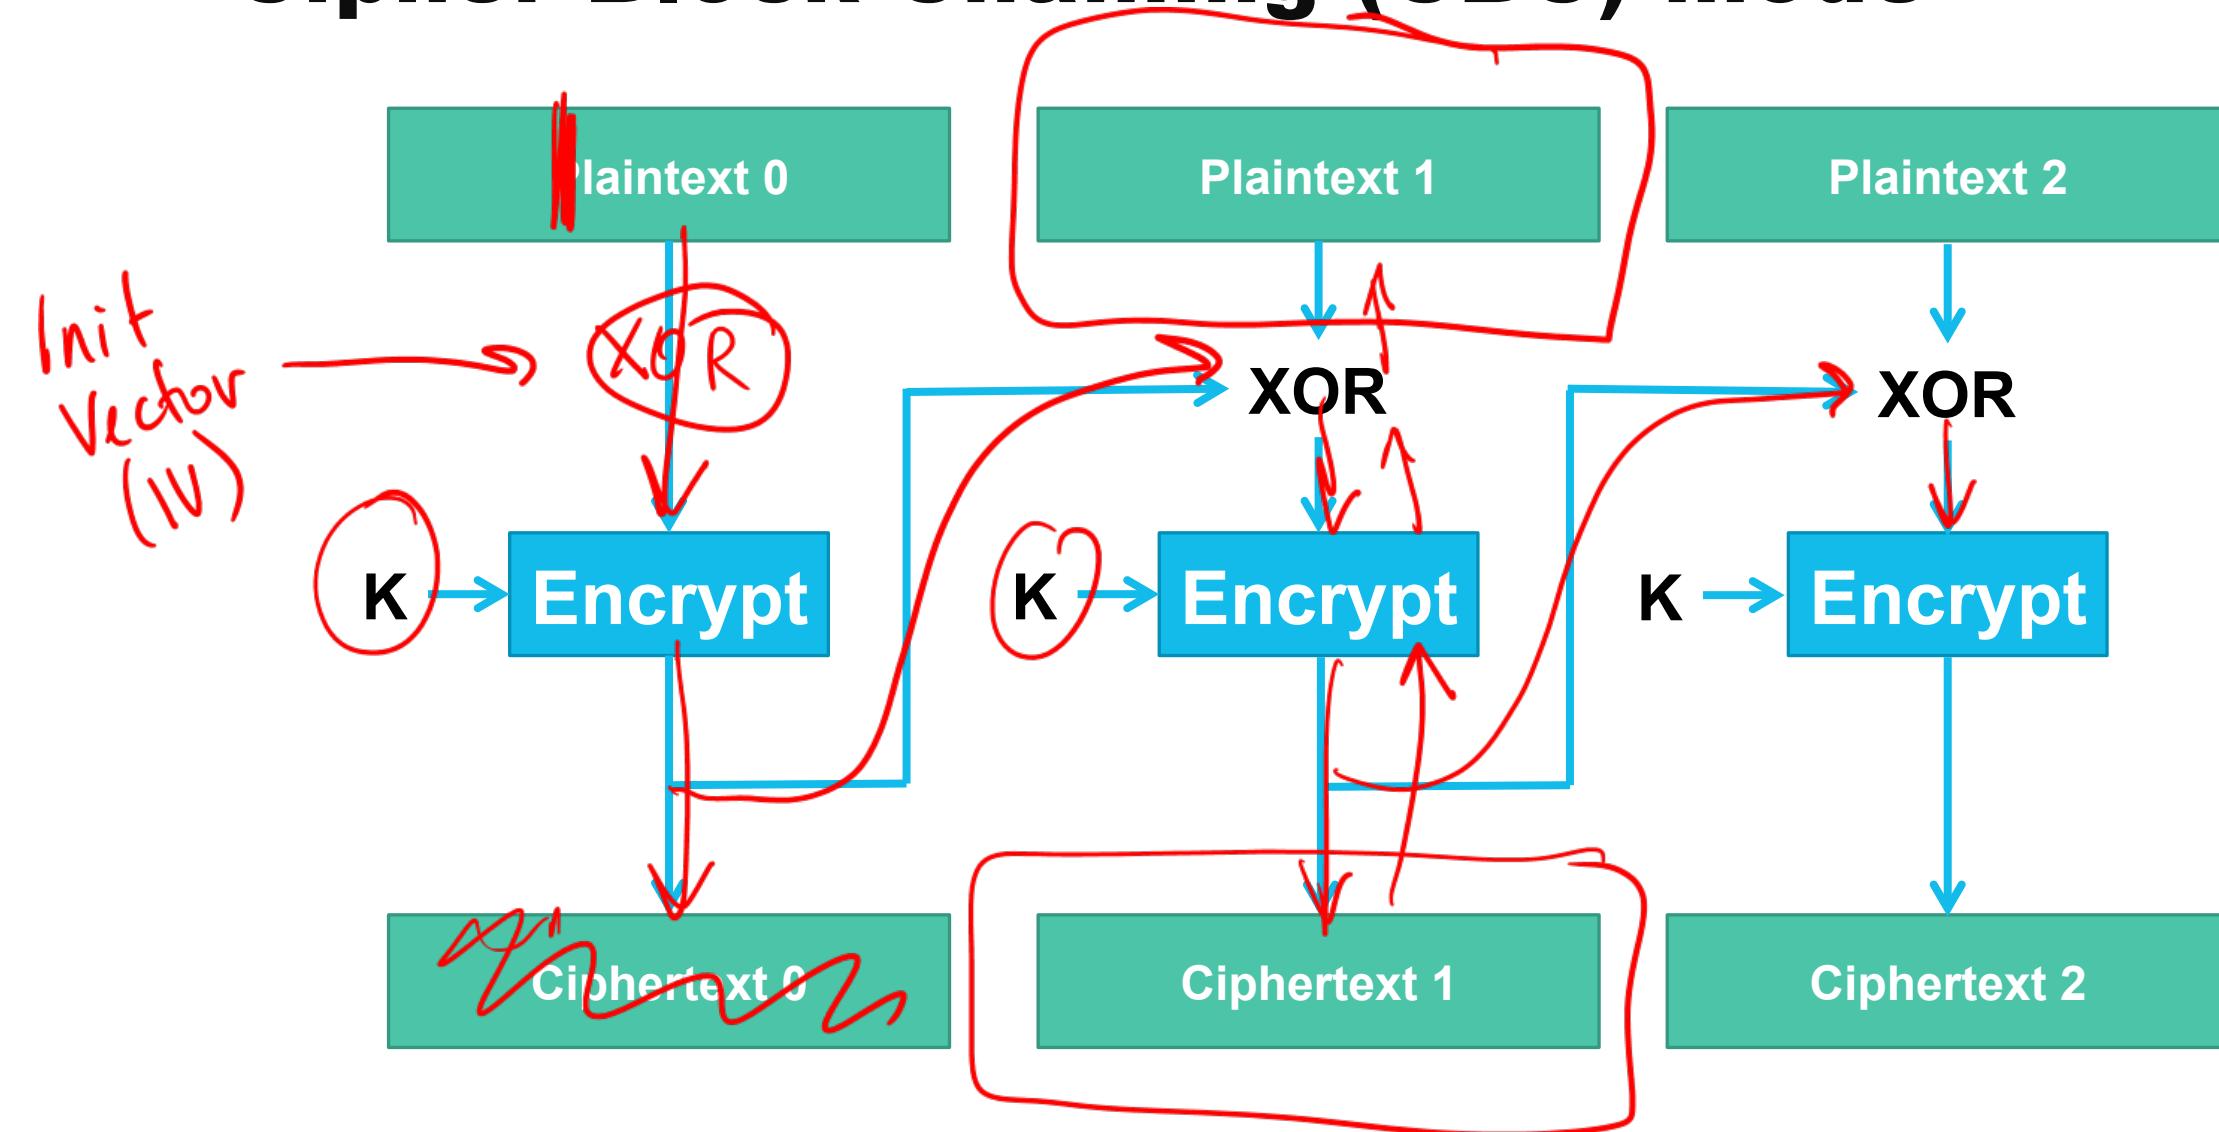
Block cipher (e.g., AES)



Encrypting more than a single AES block: Electronic Code Book (ECB) mode



Cipher Block Chaining (CBC) mode

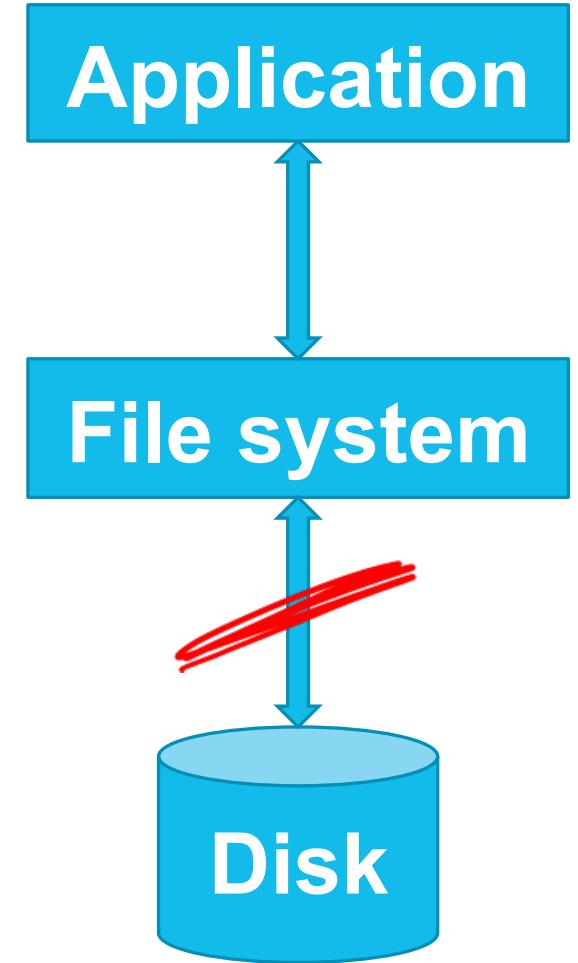


Authenticating to prevent tampering

- CBC mode: every ciphertext is “valid”
 - Cannot tell if adversary tampered with the encrypted data
- Message authentication code (MAC)
 - Rough intuition: also store $\text{MAC} = H(\text{ciphertext} \parallel \text{key})$
 - If adversary tampers with ciphertext, can’t compute H : no key
 - More efficient (and secure) constructions for authenticated encryption: UFE, OCB, CCM, ...

Constraints of disk encryption

- File system expects 4KB sectors
- Disk provides 4KB sectors
- No extra space for IV or MAC!



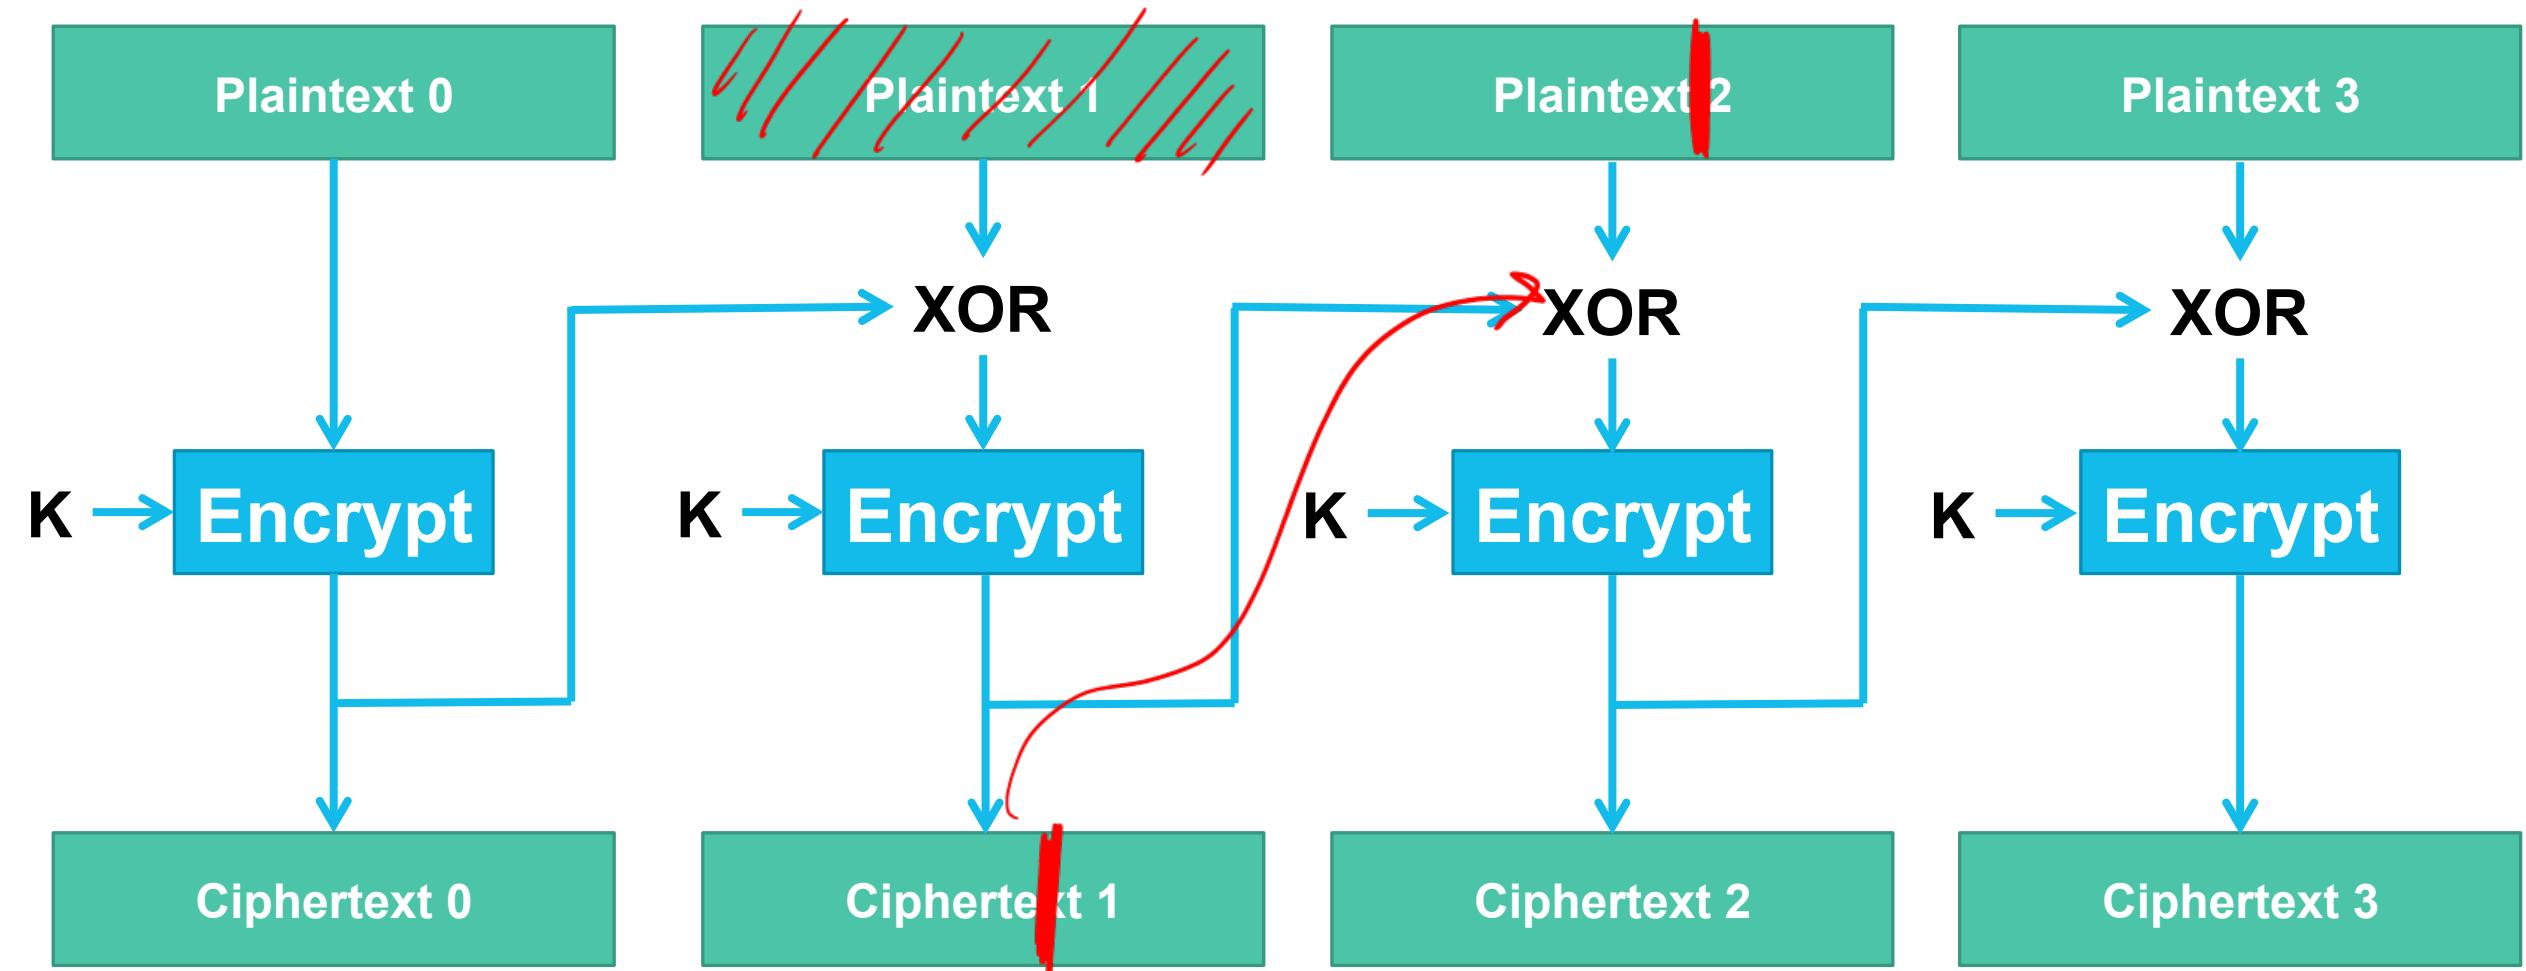
Segment 7: BitLocker data encryption

- Poor-man's authentication
- Diffusion properties of CBC
- Pseudo-random IVs
- Summary

Poor-man's authentication

- No space on disk to store per-sector MAC
- Goal: tampering with ciphertext produces “garbage” plaintext, not useful to adversary
- Good enough for our threat model
 - Adversary cannot replace login binary or password file on disk with something meaningful

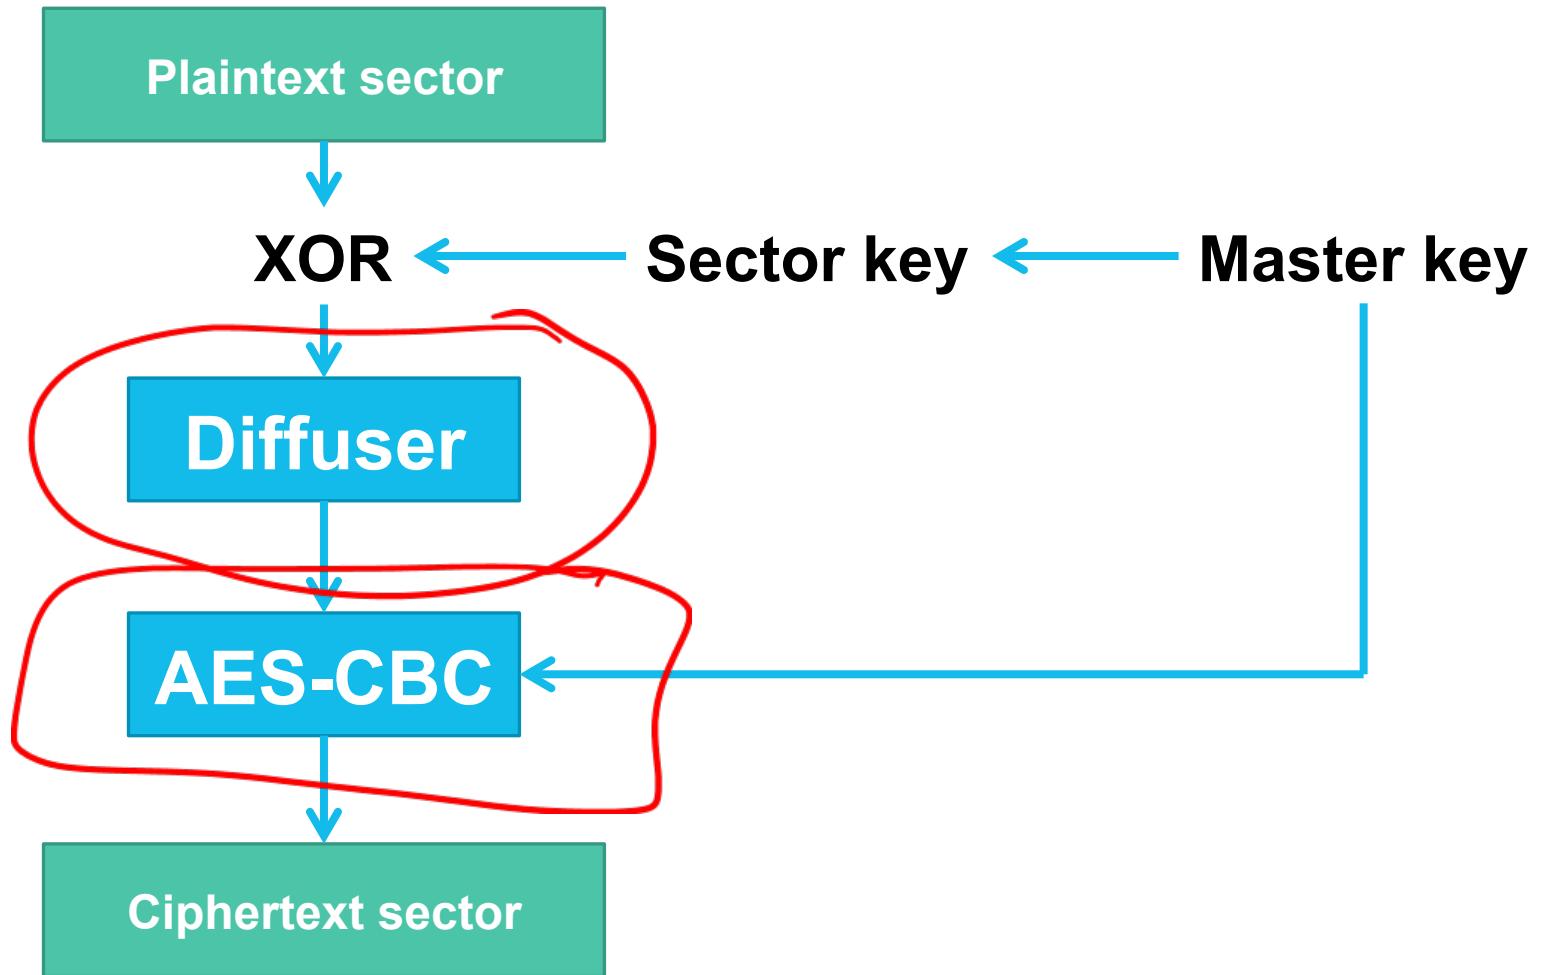
Diffusion properties of CBC



Improving diffusion with better modes

- More sophisticated block cipher modes
 - CMC, EME, XEX, ...
 - Requires two passes over the data
 - Likely the right choice now
 - Was too expensive when BitLocker was designed (2006)

BitLocker's improved diffusion



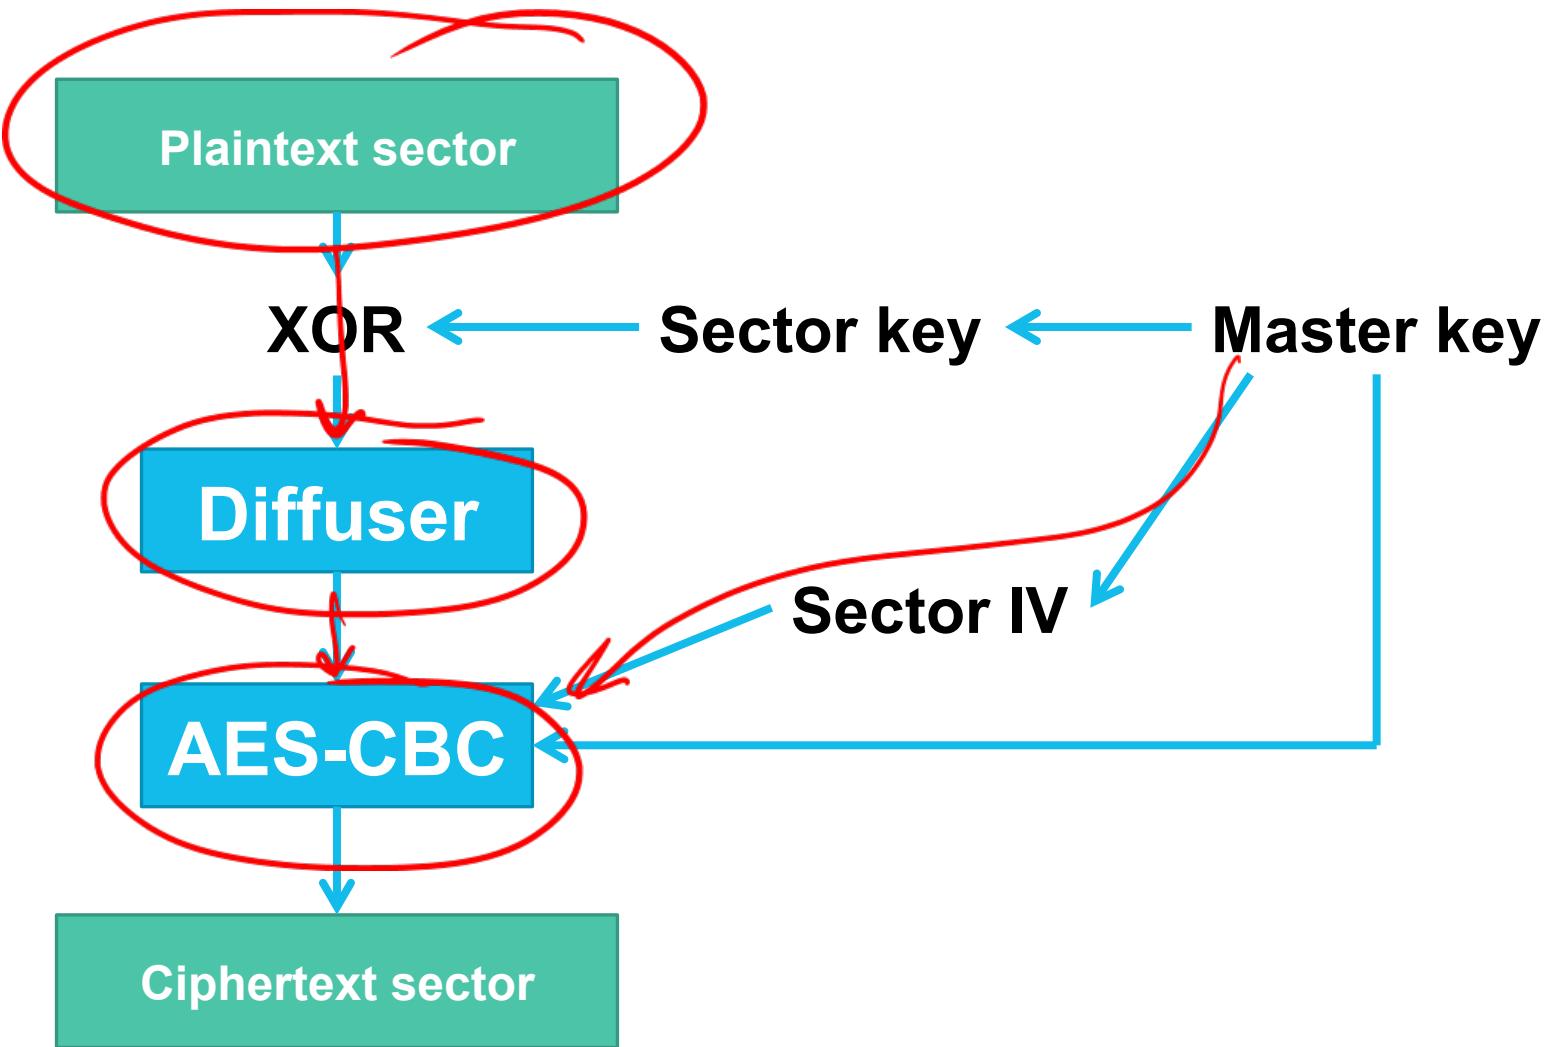
Pseudo-random IVs

- Need IV for CBC mode to randomize
- No space to store per-sector random IV
- Approach: compute pseudo-random IV

$$\text{IV}(\text{sector}\#) = \text{AES-Encrypt}(K, \text{sector}\#)$$

- Adversary doesn't know K, can't compute IV
 - One downside: repeated writes to same sector re-use IV

BitLocker's overall encryption plan



Summary

- Explored challenges in disk encryption
 - Key management, data encryption approaches, disk constraints
- Case study: Microsoft's BitLocker
 - Trade-offs to balance usability, performance, and security
 - Clever use of available TPM hardware
 - Poor-man's authentication deals with disk constraints
 - Careful deployment of new cryptographic constructions
 - Enables using existing Windows user authentication

THANK YOU

Nickolai Zeldovich

Associate Professor

